

# The perfectly synchronized round-based model of distributed computing

Carole Delporte-Gallet <sup>a</sup>, Hugues Fauconnier <sup>a</sup>, Rachid Guerraoui <sup>b</sup>,  
Bastian Pochon <sup>b,\*</sup>

<sup>a</sup>*LIAFA Institute, Université Denis Diderot, F-75251 Paris 5, France*

<sup>b</sup>*School of Computer and Communication Sciences, EPFL, CH-1015 Lausanne, Switzerland*

Received 9 February 2004; revised 5 July 2006

Available online 27 December 2006

---

## Abstract

The perfectly synchronized round-based model provides the powerful abstraction of crash-stop failures with atomic and synchronous message delivery. This abstraction makes distributed programming very easy. We describe a technique to automatically transform protocols devised in the perfectly synchronized round-based model into protocols for the crash, send omission, general omission or Byzantine models. Our transformation is achieved using a round shifting technique with a constant time complexity overhead. The overhead depends on the target model: crashes, send omissions, general omissions or Byzantine failures. Rather surprisingly, we show that no other automatic non-uniform transformation from a weaker model, say from the traditional crash-stop model (with no atomic message delivery), onto an even stronger model than the general-omission one, say the send-omission model, can provide a better time complexity performance in a failure-free execution.

© 2006 Elsevier Inc. All rights reserved.

**Keywords:** Abstraction; Simulation; Distributed systems; Fault-tolerance; Synchronous system models; Complexity

---

---

\* Corresponding author. Present address: EPFL IC IIF LPD, Station 14, CH-1015 Lausanne, Switzerland.  
*E-mail address:* Bastian.Pochon@epfl.ch (B. Pochon).

## 1. Introduction

### 1.1. Motivations

Distributed programming would be easy if one could assume a *perfectly synchronized round-based model* where the processes would share, after every round, the same view of the distributed system state. Basically, computation would proceed in a round-by-round way, with the guarantee that, in every round, a message sent by a correct process is received by all processes, and a message sent by a faulty process is either received by all or by none of the processes. All processes that reach the end of a round would have the same view of the system state.

Unfortunately for the programmers, and fortunately for the distributed computing research community, the assumption that all processes have the same view of the system state does not hold in practice. In particular, the illusion of a perfectly synchronized world breaks because messages sent over a network might be subject to partial delivery or message loss, typically because of a buffer overflow at a router, or due to a crash failure, resulting from the crash of some computer hosting processes involved in the distributed computation.

It is of course legitimate to figure out whether we could provide the programmer with the simple view of a perfectly synchronized world, and translate, behind the scenes, distributed protocols devised in such an ideal model into more realistic and weaker models. After all, the job of a computer scientist is usually about providing programming abstractions that hide low level details, so why not try to provide those that really facilitate the job of the programmer of distributed applications.

The very fact that the abstraction of a perfectly synchronized round-based model has not already been made available to programmers through popular programming middleware, even after several decades of research in distributed computing, might indicate that its implementation might turn out to be significantly involved. Indeed, a closer look at the semantics of the perfectly synchronized round-based (PSR) model reveals that what needs to be implemented is actually a succession of instances of an agreement algorithm, more precisely an algorithm solving the interactive consistency (IC) problem [18]. Indeed, this is the key to provide processes with the same view of the system at the end of every round. Roughly speaking, in the IC problem, each process is supposed to propose a value and eventually decide on a vector of values, such that the following properties are satisfied: termination (i.e., every correct process eventually decides on a vector), validity (i.e., the  $j$ th component of any decided vector by a correct process is the value proposed by process  $p_j$  if  $p_j$  is correct), and agreement (i.e., no two correct processes decide on different vectors).

The relationship between interactive consistency and the perfectly synchronized round-based model highlights two issues. The first has to do with feasibility. On the one hand, to implement the PSR abstraction over a given model, one needs to make some synchrony assumptions on the model (e.g., the interactive consistency problem is not solvable in an eventually synchronous model [10]), and the coverage of these assumptions might simply not be sufficient for certain distributed environments. The second issue has to do with performance. Even when the PSR abstraction can be implemented, the cost of its implementation might be too high. That is, devising a distributed protocol over PSR, and relying on the implementation of PSR to automatically generate a distributed protocol in a weaker model might have a significant overhead with respect to devising the protocol directly in the latter model.

The lack of any evidence concerning the exact overhead of implementing the PSR abstraction was the motivation of this work. More precisely, the motivation was to figure out whether we can come up with an efficient implementation, in terms of time complexity, of the PSR abstraction over synchronous round-based models with various types of failures, ranging from simple crash failures [12] to more general Byzantine failures [18,15], including send-omissions [12] and general-omissions [16].

## 1.2. Background

The PSR abstraction is known to have implementations in all the models mentioned above, but the inherent cost of these implementations in either of these models was unclear. The lack of any result on the cost of implementing PSR might seem surprising given the amount of work that was devoted to devising optimal agreement algorithms over various models, including the omission model and the Byzantine model.

(1) In particular, we do know that, in terms of round complexity, there is a tight lower bound on implementing interactive consistency in a synchronous round-based model where  $t$  processes can crash is  $t + 1$  [7]. The result is derived for the model with crashes, and thus also holds for send-omissions, general-omissions, and Byzantine failures. The result says that  $t + 1$  rounds of, say, the general-omission model are needed for all correct processes to reach a decision about the new global state of the distributed system (i.e., the decision vector). If, pretty much like in state machine replication [13,21], we implement PSR as a sequence of instances of interactive consistency, then the  $t + 1$  cost would add up. In other words,  $K(t + 1)$  rounds would be needed to implement  $K$  rounds of PSR.

One might wonder whether algorithms that are early deciding [14,5] would decrease this cost. Indeed, these algorithms need fewer rounds for processes to decide when only  $f$  failures occur, out of the total number  $t$  of failures that are tolerated. These algorithms however do not guarantee a simultaneous decision from all the processes [4], even from the correct processes only. In such a case, it would then be necessary to delay the simulation of the next PSR round until each process reaches the next multiple of  $t + 1$  rounds. In other words,  $K(t + 1)$  rounds would again be needed to implement  $K$  rounds of PSR.

(2) Implementing a synchronous round-based model with crash failures [12] (*crash-stop* model) over various weaker models, such as the omission model, has been the subject of several investigations, e.g., [1,17]. These can be viewed as implementing an abstraction that is weaker than PSR. (PSR prevents a message from being received by some but not all the processes, whereas the crash-stop model does not, in case the sender crashes.) The idea underlying the implementation proposed in [17], for instance for the omission model, is that of *doubling* rounds. Roughly speaking, any round of the crash-stop model is simulated with two rounds of the omission model. Hence,  $2K$  rounds of the omission model are needed to simulate  $K$  rounds of the crash-stop model.

In the first case where we use a sequence of interactive consistency instances, or in the second case where we mask failures by doubling rounds, we end up with multiplicative factor overheads, and even if we try to implement the weaker crash-stop abstraction along the lines of [17]. In fact, if we implement PSR directly on the crash-stop model (used as an intermediate model), and use the transformation of [17], we end up with a cost of  $K(f + 1)$  rounds of the omission model for  $K$  rounds

of the PSR model with  $f$  actual failures. Is this multiplicative factor inherent to implementing PSR over an omission model? Or could we devise a *shifting* implementation with an additive factor, i.e.,  $K + S$  with  $S$  a constant? At first glance, this would be counter-intuitive because it would mean devising a more efficient implementation than [17] for an abstraction that is strictly stronger.

### 1.3. Contributions

This paper presents a time-efficient shifting technique to implement the PSR abstraction over the synchronous round-based message-passing models with crash failures, send-omissions, general-omissions and Byzantine failures (for  $t < n/3$ ):  $K$  rounds of PSR require at most  $K + t$  rounds of the model with more severe failures, when  $t$  failures are tolerated. That is, with an additive factor  $S = t$ . This is clearly optimal because PSR solves interactive consistency in one round, and this costs at least  $t + 1$  rounds in either model (with crash failures, send or general omissions, or Byzantine failures) [7]. In other words, any shifting transformation technique from the PSR model to the omission model has to pay the cost of  $t$  additional rounds.

This paper gives both a uniform and a non-uniform shifting transformation. Intuitively speaking, a uniform transformation ensures that any process, be it correct or faulty, simulates a correct behavior according to the original algorithm (i.e., the algorithm to be transformed), or nothing at all, whereas a non-uniform transformation does not ensure faulty processes simulate a correct behavior with respect to the original algorithm. For both transformations, we make the details clear about the underlying interactive consistency algorithms that are used, respectively, uniform and non-uniform. Our shifting transformations do not necessarily require that all processes decide simultaneously within each underlying interactive consistency instance, hence the use of early deciding algorithms is possible. By considering an early deciding non-uniform interactive consistency algorithm, we show that our shifting transformation works in “real-time” in a failure-free execution, i.e., the transformed algorithm executes as fast as the original algorithm. In this precise case, it is clear that the transformation is optimal, since  $K$  round of PSR are transformed into  $K$  rounds of the target model.

We precisely define the general notion of transformation and then describe our novel shifting transformation technique. Beforehand, we introduce the necessary machinery to formulate the definitions of simulation and transformation. The key idea of our technique is that a round in the weak model (crash, send or general omission, Byzantine failure), is involved in the simulation of more than one round of PSR. This is also the source of some tricky algorithmic issues that we had to address.

## 2. Model

### 2.1. Processes

We consider a finite set  $\Omega$  of  $n$  processes  $\{p_0, \dots, p_{n-1}\}$ , that communicate by point-to-point message-passing. We assume that processes are fully connected. A process is characterized by its *local state* and we denote by  $S$  the set of possible states of any process. Processes interact in a synchronous, round-based computational way. Let  $\mathcal{R} = \mathbb{N}^*$  be the set of round numbers (strictly positive, integer

numbers). We denote by  $\mathcal{M}$  the set of messages that can be sent, and by  $\mathcal{M}' = \mathcal{M} \cup \{\perp\}$  the set of messages that can be received.  $\perp$  is a special value that indicates that no message has been received. The primitive `send()` allows a process to send a message to the processes in  $\Omega$ . The primitive `receive()` allows a process to receive a message sent to it that it has not yet received. We assume that each process receives an input value from the external world, at the beginning of every round, using the primitive `receiveInput()`. We denote by  $\mathcal{I}$  the set of input values that can be received, for all processes. An *input pattern* is a function  $I : \Omega \times \mathcal{R} \rightarrow \mathcal{I}$ . For any given process  $p_i$  and round number  $r$ ,  $I(i, r)$  represents the input value that  $p_i$  receives at the beginning of round  $r$ . For any given set of input values  $\mathcal{I}$ , we denote by  $\Gamma_{\mathcal{I}}$  the set composed of all input patterns over  $\mathcal{I}$ . For the sake of simplicity, we assume that input values do not depend on the state of processes. In Section 5, we discuss an extension where this assumption is relaxed.

Roughly speaking, in each synchronous round  $r$ , every process goes through four, non-atomic steps (in particular, the processes do not have any atomic broadcast primitive for executing the second step). In the first step, the process receives an external input value. In the second step, the process sends the (same) message to all processes (including itself). In the third step, the process receives all messages sent to it. The fourth step is a local computation to determine the next local state of the process.

Throughout the paper, if a variable  $v$  appears in the local state of all processes, we denote by  $v_i$  the variable at process  $p_i$ , and by  $v_i^r$  the value of  $v$  after  $p_i$  has executed round  $r$ , but before  $p_i$  has started executing round  $r + 1$ . For convenience of notation,  $v_i^0$  denotes the value of  $v$  at process  $p_i$  after initialization, before  $p_i$  takes any step.

## 2.2. Protocols

The processes execute a *protocol*  $\Pi = \langle \Pi_0, \dots, \Pi_{n-1} \rangle$ . Each process  $p_i$  executes a state machine  $\Pi_i$ , defined as a triple  $\langle s_i, T_i, O_i \rangle$ , respectively, an initial state, a state transition function and a message output function. We assume that, at any process  $p_i$ , the corresponding state machine is initialized to  $s_i$ . The message output function  $O_i : \mathcal{S} \times \mathcal{I} \times \mathcal{R} \rightarrow \mathcal{M}$  generates the message to be sent by process  $p_i$  during round  $r$ , given its state at the beginning of round  $r$ , an external input value, and the round number. Note that, throughout this paper, we assume for presentation simplicity that processes always have a value to send, and we reserve the symbol  $\perp$  for the very case where a message is not received, as the result of a failure. The state transition function  $T_i : \mathcal{S} \times (\mathcal{M}')^n \times \mathcal{R} \rightarrow \mathcal{S}$  outputs the new state of process  $p_i$ , given the current state of  $p_i$ , the messages received during the round from all processes (possibly  $\perp$  if a message is not received) and the current round number.

We introduce three functions for describing whether the execution of any protocol by any process is correct or deviate from the one intended. In the following functions,  $N$  denotes the interval of integer values  $[1, n]$ , corresponding to process identifiers in  $\Omega$ .

- $ST : N \times \mathcal{R} \cup \{0\} \rightarrow \mathcal{S}$  is a function such that, for any process  $p_i$  and round  $r$ ,  $ST(i, r)$  is the state of process  $p_i$  at the end of round  $r$ . (Slightly abusing the notation, we define  $ST(i, 0) = s_i$  for any process  $p_i$ .)
- $MS : N \times N \times \mathcal{R} \rightarrow \mathcal{M}'$  is a function such that, for any processes  $p_i, p_j$  and round  $r$ ,  $MS(i, j, r)$  is the message sent by  $p_i$  to  $p_j$  in round  $r$ , or  $\perp$  if  $p_i$  fails to send a message to  $p_j$  in round  $r$ .

- $MR : N \times N \times \mathcal{R} \rightarrow \mathcal{M}'$  is a function such that, for any processes  $p_i, p_j$  and round  $r$ ,  $MR(i, j, r)$  is the message received by  $p_i$  from  $p_j$  in round  $r$ , or  $\perp$  if  $p_i$  fails to receive a message from  $p_j$  in round  $r$ . In the following,  $MR(i, r)$  denotes the vector of all the messages received by  $p_i$  in round  $r$ , i.e.,  $MR(i, r) = [MR(i, 1, r), \dots, MR(i, n, r)]$ .

### 2.3. Correctness

When we make no assumption whatsoever about the behavior of any process  $p_i$ , we consider that  $p_i$  behaves correctly, i.e.,  $p_i$  follows the state machine  $\Pi_i$  assigned to it. Here we define the correct behavior of any process  $p_i$  more formally.

Any process  $p_i$  is *correct up to round  $r$* ,  $r \geq 1$ , if for any  $r'$ ,  $1 \leq r' \leq r$ , and any input pattern  $I$ :

- $p_i$  does not fail in sending its message

$$(\forall p_j \in \Omega)(MS(i, j, r') = O_i(ST(i, r' - 1), I(i, r'), r')),$$

- $p_i$  does not fail in receiving any message

$$(\forall p_j \in \Omega)(MR(i, j, r') = MS(j, i, r')),$$

- $p_i$  makes a correct state transition ( $p_i$  does not crash)

$$ST(i, r') = T_i(ST(i, r' - 1), MR(i, r'), r').$$

By definition, any process is correct up to round 0.

### 2.4. Failures

If any process  $p_i$  does not follow the state machine  $\Pi_i$  assigned to  $p_i$  in any round  $r$ , i.e.,  $p_i$  is correct up to round  $r - 1$  and is not correct up to round  $r$ ,  $p_i$  is faulty in round  $r$  and may fail by either of the following types of failure. (For the sake of clarity, we indicate here the complete behavior of  $p_i$  in round  $r$ , not only the faulty part.)

*Atomic failure.* A process  $p_i$  that commits an atomic failure in round  $r$  can either crash before sending its message to all or after sending its message to all in round  $r$ . Processes do not recover after an atomic failure: a process that crashes due to an atomic failure in round  $r$  does not send nor receive any message in any subsequent round  $r' > r$ . More formally,

- $p_i$  either crashes before sending any message:
  - $p_i$  does not send any message to any process

$$(\forall p_j \in \Omega)(MS(i, j, r) = \perp),$$

- $p_i$  does not receive any message from any process

$$(\forall p_j \in \Omega)(MR(i, j, r) = \perp),$$

- or  $p_i$  crashes after sending a message to all and before receiving any message:
  - $p_i$  sends a message to all processes

$$(\forall p_j \in \Omega)(MS(i, j, r) = O_i(ST(i, r - 1), I(i, r), r)),$$

- $p_i$  does not receive any message from any process

$$(\forall p_j \in \Omega)(MR(i, j, r) = \perp).$$

In either cases,  $p_i$  does not perform any step after crashing:

- $p_i$  does not send nor receive any message

$$(\forall r' > r)(\forall p_j \in \Omega)(MS(i, j, r') = \perp \wedge MR(i, j, r') = \perp),$$

- $p_i$  does not perform any state transition

$$(\forall r' \geq r)(ST(i, r') = ST(i, r - 1)).$$

*Crash failure.* A process  $p_i$  that commits a crash failure in a round  $r$ —or simply that crashes in round  $r$ —can either (i) send a message to a subset of the processes, crash, not receive any message, or (ii) send a message to all, receive a subset of the messages sent to it, and crash. Processes do not recover after crashing: a process that crashes in round  $r$  does not send nor receive any message in any subsequent round  $r' > r$ . More formally,

- $p_i$  either sends its message to a subset of the processes, crashes, and does not receive any message:
  - $p_i$  sends its message to  $p_j$  or nothing at all

$$(\forall p_j \in \Omega)(MS(i, j, r) = O_i(ST(i, r - 1), I(i, r), r) \vee MS(i, j, r) = \perp),$$

- $p_i$  does not receive any message

$$(\forall p_j \in \Omega)(MR(i, j, r) = \perp),$$

- or  $p_i$  sends its message to all processes, receives the message from a subset of the processes, and crashes:

- $p_i$  sends its message to all processes

$$(\forall p_j \in \Omega)(MS(i, j, r) = O_i(ST(i, r - 1), I(i, r), r)),$$

- $p_i$  receives a message from a subset of the processes

$$(\forall p_j \in \Omega)(MR(i, j, r) = MS(j, i, r) \vee MR(i, j, r) = \perp).$$

In either cases,  $p_i$  does not perform any step after crashing:

- $p_i$  does not send nor receive any message

$$(\forall r' > r)(\forall p_j \in \Omega)(MS(i, j, r') = \perp \wedge MR(i, j, r') = \perp),$$

- $p_i$  does not perform any state transition

$$(\forall r' \geq r)(ST(i, r') = ST(i, r - 1)).$$

*Send-omission failure.* A process  $p_i$  that commits a send-omission in a round  $r$  fails to send its message in that round to a subset of processes in the system. More formally

$$(\forall p_j \in \Omega)(MS(i, j, r) = O_i(ST(i, r - 1), I(i, r), r) \vee MS(i, j, r) = \perp) \wedge (\exists p_j \in \Omega)(MS(i, j, r) = \perp).$$

*Receive-omission failure.* A process  $p_i$  that commits a receive-omission in a round  $r$  fails to receive a message from a subset of processes in the system. More formally

$$(\forall p_j \in \Omega)(MR(i, j, r) = MS(j, i, r) \vee MR(i, j, r) = \perp) \wedge (\exists p_j \in \Omega)(MR(i, j, r) = \perp \wedge MS(j, i, r) \neq \perp).$$

*General-omission failure.* A process  $p_i$  that commits a general omission in a round  $r$  if  $p_i$  commits either a send- and/or a receive-omission failure in round  $r$ .

*Byzantine failure.* A process  $p_i$  that commits a Byzantine failure in round  $r$  may arbitrarily deviate from its protocol, there is no message authentication mechanism:  $p_i$  sends any message, alters any message that  $p_i$  has received, or relays spurious messages that appear to be from other processes. More formally, a process  $p_i$  that commits a Byzantine failure in round  $r$  performs at least one of the following items in round  $r$  and behaves correctly for the rest of round  $r$ :

- $p_i$  fails to send correctly to at least one process

$$(\exists p_j \in \Omega)(MS(i, j, r) \neq O_i(ST(i, r - 1), I(i, r), r)),$$

- $p_i$  fails to receive correctly from at least one process

$$(\exists p_j \in \Omega)(MR(i, j, r) \neq MS(j, i, r)),$$

- $p_i$  makes an incorrect state transition

$$(ST(i, r) \neq T_i(ST(i, r - 1), MR(i, r), r)),$$



## 2.5. Runs

A *run* corresponds to an execution of a protocol, and is defined as a tuple  $\langle I, ST, MS, MR \rangle$ , where  $I$  is the input pattern observed in the run,  $ST$  is the state function,  $MS$  represents the messages sent, and  $MR$  the messages received.

Any process  $p_i$  is *correct in run  $R$*  if  $p_i$  is correct up to round  $r$ , for any  $r \geq 0$ . A process that is not correct in run  $R$  is *faulty in  $R$* . Let  $correct(R, r)$ ,  $r \geq 1$ , denote the set of processes correct up to round  $r$  in run  $R$  (all the processes are correct up to round 0). The set of correct processes in run  $R$  is  $correct(R) = \cup_{r \geq 0} correct(R, r)$ , whereas the set of faulty processes in run  $R$  is  $faulty(R) = \Omega \setminus correct(R)$ .

## 2.6. System models and problem specifications

A *system model*, or *model*, is the particular set of all runs that can occur under some conditions (for any protocol). Hence a system model may be defined as a set of conditions that its runs must satisfy. We denote by  $R(\Pi, M, \Gamma_{\mathcal{I}})$  the set of all runs produced by protocol  $\Pi$  in system model  $M$  and input pattern in  $\Gamma_{\mathcal{I}}$ . A problem specification, or *problem*,  $\Sigma$  is defined as a predicate on runs.

**Definition 1.** A protocol  $\Pi$  *solves* a problem  $\Sigma$  in system model  $M$  with input pattern in  $\Gamma_{\mathcal{I}}$  if and only if  $(\forall R \in R(\Pi, M, \Gamma_{\mathcal{I}}))(R \text{ satisfies } \Sigma)$ .

A model  $M$  is defined as a particular set of runs. In particular, we define six distinct models:

- Model  $PSR(n, t)$  (*Perfectly synchronized round*) is defined by all runs over  $n$  processes where at most  $t < n$  processes are subject to atomic failures, and the remaining processes are correct.
- Model  $Crash(n, t)$  is defined by all runs over  $n$  processes where at most  $t < n$  processes are subject to crash failures only, and the remaining processes are correct.
- Model  $Omission(n, t)$  is defined by all runs over  $n$  processes where at most  $t < n$  processes are subject to crash failures or send-omission failures in some round, and the remaining processes are correct.
- Models  $General(n, t)$  and  $General - MAJ(n, t)$  are defined by all runs over  $n$  processes where at most  $t$ , where  $t < n$  processes, respectively,  $t < n/2$  processes, are subject to crash failures, send- and/or receive-omission failures in some round, and the remaining processes are correct.
- Model  $Byzantine(n, t)$  is defined by all runs over  $n$  processes where at most  $t < n/3$  processes are subject to Byzantine failures, and the remaining processes are correct.

We say that a model  $M_s$  is *stronger* than a model  $M_w$ , and we write  $M_s \geq M_w$ , if and only if  $M_s \subseteq M_w$ . We say that a model  $M_s$  is *strictly stronger* than  $M_w$ , and we write  $M_s \succ M_w$ , if and only if  $M_s \geq M_w$  and  $M_w \not\geq M_s$ . Weaker and strictly weaker relations are defined accordingly. From the equations above, it is clear that  $PSR(n, t) \geq Crash(n, t) \geq Omission(n, t) \geq General(n, t)$ .

For any run  $R$ , in any model, we denote by  $f$  the effective number of faulty processes in  $R$ , i.e.,  $f = |faulty(R)|$ .

### 3. Simulation and transformation

The notions of simulation and transformation, although intuitive, require a precise definition. In particular, some problems in a given model cannot be transformed into another model, simply because they cannot be solved in the second model.

Consider two models  $M_s$  and  $M_w$ , such that  $M_s \geq M_w$ . A transformation  $\mathcal{T}$  takes any protocol  $\Pi_s$  designed to run in the strong model  $M_s$  and converts it into a protocol  $\Pi_w = \mathcal{T}(\Pi_s)$  that runs correctly in the weak model  $M_w$ . For example,  $M_s$  could be *PSR* and  $M_w$  could be *Crash*. To avoid ambiguities, we call a round in the weak model  $M_w$ , a *phase*.

The transformation of a protocol  $\Pi_s$  in  $M_s$  to a protocol  $\Pi_w$  in  $M_w$  is defined through a *simulation function*,  $Sim$ , which simulates a run of  $\Pi_s$  by a run of  $\Pi_w$ . In [3], the authors present a problem, called the strong dependent decision (SDD) problem, which is solvable in a synchronous model, and show that this problem does not admit any solution in an asynchronous model augmented with a Perfect failure detector [2] when one process can crash. This seems to contradict the fact that algorithms designed for the former model can be run in the latter [17]. The contradiction is in apperance only, and depends on how we define the notion of simulation.

For any process  $p_i$  executing a protocol  $\Pi_w$  in  $M_w$  simulating  $\Pi_s$ , the local state  $s$  of  $p_i$  contain variables  $s.states_i$  and  $s.ss_i$ , which maintain the simulated states of protocol  $\Pi_s$ . Indeed, in contrast with the doubling technique of [17] where each state of the run in  $\Pi_w$  simulates *at most* one state of a run in  $\Pi_s$ , we do not restrict our transformation to simulate only one state of a run of  $\Pi_s$  in a state of the run of  $\Pi_w$ . More precisely,  $s.states_i$  is a set of round numbers, such that, at the end of any phase  $x$ , for any round  $r$  in  $s.states_i$ ,  $s.ss_i[r]$  gives the  $r$ -th simulated state, i.e., the simulated state at the end of round  $r$  ( $s.states_i^0 = \{0\}$ ,  $s.ss_i^0[0] = s_i$ ). We now give the formal definitions of our transformation notions, over an arbitrary set of input values  $\mathcal{I}$ .

We first define the notion of non-uniform transformation, and then use this definition to define the notion of uniform transformation.

**Definition 2.** An algorithm  $\mathcal{T}$  is called a non-uniform transformation from model  $M_s$  to model  $M_w$ , with input pattern in  $\Gamma_{\mathcal{I}}$ , if there is a corresponding simulation function  $Sim$  and a function  $f : \mathcal{R} \rightarrow \mathcal{R}$ , with the following property: for any protocol  $\Pi_s$  and any run  $R_w$  of  $\Pi_w = \mathcal{T}(\Pi_s)$  running in  $M_w$  with input pattern  $I_w$ ,  $Sim$  maps run  $R_w = \langle I_w, ST_w, MS_w, MR_w \rangle$  onto a corresponding simulated run  $R_s = Sim(R_w)$  such that

- (i)  $R_s = \langle I_s, ST_s, MS_s, MR_s \rangle$  and  $R_s \in R(\Pi_s, M_s, \Gamma_{\mathcal{I}})$ ,
- (ii)  $correct(R_w) \subseteq correct(R_s)$ ,
- (iii)  $(\forall r \in \mathcal{R})(\forall p_i \in correct(R_w))(I'(i, r) = I(i, r))$ ,
- (iv)  $(\forall x \in \mathcal{R})(\forall p_i \in correct(R_w))(\forall r \in ST_w(i, x).states)$   
 $(ST_w(i, x).ss[r] = ST_s(i, r))$ ,
- (v)  $(\forall r \in \mathcal{R})(\forall p_i \in correct(R_s))(\exists c \leq f(r))(r \in ST_w(i, c).states)$ ,
- (vi)  $(\forall r, r' \in \mathcal{R}, r \neq r')(\forall p_i \in correct(R_w))$   
 $(x \in ST_w(i, r).states \cap ST_w(i, r').states \Rightarrow$   
 $ST_w(i, r).ss[x] = ST_w(i, r').ss[x])$ ,
- (vii)  $(\forall x \in \mathcal{R})(\forall p_i \in correct(R_w))(\forall r \in ST_w(i, x).states)(\forall r' < r)$   
 $(r' \in \cup_{k=0}^x ST_w(i, k).states)$ .

**Definition 3.** An algorithm  $\mathcal{T}$  is called a uniform transformation from model  $M_s$  to model  $M_w$ , with input pattern in  $\Gamma_{\mathcal{T}}$ , if  $\mathcal{T}$  is a non-uniform transformation from  $M_s$  to  $M_w$  with simulation function  $Sim$  and function  $f$  satisfying the properties of a non-uniform transformation and such that, for any protocol  $\Pi_s$  and any run  $R_w$  of  $\Pi_w = \mathcal{T}(\Pi_s)$  running in  $M_w$  with input pattern  $I_w$ ,  $Sim$  maps run  $R_w = \langle I_w, ST_w, MS_w, MR_w \rangle$  onto a corresponding simulated run  $R_s = Sim(R_w)$ , the additional following properties are also satisfied:

- (iii')  $I_w = I_s$ ,
- (iv')  $(\forall x \in \mathcal{R})(\forall p_i \in \Omega)(\forall r \in ST_w(i, x).states)(ST_w(i, x).ss[r] = ST_s(i, r))$ ,
- (vi')  $(\forall r, r' \in \mathcal{R}, r \neq r')(\forall p_i \in \Omega)$   
 $(x \in ST_w(i, r).states \cap ST_w(i, r').states \Rightarrow$   
 $ST_w(i, r).ss[x] = ST_w(i, r').ss[x])$ ,
- (vii')  $(\forall x \in \mathcal{R})(\forall p_i \in \Omega)(\forall r \in ST_w(i, x).states)(\forall r' < r)$   
 $(r' \in \cup_{k=0}^x ST_w(i, k).states)$ .

The difference between the definitions of a non-uniform and a uniform transformation concerns properties (iii), (iv), (vi), and (vii) in the non-uniform case, denoted, respectively (iii'), (iv'), (vi'), and (vii') in the uniform case. For a non-uniform transformation, the corresponding properties must be satisfied by the processes that are correct in the underlying run  $R_w$  of  $M_w$ , whereas for a uniform transformation, the corresponding properties must be satisfied by the processes that are correct in the simulated run  $R_s$  of  $M_s$ . These processes include those that are correct in  $R_w$  by Property (ii).

Property (i) states that the simulated run should be one of the runs of the simulated protocol. Property (ii) forces a correct process to be correct in the simulated run (though a faulty process may appear correct in the simulated run). Properties (iii) and (iii') state that the input pattern is preserved by the simulation. Properties (iv) and (iv') state that any simulated state is correct w.r.t.  $\Pi_s$ . Property (v) forces the simulation to accomplish progress. Properties (vi) and (vi') state that each state of  $\Pi_s$  is simulated in at most one manner. Properties (vii) and (vii') force a process to simulate states sequentially w.r.t.  $\Pi_s$ .

With a non-uniform transformation, any process  $p_i$  that is faulty in any run  $R_w$  of  $M_w$  is not required, according to Definition 2, to make any progress nor to guarantee any property on the states that  $p_i$  simulates. In particular, if the underlying model  $M_w$  permits it, e.g.,  $M_w$  is *Byzantine*,  $p_i$  may behave arbitrarily and may simulate states that are not consistent in  $M_s$  with the simulation of the correct processes in  $R_w$ . Nevertheless Property (i) ensures the correctness of the simulation w.r.t.  $M_s$  at all times. Roughly speaking, all the processes correct in  $R_w$  maintain a simulated state for any process faulty in  $R_w$ , that is consistent in  $M_s$ . In Section 5 we will see that we indeed define the simulation function  $Sim$  through the processes that are correct in  $R_w$  for a non-uniform transformation.

Apart from Property (iii), our definition encompasses the notion of simulation of [17], although the notion of input pattern does not appear in [17]. In the transformation of [17] from *Crash* to *Omission*, each round is transformed in two phases, which can be defined with  $f(r) = 2r$ , and  $c = 2r$  in (v). This implies that  $ST_w(i, 2x).states = \{x\}$  and  $ST_w(i, 2x + 1).states = \emptyset$ .

In the following definition, we recall the notion of *effectively solving* [17] a problem, to indicate that the resolution is obtained through a simulation function.

**Definition 4.** For any given function  $Sim$ ,  $\Pi_w$  effectively solves problem  $\Sigma$  in model  $M_w$  with input pattern in  $\Gamma_{\mathcal{T}}$  if and only if, for any run  $R \in R(\Pi_w, M_w, \Gamma_{\mathcal{T}})$ ,  $Sim(R)$  satisfies  $\Sigma$ .

The next proposition follows from definitions 2, 3, and 4.

**Proposition 5.** Let  $\Pi_s$  be any protocol that solves specification  $\Sigma$  in model  $M_s$ . If  $\mathcal{T}$  is a transformation from  $M_s$  to  $M_w$  and  $Sim$  the corresponding simulation function, then protocol  $\mathcal{T}(\Pi_s)$  effectively solves  $\Sigma$  in model  $M_w$ .

#### 4. Interactive consistency algorithms

We consider in this section the interactive consistency (IC) problem [18] that is solved to simulate a single round in our transformation. Roughly speaking, in the IC problem, each process  $p_i$  is supposed to propose an initial value and eventually decide on a vector of values.

We use two specifications of interactive consistency: a uniform specification and a non-uniform specification. Non-uniform interactive consistency is the original problem as defined in [18].

In uniform IC, each process  $p_i$  is supposed to propose a value  $v_i$  and eventually decide on a vector of values  $V_i$  such that

*Termination:* every correct process eventually decides,

*Uniform agreement:* no two decided vectors differ,

*Validity:* for any decision vector  $V$ , the  $j$ th component of  $V$  is either the value proposed by  $p_j$  or  $\perp$ , and is  $\perp$  only if  $p_j$  fails.

In non-uniform IC, each correct process  $p_i$  is supposed propose a value  $v_i$  and eventually decide on a vector of values  $V_i$  such that

*Termination:* (similar as for uniform IC) every correct process eventually decides,

*Agreement:* no two vectors decided by *correct* processes differ,

*Validity:* for any vector  $V$  decided by a *correct* process, the  $j$ th component of  $V$  is either the value proposed by  $p_j$  or  $\perp$ , and is  $\perp$  only if  $p_j$  fails.

To be self-contained, this paper presents several IC algorithms:

- for non-uniform IC, the paper presents, in Fig. 1, an early stopping interactive consistency algorithm in *General*, *Omission*, and *Crash*, for all three models with  $t < n$ . In this algorithm, and for any run  $R$  of *General*, *Omission* or *Crash*, all the correct processes decide by the end of round  $f + 1$  in run  $R$  in which at most  $f \leq t$  processes fail (this is a tight lower bound due to [14]), and halt by the end of round  $\min(f + 2, t + 1)$  in run  $R$  (this is a tight lower bound due to [5]). In any failure-free run ( $f = 0$ ), all the processes are correct and decide by the end of round 1 and halt by the end of round 2.
- for uniform IC, the paper presents two algorithms, respectively, for *Omission* and *Crash* ( $t < n$ ) in Fig. 2 and for *General-MAJ* ( $t < n/2$ ) in Fig. 3. In both algorithms, and for any run  $R$  of *Omission* or *Crash*, respectively *General-MAJ*, all the correct processes decide and halt at the end of round  $t + 1$  in run  $R$  (this is a lower bound due to [7]).

At process  $p_i$ :

```

1:  $quiet_i := \emptyset$ 
2:  $V_i := [\top, \dots, v_i, \dots, \top]$ 

3: for  $r$  from 1 to  $t + 1$  do                                     {For each round  $r$ }
4:   send  $V_i$  to all processes                                   {Send phase}
5:   if  $(\forall k : V_i[k] \neq \top)$  then halt

6:    $\forall p_j \in \Omega \setminus quiet_i$  : if receive  $V_j$  then           {Receive phase}
7:      $\forall k : \text{if } (V_i[k] = \top) \text{ then } V_i[k] := V_j[k]$ 
8:   else
9:      $quiet_i := quiet_i \cup \{p_j\}$ 

10:  if  $|quiet_i| < r$  then
11:     $\forall k : \text{if } V_i[k] = \top \text{ then } V_i[k] := \perp$ 
12:  if  $(\forall k : V_i[k] \neq \top)$  then decide( $V_i$ )

13:  $\forall k : \text{if } V_i[k] = \top \text{ then } V_i[k] := \perp$                  {If not decided in round  $t + 1$ }
14: decide( $V_i$ )
15: halt

```

Fig. 1. Non-uniform interactive consistency in *General* ( $t < n$ ).

In *Byzantine*, we refer the reader to the following algorithms from the literature solving non-uniform IC: [18] (with exponentially growing message size), or [9] (with polynomially growing message size). It is not possible to solve uniform IC in *General* ( $t < n$ ) [17,19] and in *Byzantine* [18].

#### 4.1. Non-uniform interactive consistency

We give in Fig. 1 an algorithm solving non-uniform IC in *General*, *Omission* and *Crash*, for  $t < n$ .

The algorithm is derived from the Terminating Reliable Broadcast algorithm of [20]. We briefly describe how the algorithm works. In this description, we focus on the value of process  $p_i$ , the same mechanism extends to the value of the other processes.

In round 1,  $p_i$  sends its value to all the processes. In later rounds, every process relays  $p_i$ 's value by sending its complete vector of values to all the processes. In parallel, each process  $p_j$  maintains a set *quiet* with the processes from which  $p_j$  does not receive a message in some round. When receiving a vector of values,  $p_i$  copies from this vector the values that  $p_i$  does not have in its own vector of values.

At process  $p_i$ :

```

1:  $V_i := [\perp, \dots, v_i, \dots, \perp]$ 
2: for  $r$  from 1 to  $t + 1$  do
3:   send  $V_i$  to all processes
4:    $\forall p_j \in \Omega$  : if receive  $V_j$  do
5:      $V_i[(j - r + 1) \bmod n] := V_j[(j - r + 1) \bmod n]$ 
6:   decide( $V_i$ )

```

Fig. 2. Interactive consistency in *Omission* ( $t < n$ ).

```

At process  $p_i$ :
1:  $halt_i := \emptyset$ 
2:  $suspect_i := \emptyset$ 
3:  $V_i := [\perp, \dots, v_i, \dots, \perp]$ 
4: for  $r$  from 1 to  $t + 1$  do
5:   send  $[V_i, halt_i]$  to all processes
6:    $\forall p_j \in \Omega \setminus halt_i$ : if receive  $[V_j, halt_j]$  then
7:     if  $p_i \in halt_j$  then
8:        $suspect_i := suspect_i \cup \{p_j\}$ 
9:     else
10:       $halt_i := halt_i \cup \{p_j\}$ 
11:   for all  $p_j \in \Omega \setminus halt_i$  and for all  $k$  do
12:     if  $V_i[k] = \perp$  then  $V_i[k] := V_j[k]$ 
13: if  $|halt_i \cup suspect_i| \leq t$  then
14:   decide( $V_i$ )

```

Fig. 3. Interactive consistency in *General* ( $t < n/2$ ).

At the end of any round  $r$ , if  $p_i$  did not receive a message from strictly less than  $r$  processes, then  $p_i$  fills the missing values in its vector of value with  $\perp$ , meaning that the process at this position is faulty. Process  $p_i$  decides on its vector of values at the end of any round when its vector is filled for each position, either with a value or with  $\perp$ . After sending its vector in the next round,  $p_i$  may halt. If after  $t + 1$  rounds,  $p_i$  still has not decided,  $p_i$  fills the missing values in its vector of values with  $\perp$  and decides on its vector before halting.

#### 4.2. Uniform interactive consistency in *Omission* and *Crash*

The algorithm in Fig. 2 solves IC in *Omission* and *Crash*. This algorithm is given in [11]. In both models we assume  $t < n$ . In the algorithm, all the processes that decide, decide after  $t + 1$  rounds. We briefly explain how the algorithm works.

In the algorithm in Fig. 2, any process  $p_i$  sends its value  $v_i$  to all the processes in round 1,  $v_i$  being embedded in the vector  $V_i$ . The value  $v_i$  is then relayed to all processes by another process in each round, until round  $t + 1$ . When any process relays a estimate value, it sends its vector of estimate values to all the processes. For any process  $p_i$ , the value of  $p_i$  is relayed successively by the processes  $p_i, \dots, p_{(i+t) \bmod n}$ , respectively, in round 1 to  $t + 1$ .

In the algorithm, the relaying mechanism is hidden in the reception phase. More precisely, when any process  $p_k$  receives a vector of estimate values in round  $r$  from  $p_j$ ,  $p_k$  only copies into its own vector of estimate values, the component that  $p_k$  is supposed to relay in that particular round  $r$ . Process  $p_k$  does not make use of any other estimate value from  $p_j$ 's vector.

In any particular round  $r$ , any process  $p_k$  relays the value of a different process (i.e., the value of process  $(k - r + 1) \bmod n$ ); whereas the value of any process  $p_i$  is relayed by a different process (i.e., process  $(i + r - 1) \bmod n$ ).

We now give an intuition of the correctness of the algorithm. The intuition is particularly simple: as  $t + 1$  processes are involved in relaying the value of any process  $p_i$ , at least one of them is correct, say  $p_j$ , and thus never commits omissions. When  $p_j$  relays the estimate value of  $p_i$ 's value to all the

processes, as  $p_j$  is correct, all processes receive  $p_j$ 's estimate value. From this round on, all processes maintain the same estimate value for  $p_i$ 's value in their vector of estimate values.

#### 4.3. Uniform interactive consistency in General-MAJ

The algorithm in Fig. 3 solves IC in *General-MAJ*. The algorithm is inspired from the uniform consensus algorithm of [6]. In *General-MAJ*, we assume  $t < n/2$ , as [17,19] that this is necessary for the problem to be solvable. In the algorithm, all the processes that decide, decide after  $t + 1$  rounds. We briefly explain how the algorithm works.

Primarily, the processes exchange vectors of estimate values, and update their own vector with the vectors received in each round. In the absence of omission, this procedure ensures that each process decides on the same vector of estimate values at the end of round  $t + 1$ . To tolerate general omission, we do not allow some faulty processes (those with insufficient information) to decide, at the end of round  $t + 1$ .

Similarly as in the algorithm for *Omission*, any process  $p_i$  maintains a set  $halt_i$  with the identity of processes from which  $p_i$  does not receive a message in this round or in a previous round. Moreover, any process  $p_i$  maintains in addition a set  $suspect_i$  with the identity of any process  $p_x$  which includes  $p_i$ 's identity in its set  $halt_x$ .  $p_i$  maintains the vector of estimate values  $V_i$ , and decides on this vector at the end of round  $t + 1$ .

### 5. Shifting transformation

We present our algorithm to transform any protocol  $\Pi$  written in *PSR* into a protocol  $\Pi'$  in a weaker model  $M_w$  such that  $\Pi'$  simulates  $\Pi$ , through a simulation function *Sim* that we give.

For any two distinct processes  $p_i$  and  $p_j$  simulating protocol  $\Pi$ , we do not necessarily assume that  $\Pi_i = \Pi_j$ . However, we will assume for the time being that  $p_i$  knows the state machine  $\Pi_j = \langle s_j, T_j, O_j \rangle$  executed by  $p_j$ . We relax this assumption in Section 6.

Our transformation works on a round basis: it transforms a single round in *PSR* into several phases in  $M_w$ . The key to its efficiency is that a phase is involved in the simulation of more than one round simultaneously. We start by giving a general definition of the notion of shifting transformation, before giving our own.

Our algorithm implements at the same time a non-uniform and a uniform transformation. Roughly speaking, the transformation algorithm relies on an underlying interactive consistency algorithm, which alone determines if the overall transformation algorithm is uniform or non-uniform. More precisely, using a non-uniform, respectively uniform, IC algorithm as the underlying IC algorithm in the transformation leads to a transformation that is also non-uniform, respectively uniform. In the rest of this section, we thus present a single transformation algorithm that uses an underlying IC algorithm that may either be uniform or non-uniform, depending on the transformation required.

From this observation, it is thus clear that the uniform transformation may work only in *Crash*, *Omission* and *General-MAJ* ( $t < n/2$ ), and not in *General* ( $t < n$ ) or *Byzantine*, since it is not possible to implement uniform IC in both models [17,19]. The non-uniform transformation works for all

models, i.e., transformation may work only in *Crash*, *Omission*, and *General-MAJ* ( $t < n/2$ ), *General* ( $t < n$ ) and *Byzantine*.

Let  $\Pi_s$  be any protocol in model  $M_s$ ,  $\mathcal{T}$  any transformation (uniform or not) from  $M_s$  to  $M_w$ , and  $\Pi_w = \mathcal{T}(\Pi_s)$  the transformed protocol. Roughly speaking, a shifting transformation is such that any process simulates *round*  $r$  of  $\Pi_s$  after a bounded number of phases counting from *phase*  $r$ . More precisely:

**Definition 6.** A non-uniform (respectively uniform) transformation  $\mathcal{T}$  from model  $M_s$  to model  $M_w$  is a non-uniform (respectively uniform) *shifting transformation* if and only if there exists a constant  $S \in \mathbb{N}$ , such that, for all  $r \in \mathcal{R}$ ,  $f(r) = r + S$ . We call  $S$  the *shift* of the transformation.

### 5.1. Algorithm

In our transformation, all processes collaborate to reconstruct the failure and input patterns of a run in *PSR*. They accomplish both reconstructions in parallel, one round after another. When processes terminate the reconstruction of the patterns for a round, they locally execute one step of the simulated protocol. If a process realizes that it is faulty in the simulated failure pattern, this process simulates a crash in *PSR*. To simulate one round in *PSR*, processes solve exactly one instance

```

1: failure := ∅                                {failure corresponds to one round of the failure pattern}
2: simulatedRound := 1                         {simulatedRound is the current simulated round number}
3: (∀j ∈ [1, n])(simst(0)[j] := s_j)           {state of protocol Π_s for p_j at the end of round 0}
4: states := {0}                               {set of rounds of protocol Π_s simulated by Π_w in the current round}
5: ss[0] := s_i                                {set of states of protocol Π_s simulated by Π_w in the current round}
6: for phase r (r = 1, 2, ...) do
7:   input := receiveInput()                    {receive input value corresponding to I(i, r)}
8:   start IC instance number r, and propose(input)
9:   execute one round of all pending IC instances
10:  states := ∅                                {has any IC instance decided?}
11:  while simulatedRound-th instance of IC has decided do
12:    states := states ∪ {simulatedRound}       {instance simulatedRound has decided}
13:    decision := decision vector of instance simulatedRound {reconstruct patterns}
14:    for each p_j ∈ Ω do                       {check input against byzantine failures}
15:      if decision[j] ≠ ⊥ and decision[j] ∉ I then
16:        decision[j] := ⊥
17:    failure := failure ∪ {p_j | decision[j] = ⊥} {ensure atomic failures only}
18:    if p_i ∈ failure then halt                 {is process p_i faulty?}
19:    for each p_j ∈ Ω do                       {adjust decision vector with previous failure pattern}
20:      if p_j ∈ failure then
21:        decision[j] := ⊥
22:    for each p_j ∈ Ω do                       {generate messages}
23:      if p_j ∉ failure then
24:        rcvd[j] := O_j(simst(simulatedRound - 1)[j], decision[j], simulatedRound)
25:      else
26:        rcvd[j] := ⊥
27:    for each p_j ∈ Ω do                       {perform state transitions}
28:      if p_j ∉ failure then
29:        simst(simulatedRound)[j] := T_j(simst(simulatedRound - 1)[j], rcvd, simulatedRound)
30:      else
31:        simst(simulatedRound)[j] := simst(simulatedRound - 1)[j]
32:    ss[simulatedRound] := simst(simulatedRound)[i]
33:    simulatedRound := simulatedRound + 1       {increment simulated round counter}
34:    if r - simulatedRound ≥ δ then halt        {is process p_i faulty?}

```

Fig. 4. Transformation algorithm (code for process  $p_i$ ).



of the interactive consistency problem. In the instance of IC, each process  $p_i$  proposes its own input value for the round. The decision vector corresponds at the same time to a round of the failure pattern, and of the input pattern.

Fig. 4 gives the transformed protocol  $\mathcal{T}(\Pi_s)$  for process  $p_i$ , in terms of  $\Pi_s$  and the input pattern  $I$ . The underlying IC algorithm may either be uniform or non-uniform, for the transformation to be uniform or non-uniform. When using the early stopping non-uniform IC algorithm in Fig. 1 for a transformation into *Crash*, *Omission*, or *General*, with  $t < n$  for all three models, each round of *PSR* is transformed into  $f + 1$  phases of the weaker model, in any run  $R$  with at most  $f \leq t$  failures. In any failure-free ( $f = 0$ ) run, we observe that the transformation of one round of *PSR* requires a single phase of the weaker mode. In this sense, the simulation outputs the results in real time.

When using a uniform IC algorithm, e.g., the algorithm in Fig. 2 when the weaker model is *Crash* or *Omission* with  $t < n$  or the algorithm in Fig. 3 when the weaker model is *General-MAJ* with  $t < n/2$ , each round of *PSR* is transformed into  $t + 1$  phases of the weaker model, in any run  $R$ .

For the sake of simplicity, the transformation algorithm is given in an operational manner (i.e., pseudo-code). During any phase, many IC instances might be running together. If the condition of the *while loop* at line 11 (“*simulatedRound-th instance of IC has decided*”) is true in a phase  $x$  of process  $p_i$ , then we denote by  $decision_i(simulatedRound)$  the decision vector for the instance of IC in line 13,  $failure_i(simulatedRound)$  the value of the variable *failure* updated in line 17,  $rcvd_i(simulatedRound)$  the value of the variable *rcvd* updated in lines 24 or 26, and  $simst_i(simulatedRound)$  the value of the variable *simst* updated in lines 29 or 31. The following proposition defines the simulation function *Sim* in our transformation.

The next proposition gives the construction of the simulation function *Sim* for both the uniform and non-uniform transformation.

**Proposition 7.** *The simulation  $Sim$  for a run of  $\mathcal{T}(\Pi_s)$ ,  $R = \langle I, ST, MS, MR \rangle$ , is defined by  $R' = \langle I', ST', MS', MR' \rangle$  as follows. Let  $p_i$  be a process in  $correct(R)$ . We consider the simulation of round  $r$  of  $R'$ , for any process  $p_j$ .*

- (i)  $I'(j, r)$  is the value  $decision_i(r)[j]$  of the  $r$ -th instance of IC.
- (ii) if  $r = 0$  then  $ST'(j, 0) = s_j$ , otherwise  $ST'(j, r) = simst_i(r)[j]$ .
- (iii) if  $p_j \in failure_i(r)$  then  $MS'(j, k, r) = \perp$ , otherwise  
 $MS'(j, k, r) = rcvd_i(r)[j]$  for any process  $p_k \in \Omega$ .
- (iv) if  $p_j \in failure_i(r)$  then  $MR'(j, r) = [\perp, \dots, \perp]$ , otherwise  
 $MR'(j, r) = rcvd_i(r)$ .

The next propositions assert the correctness of the non-uniform, respectively, uniform, transformation.

**Proposition 8.** *The algorithm of Fig. 4 (used in conjunction with an underlying non-uniform IC algorithm) is a non-uniform shifting transformation from  $PSR(n, t)$  to  $Crash(n, t)$ ,  $Omission(n, t)$ ,  $General(n, t)$  ( $t < n$ ), or  $Byzantine(n, t)$  ( $t < n/3$ ) where the shift  $S$  is the number of rounds needed to solve non-uniform interactive consistency in  $Crash(n, t)$ ,  $Omission(n, t)$ ,  $General(n, t)$  ( $t < n$ ), or  $Byzantine(n, t)$  ( $t < n/3$ ).*

**Proposition 9.** *The algorithm of Fig. 4 (used in conjunction with an underlying uniform IC algorithm) is a uniform shifting transformation from  $PSR(n, t)$  to  $Crash(n, t)$ ,  $Omission(n, t)$  ( $t < n$ ), or*

*General – MAJ*( $n, t$ ) ( $t < n/2$ ) where the shift  $S$  is the number of rounds needed to solve uniform interactive consistency in *Crash*( $n, t$ ), *Omission*( $n, t$ ) ( $t < n$ ), *General – MAJ*( $n, t$ ) ( $t < n/2$ ).

In this section we prove Proposition 9, but Proposition 8 would be proved in the exact same manner.

To prove Proposition 9, we first show that the construction of function *Sim* in Fig. 4 is consistent with Proposition 7. We proceed through a series of lemmas.

**Lemma 10.** *For any run  $R$  and any process  $p_i$  in  $\text{correct}(R)$ ,  $p_i$  never halts, and decides in all IC instances.*

**Proof.** A process  $p_i$  that is correct in  $R$  exists since  $t < n$  for *Crash*( $n, t$ ), *Omission*( $n, t$ ) or *General-MAJ*( $n, t$ ). Thus,  $p_i$  does not halt in lines 18 or 34, as  $p_i \in \text{correct}(R)$ . Process  $p_i$  is correct and never crashes. By the termination property of IC,  $p_i$  always decides in any instance of IC. Thus  $p_i$  never halts in line 34. Moreover, by the validity property of IC, in any decision vector,  $\text{decision}[i] \neq \perp$ . Thus  $p_i$  never halts in line 18.  $\square$

**Lemma 11.** *For any run  $R$ , any process  $p_i$  in  $\text{correct}(R)$ , any process  $p_j$ , and any round  $r$  such that  $p_j$  decides in the  $r$ -th instance of IC, we have the following properties:*

- $\text{decision}_i(r) = \text{decision}_j(r)$
- $\text{failure}_i(r) = \text{failure}_j(r)$
- $\text{rcvd}_i(r) = \text{rcvd}_j(r)$
- $\text{simst}_i(r) = \text{simst}_j(r)$

**Proof.** By Lemma 10,  $p_i$  decides in all IC instances. By the agreement property of IC, the decision is the same for  $p_i$  and  $p_j$ . In the algorithm,  $p_j$  decides in the  $r$ -th instance of IC if and only if  $p_j$  has decided in all previous instances. We show the three last items by induction on  $r$ . Initially, because of initialization, the properties are true for  $r = 0$ . Assume the properties are true up to round  $r - 1$ . When  $p_i$  decides in the  $r$ -th instance, it adds a set of processes to  $\text{failure}_i$ . By the agreement property of IC,  $p_i$  and  $p_j$  add the same set. By induction,  $\text{failure}_i(r) = \text{failure}_j(r)$ . As  $\text{decision}(r)$  and  $\text{failure}(r)$  are the same for all processes for which they are defined, by induction hypothesis, we have  $\text{rcvd}_i(r) = \text{rcvd}_j(r)$  and  $\text{simst}_i(r) = \text{simst}_j(r)$ .  $\square$

We can define the simulation through the value of the variables of any correct process. Consider any run  $R$  and let  $p_k$  be a correct process in  $R$  (we know there exists at least one as  $t < n$ ). We define the simulation through the variables of  $p_k$ .

**Lemma 12.** *If there exists a round  $r$  and a process  $p_i$  such that  $i \in \text{failure}_k(r)$  then, for any  $r' \geq r$ ,  $\text{simst}_k(r')[i] = \text{simst}_k(r)[i]$ .*

**Proof.** Directly from the transformation algorithm.  $\square$

**Lemma 13.** *For any process  $p_i$ , any round  $r$  such that  $\text{decision}_i(r)$  and  $\text{decision}_i(r + 1)$  occur,  $\text{failure}_i(r) \subseteq \text{failure}_i(r + 1)$ .*

**Proof.** From the transformation algorithm,  $\text{failure}_i$  always increases.  $\square$

**Lemma 14.**  $\text{correct}(R) \subseteq \text{correct}(R')$ .

**Proof.** By Lemma 10, all correct processes decide in all instances of IC. By termination of IC, no correct processes ever halt in line 34. By the validity property of IC, the decision value is not  $\perp$  for any correct process in any decision vector. Thus no correct process ever halts in line 18. Therefore all correct processes in  $R$  are correct in  $R'$ .  $\square$

**Lemma 15.** For any process  $p_i$  and any protocol  $\Pi_s$  to simulate, let  $\langle s_i, T_i, O_i \rangle$  be the state machine for  $p_i$ . For any round  $r$  and any  $p_j \in \Omega$ , we have:

- (1)  $ST'(i, 0) = s_i$ .
- (2) if  $p_i \in \text{correct}(R', r)$  then  $MS'(i, j, r) = O_i(ST'(i, r - 1), I'(i, r), r)$ , otherwise  $MS'(i, j, r) = \perp$ .
- (3) if  $p_i \in \text{correct}(R', r)$  then  $MR'(i, j, r) = MS'(j, i, r)$ , otherwise  $MR'(i, j, r) = \perp$ .
- (4) if  $p_i \in \text{correct}(R', r)$  then  $ST'(i, r) = T_i(ST'(i, r - 1), [MS'(1, i, r), \dots, MS'(n, i, r)], r)$ , otherwise  $ST'(i, r) = ST'(i, r - 1)$ .

**Proof.** (1) is immediate, from the initialization of the variable  $\text{simst}_k[i](0)$ . We prove (2), (3) and (4) by induction. For the case  $r = 1$ , and  $p_i \notin \text{correct}(R', 1)$ , then  $p_i \in \text{failure}_k(1)$ . By the algorithm  $\text{rcvd}_k(1)[i] = \perp$ , and by properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, 1) = \text{rcvd}_k(1)[i]$  and  $MR'(i, j, 1) = \text{rcvd}_k(1)[j]$ . By Lemma 12 and (1), we have  $ST'(i, 1) = ST'(i, 0)$ . If  $p_i \in \text{correct}(R', 1)$ , then  $p_i \notin \text{failure}_k(1)$ . By the properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, 1) = \text{rcvd}_k(1)[i]$  and  $MR'(i, j, 1) = \text{rcvd}_k(1)[j]$ . By line 24 of the algorithm,  $\text{rcvd}_k(1)[i] = O_i(\text{simst}_k(0)[i], \text{decision}_k[i], 1)$ , and so  $MS'(i, j, 1) = O_i(ST'(i, 0), I'(i, 1), 1)$ . By line 24 of the algorithm,  $\text{rcvd}_k(1)[j] = MS'(j, k, 1)$ , and so  $MR'(i, j, 1) = MS'(j, i, 1)$ . By line 29 of the algorithm,  $\text{simst}_k(1)[i] = T_i(\text{simst}_k(0)[i], \text{rcvd}_k, 1)$ , and  $ST'(i, 1) = T_i(ST'(i, 0), [MS'(1, i, 1), \dots, MS'(n, i, 1)], 1)$ .

Assume the properties (2) and (4) are true up to round  $r - 1$ . If  $p_i \notin \text{correct}(R', r)$ , then  $p_i \in \text{failure}_k(r)$ . By the properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, r) = \perp$ , and  $MR'(i, j, r) = [\perp, \dots, \perp]$ . By Lemma 12 and (1), we have  $ST'(i, r) = ST'(i, r - 1)$ . If  $p_i \in \text{correct}(R', r)$ , then by definition and by the transformation algorithm,  $p_i \notin \text{failure}_k(r)$ . By the properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, r) = \text{rcvd}_k(r)[i]$  and  $MR'(i, j, r) = \text{rcvd}_k(r)[j]$ . By line 24 of the algorithm,  $\text{rcvd}_k(r)[i] = O_i(\text{simst}_k(r - 1)[i], \text{decision}_k[i], r)$ , and so  $MS'(i, j, r) = O_i(ST'(i, r - 1), I'(i, r), r)$ . By line 29 of the algorithm,  $\text{simst}_k(r)[i] = T_i(\text{simst}_k(r - 1)[i], \text{rcvd}_k, r)$ , and so  $ST'(i, r) = T_i(ST'(i, r - 1), [MS'(1, i, r), \dots, MS'(n, i, r)], r)$ .  $\square$

By Lemmas 11 and 15, the function  $\text{Sim}$  is well defined, and consistent with Proposition 7.

**Lemma 16.**  $R'$  is a run of  $\Pi_s$ .

**Proof.** Lemmas 11 and 13 show that  $R'$  is in  $PSR$  if  $R$  is in  $M$ , where  $M \in \{\text{Crash}, \text{Omission}, \text{General-MAJ}\}$ . Lemma 15 shows that the functions  $ST'$ ,  $MS'$ , and  $MR'$  consistently define a run of  $\Pi_s$ , with input value  $I'$ .  $\square$

**Lemma 17.** Let  $x$  be any phase,  $p_i$  any process and  $r$  any round. If  $r \in ST(i, x).states$  then  $ST(i, x).ss[r] = ST'(i, r)$ .

**Proof.** Each time a round  $r$  is added to  $states_i$  (line 12),  $ss_i[r] = \text{simst}_i(r)[i]$  (line 32). By Lemma 11 and Proposition 7,  $\text{simst}_i(r)[i] = ST'(i, r)$ .  $\square$

**Lemma 18.** *For any round  $r$  and any process  $p_i$  in  $\text{correct}(R')$ , if it exists  $c$  such that  $r \in ST(i, c).states$  then  $ST(i, c).ss[r] = ST'(i, r)$ .*

**Proof.** If such a  $c$  exists, then  $p_i$  has decided in the  $r$ -th instance of IC, and by the algorithm in Figure 4, in each previous instance. By Lemma 17, we have  $ST(i, c).ss[r] = ST'(i, r)$ .  $\square$

**Lemma 19.** *Let  $p_i$  be any process in  $\text{correct}(R')$  and  $r$  any round. There exists a unique  $c \leq r + S$  such that  $r \in ST(i, c).states$  and  $ST(i, c).ss[r] = ST'(i, r)$ .*

**Proof.** By Lemma 10,  $p_i$  decides in all instance of IC. In particular, for the  $r$ -th instance,  $p_i$  decides in phase  $c$ . We pose  $S$  as the number of phases for any instance of IC to decide, and thus  $c \leq S + r$ . We have  $r \in ST(i, c).states$ , and by Lemma 18,  $ST(i, c).ss[r] = ST'(i, r)$ .  $\square$

**End of proof of Proposition 9.** We show that our function *Sim*, as defined by Proposition 7, satisfies the seven properties of Definition 3, with  $f(r) = r + S$ , where  $S$  is the number of phases to solve IC in  $M_w$ . Hence,  $S = t + 1$  for *Crash*, *Omission*, and *General-MAJ*.

Lemma 16 shows that  $\text{Sim}(R)$  is a run of  $\Pi_s$ , which implies Property (i). Lemma 14 shows that  $\text{correct}(R) \subseteq \text{correct}(R')$ , which implies Property (ii). By the definition of the simulation, we have  $I = I'$ , which implies Property (iii). Property (iv) follows from Lemma 17. By Lemma 10, any correct process  $p_i$  in  $R$  decides in any instance of IC. In particular, in the  $r$ -th instance of IC, it decides at most at phase  $r + S$ . Lemma 19 shows that for any round  $r$  and any process  $p_i$  in  $\text{correct}(R')$ , there exists  $c \leq r + S$  such that  $r \in ST(i, c).states$  and  $ST(i, c).ss[r] = ST'(i, r)$  (Property (v)) and that this  $c$  is unique (Property (vi)). If a process adds any round  $r$  in its set *states*, by the algorithm, it has decided the  $r$ -th instance of IC, and all previous instances of IC as well. This implies Property (vii).

## 6. Transformation extension

We give in this section an extension of the uniform transformation of Fig. 4. In the transformation of Fig. 4, the processes only need to send their input value in a phase, because the protocol itself can be locally simulated by other processes. We assume now that the processes do not know the state machine simulated by any other process. As a result, any process  $p_i$  needs to send, in addition to the message of the previous transformation, the content of the message it would normally send in the simulated protocol, i.e., the output of function  $O_i$ . Nevertheless, as with our previous transformation, we would like to start the simulation of a round before the decision of all previous simulations are known. Thus  $p_i$  cannot know in which *precise* state of the protocol it should be at the time it has to generate a message (remember that the current state is a parameter of the message output function  $O_i$ ).

More precisely, consider any process  $p_i$  simulating a run  $R' = \langle I', ST', MS', MR' \rangle$  of *PSR*. The idea of the extended transformation is to maintain, for  $p_i$ , all simulated states of  $ST'$  that are coherent with previous (terminated) simulations, and which only depend on the outcome of on-going (not yet terminated) simulations. Hereafter, these states are called the *extended set of states* and denoted by *es*. For any two processes  $p_i$  and  $p_j$  simulating the execution of protocol  $\Pi$  in *PSR*, we denote by  $m_j$  the message  $p_j$  sends to  $p_i$  in round  $r$ . Before the end of round  $r$  simulation, i.e., in any phase  $r'$  such that  $r \leq r' \leq r + \delta - 2$  where  $\delta$  is the number of phases for the  $r$ -th IC instance to decide,  $p_i$

does not know the decision value corresponding to  $p_j$ 's proposal: (1) as long as  $p_i$  has not received  $m_j$ , the decided value can be any value in  $\mathcal{M}'$  (including  $\perp$ ), and (2) if  $p_i$  receives  $m_j$ , the decided value can either be  $m_j$  or  $\perp$ . To be able to start the next instance in the next phase,  $p_i$  generates a new extended set of states. To generate this set of states,  $p_i$  computes  $T_i$  on every state in the current set of states, with every possible combination of messages received in phase  $r$  (i.e.,  $\perp$  values are successively substituted by any value of  $\mathcal{M}$ , and any received value successively substituted with  $\perp$ ). To each state in the extended set of states corresponds a message of  $\Pi_i$  to be sent in round  $r$  by  $p_i$ . These messages are gathered in a set, hereafter called the *extended message* and denoted by  $em$ .

For example, consider the case of the *Crash*(3,2) model with  $\mathcal{I} = \mathcal{M} = \{0, 1\}$ . After phase 1, process  $p_1$  gathers the received values in the vector  $[10 \perp]$ . The possible combinations of messages are  $[10 \perp]$ ,  $[100]$ ,  $[101]$ ,  $[1 \perp \perp]$ ,  $[1 \perp 0]$ ,  $[1 \perp 1]$ ,  $[\perp 0 \perp]$ ,  $[\perp 00]$ ,  $[\perp 01]$ ,  $[\perp \perp 0]$ , and  $[\perp \perp 1]$ . Process  $p_1$  generates the extended set of states by applying function  $T_1$  on each combination of messages.

Fig. 5 presents our extended transformation algorithm. For the sake of clarity, we ignore possible optimizations in this algorithm (e.g., any process can reduce the number of possible states as it receives more values from other processes). In Fig. 5, we denote by  $rcvd[r]$  the messages of instance  $r$  received in phase  $r$ . We assume without loss of generality that any process sends in any phase of the underlying IC algorithm, the value it proposes to this instance.

Let  $p_i$  be any process simulating state machine  $\Pi_i = \langle s_i, T_i, O_i \rangle$ . We consider the transformation algorithm at the beginning of phase  $r$ .

### 6.1. Message generation

At the beginning of phase  $r$ ,  $p_i$  receives an input value  $input = I(i, r)$ , and computes a new extended set of states  $es'$  and the corresponding extended message  $em$ . A tuple in  $em$  is of the form  $\langle num(st), rec, num(st'), m \rangle$ , and contains (i) the identifier  $num(st)$  of a possible state  $st$  of  $p_i$  at the beginning of round  $r - 1$ , (ii) a combination  $rec$  of messages received by  $p_i$  in phase  $r - 1$ , (iii) the identifier  $num(st')$  of the state  $st'$  of  $p_i$  at the beginning of round  $r$ , such that  $st' = T_i(st, rec, r - 1)$ , (iv) the message sent in round  $r$ , i.e.,  $m = O_i(st', I(i, r), r)$ . For each state  $st$  in the current extended set of states  $es$ , and for any combination  $rec$  of messages (according to the extended messages of phase  $r - 1$ ),  $p_i$  computes the next state  $st'' = T_i(st, rec, r)$  (whenever  $p_i$  includes a new state  $st'$  in  $es'$ , it associates a unique identifier  $num(st')$  with  $st'$ ), and the corresponding message  $m = O_i(st', input, r)$ .  $p_i$  sends  $em$  and the extended messages of other running IC instances in phase  $r$ .

### 6.2. Simulation

In the following, the variable *simulatedRound* denotes the next round to be simulated (we consider that the simulation has been performed up to round *simulatedRound*-1). Each process  $p_i$  maintains (1) the simulated state of machine  $\Pi_i$  at the end of round *simulatedRound*-1 (denoted by  $ss[\text{simulatedRound} - 1]$ ), and (2) the identifier associated with the state currently simulated at each process  $p_j$ , at the end of round *simulatedRound*-1, denoted by  $sim[j]$ .

If the condition of the *while* loop at line 22 ("*simulatedRound*-th instance of IC has decided") is true in a phase  $x$  at process  $p_i$ , then  $decision_i(\text{simulatedRound})$  denotes the decided vector of messages at line 24,  $failure_i(\text{simulatedRound})$  the value of the variable *failure* updated in line 25, and  $trueRcvd_i(\text{simulatedRound})$  the value of the variable *trueRcvd* updated in line 39. Process  $p_i$

uses the decided vector  $decision_i(simulatedRound)$  to update the simulated state of machine  $\Pi_i$ , i.e.,  $p_i$  adds  $simulatedRound$  in  $states$  and computes  $ss[simulatedRound]$ . More precisely,

- (1)  $p_i$  computes the messages  $trueRcvd$ : (1a) if  $p_j \in failure$  or  $decision[j] = \perp$ , then  $trueRcvd[j] = \perp$ , otherwise (1b)  $p_i$  searches for the tuple  $\langle sim[j], M, *, * \rangle$  in the extended message of  $p_j$  (generated at phase  $simulatedRound$ ), where  $M$  is the set of messages received in round  $simulatedRound-1$  (i.e., the previous value of  $trueRcvd$ ). Let  $\langle sim[j], M, s, m \rangle$  be this tuple.  $sim[j]$  is updated with  $s$  and  $trueRcvd[j]$  with  $m$ .
- (2)  $p_i$  updates  $ss[simulatedRound]$  with the state  $T_i(ss[simulatedRound - 1], trueRcvd, simulatedRound)$ .

If any value in the vector  $decision$  is  $\perp$ , then the corresponding process is added to  $failure$ . If any process adds itself to  $failure$ , it stops. In the algorithm, at line 43,  $\delta$  denotes the maximum number of phases for the underlying IC algorithm to decide in the system model in which the transformation algorithm is running. Indeed, a process that does not decide in an IC instance in  $\delta$  phases is faulty, and thus stops taking part to the simulation.

The following propositions assert the correctness of the extension of our transformation. We introduce here the extension for the uniform transformation, in *Crash* and *Omission* ( $t < n$ ), and *General-MAJ* ( $t < n/2$ ).

**Proposition 20.** *The simulation  $Sim$  for a run of  $\mathcal{T}(\Pi_s)$ ,  $R = \langle I, ST, MS, MR \rangle$  is defined by the run  $R' = \langle I', ST', MS', MR' \rangle$  as follows. Let  $p_i$  be a process in  $correct(R)$ . We consider the simulation of round  $r$  of  $R'$ , for any process  $p_j$ .*

- (i)  $I' = I$ .
- (ii) if  $p_j \in failure_i(r)$  then  $p_j \notin correct(R', r)$   
otherwise  $p_j \in correct(R', r)$ .
- (iii)  $ST'(i, 0) = s_i$  and  $ST'(i, r) = ss_i[r]$ . For any process  $p_j$  (including  $p_i$ ) not in  $failure_i(r)$ ,  $ST'(j, r)$  is the state of  $p_j$  at the end of round  $r$ , such that  $ST'(j, 0) = s_j$  and  $ST'(j, x) = T_j(ST'(j, x-1), trueRcvd_i(x), x)$ , for each  $x$  from 1 to  $r$ . Otherwise, for any  $p_j$  in  $failure_i(r)$ ,  $ST'(j, r) = ST'(j, r-1)$ .
- (iv) if  $p_j \in failure_i(r)$  then  $MS'(j, k, r) = \perp$ ,  
otherwise  $MS'(j, k, r) = trueRcvd_i(r)[j]$  for any process  $p_k \in \Omega$ .
- (v) if  $p_j \in failure_i(r)$  then  $MR'(j, r) = [\perp, \dots, \perp]$ ,  
otherwise  $MR'(j, r) = trueRcvd_i(r)$ .

**Proposition 21.** *The algorithm of Fig. 5 is a uniform shifting transformation (with an underlying uniform IC algorithm) from  $PSR(n, t)$  to  $Crash(n, t)$ ,  $Omission(n, t)$  ( $t < n$ ),  $General - MAJ(n, t)$  ( $t < n/2$ ) where the shift  $S$  is number of rounds needed to solve uniform interactive consistency in  $Crash(n, t)$ ,  $Omission(n, t)$  ( $t < n$ ), or  $General - MAJ(n, t)$  ( $t < n/2$ ).*

The same idea can be applied when input values can depend on the state of the processes, and there are finitely many possible input values (i.e.,  $|\mathcal{I}| < \infty$ ). Using the technique described above, a process anticipates on the different input values that it can receive, to start the next simulations.

```

1: failure :=  $\emptyset$  {failure corresponds to one round of the failure pattern}
2: simulatedRound := 1 {simulatedRound is the current simulated round number}
3: states :=  $\{0_i\}$ ; ss[0] :=  $s_i$  {set of states of protocol  $\Pi_s$  which are simulated by  $\Pi_w$  in the current round}
4:  $(\forall j \in [1, n])(\text{sim}[j] := 0)$  {all processes start in the 0-th state}
5: number := 1 {by convention  $\text{num}(\text{ss}[0]) = 0$ }

6: for phase r (r = 1, 2, ...) do
7:   input := receiveInput() {receive input value corresponding to  $I(i, r)$ }
8:   if r = 1 then
9:     es :=  $\{s_i\}$ ; em :=  $\{\langle -, -, 0, O_i(s_i, \text{input}, r) \rangle\}$ 
10:   else
11:     es' :=  $\emptyset$ ; em :=  $\emptyset$ 
12:     for any possible combination rec of n messages of rcvd[r - 1] do
13:       for any possible state st of es do
14:         st' :=  $T_i(st, \text{rec}, r)$ ; num(st') := number; number := number + 1
15:         es' := es'  $\cup \{st'\}$ 
16:         em := em  $\cup \{\langle \text{num}(st), \text{rec}, \text{num}(st'), O_i(st', \text{input}, r) \rangle\}$ 
17:       es := es'

18:   start instance r, and propose(em)
19:   execute one phase of all other running instances
20:   rcvd[r] := extended messages of instance r

21:   states :=  $\emptyset$ 
22:   while simulatedRound-th instance of IC has decided do {has any IC instance decided?}
23:     states := states  $\cup \{\text{simulatedRound}\}$  {instance simulatedRound has decided}
24:     decision := decision vector of instance simulatedRound {reconstruct patterns}
25:     failure := failure  $\cup \{p_j \mid \text{decision}[j] = \perp\}$  {ensure atomic failures}
26:     if  $p_i \in \text{failure}$  then {is process  $p_i$  faulty?}
27:       halt { $p_i$  does not perform any step}
28:     for each  $p_j \in \Omega$  do {adjust decision vector with previous failure pattern}
29:       if  $p_j \in \text{failure}$  then
30:         decision[j] :=  $\perp$ 
31:       for each  $p_j \in \Omega$  do {compose the messages of the round}
32:         if  $p_j \in \text{failure}$  then
33:           tmpRcvd[j] :=  $\perp$ 
34:         else
35:           if simulatedRound = 1 then
36:             tmpRcvd[j] := m such that  $\langle *, *, 0, m \rangle \in \text{decision}[j]$ 
37:           else
38:             let k and m such that  $\langle \text{sim}[j], \text{trueRcvd}, k, m \rangle \in \text{decision}[j]$ 
39:             tmpRcvd[j] := m; sim[j] := k
40:           trueRcvd := tmpRcvd
41:           ss[simulatedRound] :=  $T_i(\text{ss}[\text{simulatedRound} - 1], \text{trueRcvd}, \text{simulatedRound})$ 
42:           simulatedRound := simulatedRound + 1

43:   if r - simulatedRound  $\geq \delta$  then halt {is process  $p_i$  faulty?}

```

Fig. 5. Extended transformation algorithm (code for process  $p_i$ ).

When the preceding simulations terminate, the input value that had correctly anticipated the state of the process is determined, and only the messages and states following from this input value are kept. The algorithm in Fig. 5 can easily be adapted to the case where input values depend on the state of processes. Note that in both of the above cases, the number of messages generated is very high.

To prove Proposition 21, we first show that the construction of function *Sim* in Fig. 5 is consistent with Proposition 20. We proceed through a series of lemmas.

**Lemma 22.** *For any run  $R$  and any process  $p_i$  in  $\text{correct}(R)$ ,  $p_i$  never halts, and decides in all IC instances.*

**Proof.** A process  $p_i$  that is correct in  $R$  exists since  $t < n$  for  $\text{Crash}(n, t)$ ,  $\text{Omission}(n, t)$  or  $\text{General} - \text{MAJ}(n, t)$ . Thus,  $p_i$  does not halt in lines 27 or 43.) Process  $p_i$  is correct and never crashes. By the termination property of IC,  $p_i$  always decides in any instance of IC within a bounded number of phases. Thus  $p_i$  never halts in line 43. Moreover, by the validity property of IC, in any decision vector,  $\text{decision}[i] \neq \perp$ . Thus  $p_i$  never halts in line 27.  $\square$

**Lemma 23.** *For any run  $R$ , any process  $p_i$  in  $\text{correct}(R)$ , any process  $p_j$ , and any round  $r$  such that  $p_j$  decides in the  $r$ -th instance of IC, we have the following properties:*

- $\text{decision}_i(r) = \text{decision}_j(r)$
- $\text{failure}_i(r) = \text{failure}_j(r)$
- $\text{trueRcvd}_i(r) = \text{trueRcvd}_j(r)$
- $\text{sim}_i(r)[j] = \text{num}(ST'(j, r))$

**Proof.** By Lemma 22,  $p_i$  decides in all IC instances. In the algorithm,  $p_j$  decides in the  $r$ -th instance of IC if and only if  $p_j$  has decided in all previous instances. By the agreement property of IC, the decision is the same for  $p_i$  and  $p_j$ . We show the three last items by induction on  $r$ . It is easy to see that these properties hold for  $r = 1$ : for the first IC instance, every process proposes its initial value; when all processes terminates the first IC instance, they decide upon the same vector. This implies  $\text{decision}_i(1) = \text{decision}_j(1)$ ,  $\text{failure}_i(1) = \text{failure}_j(1)$ ,  $\text{trueRcvd}_i(1) = \text{trueRcvd}_j(1)$ , and  $\text{sim}_i(1)[j] = \text{num}(ST'(j, 1))$ , for any  $p_i \neq p_j$ . Assume the properties hold up to round  $r - 1$ . When  $p_i$  decides in the  $r$ -th instance, it adds a set of processes to  $\text{failure}_i$ . By the agreement property of IC,  $p_i$  and  $p_j$  add the same set of processes. By induction,  $\text{failure}_i(r) = \text{failure}_j(r)$ . As  $\text{decision}(r)$  and  $\text{failure}(r)$  are the same for all processes for which they are defined, by induction hypothesis, we have  $\text{trueRcvd}_i(r) = \text{trueRcvd}_j(r)$ , and  $\text{sim}_i(r)[j] = \text{num}(ST'(j, r))$ .  $\square$

We may once again define the simulation through the value of the variables of any correct process. Consider any run  $R$ . In the following,  $p_k$  denotes a correct process in  $R$  (we know there exists at least one as  $t < n$ ). We define some parts of the simulation through the variables of  $p_k$ . Note that, in contrast with the proof of Proposition 9, we cannot prove Proposition 21 only through the variables of  $p_k$ , because in the algorithm of Fig. 5, for instance,  $p_k$  does not keep the states in which the other processes are.

**Lemma 24.** *For any process  $p_i$ , any round  $r$  such that  $\text{decision}_i(r)$  and  $\text{decision}_i(r + 1)$  occur,  $\text{failure}_i(r) \subseteq \text{failure}_i(r + 1)$ .*

**Proof.** From the algorithm,  $\text{failure}_i$  always increases.  $\square$

**Lemma 25.**  $\text{correct}(R) \subseteq \text{correct}(R')$ .

**Proof.** By Lemma 22, all correct processes decide in all instances of IC. Thus they never halt in line 43. By the validity property of the underlying IC algorithm, the decision value is not  $\perp$  for any correct process in any decision vector. Thus no correct process ever halts in line 27. Therefore all correct processes in  $R$  are correct in  $R'$ .  $\square$



**Lemma 26.** For any process  $p_i$  and any protocol  $\Pi_s$  to simulate, let  $\langle s_i, T_i, O_i \rangle$  be the state machine for  $p_i$ . For any round  $r$  and any  $p_j \in \Omega$ , we have:

- (1)  $ST(i, 0) = s_i$ .
- (2) if  $p_i \in \text{correct}(R', r)$  then  $MS'(i, j, r) = O_i(ST'(i, r-1), I'(i, r), r)$ , otherwise  $MS'(i, j, r) = \perp$ .
- (3) if  $p_i \in \text{correct}(R', r)$  then  $MR'(i, j, r) = MS'(j, i, r)$ , otherwise  $MR'(i, j, r) = \perp$ .
- (4) if  $p_i \in \text{correct}(R', r)$  then  $ST'(i, r) = T_i(ST'(i, r-1), [MS'(1, i, r), \dots, MS'(n, i, r)], r)$ , otherwise  $ST'(i, r) = ST'(i, r-1)$ .

**Proof.** (1) is immediate, by the definition of  $ST'$  in Proposition 20. We prove (2) and (26) by induction. For the case  $r = 1$  and  $p_i \notin \text{correct}(R', 1)$ , then  $p_i \in \text{failure}_k(1)$ . By the algorithm  $\text{trueRcvd}_k(1)[i] = \perp$ , and by properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, 1) = \text{trueRcvd}_k(1)[i]$  and  $MR'(i, j, 1) = \text{trueRcvd}_k(1)[j]$ . By the algorithm and (26),  $ST'(i, 1) = ST'(i, 0)$ . If  $p_i \in \text{correct}(R', 1)$ , then  $p_i \notin \text{failure}_k(1)$ . By the properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, 1) = \text{trueRcvd}_k(1)[i]$  and  $MR'(i, j, 1) = \text{trueRcvd}_k(1)[j]$ . By line 39 of the algorithm,  $\text{trueRcvd}_k(1)[i] = O_i(ST'(i, 0), \text{decision}_k[i], 1)$ , and so  $MS'(i, j, 1) = O_i(ST'(i, 0), I'(i, 1), 1)$ . By line 39 of the algorithm,  $\text{trueRcvd}_k(1)[j] = MS'(j, k, 1)$ , and so  $MR'(i, j, 1) = MS'(j, i, 1)$ . By line 41 of the algorithm,  $ss_k(1)[k] = T_k(ss_k(0)[k], \text{trueRcvd}_k(1), 1)$ , and  $ST'(k, 1) = T_k(ST'(k, 0), [MS'(1, k, 1), \dots, MS'(n, k, 1)], 1)$ .

Assume the properties (26) and (26) are true up to round  $r-1$ . If  $p_i \notin \text{correct}(R', r)$ , then  $p_i \in \text{failure}_k(r)$ . By the properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, r) = \perp$ , and  $MR'(i, j, r) = [\perp, \dots, \perp]$ . By the algorithm and (26)  $ST'(i, r) = ST'(i, r-1)$ . If  $p_i \in \text{correct}(R', r)$ , then  $p_i \notin \text{failure}_k(r)$ . By the properties (iii) and (iv) of the definition of the simulation,  $MS'(i, j, r) = \text{trueRcvd}_k(r)[i]$  and  $MR'(i, j, r) = \text{trueRcvd}_k(r)[j]$ . By line 39 of the algorithm, and by the induction hypothesis,  $\text{trueRcvd}_k(r)[i] = O_i(ST'(i, r-1), \text{decision}_k[i], r)$ , and so  $MS'(i, j, r) = O_i(ST'(i, r-1), I'(i, r), r)$ . By line 39 of the algorithm,  $\text{trueRcvd}_k(r)[j] = MS'(j, k, r)$ , and so  $MR'(i, j, r) = MS'(j, i, r)$ . By line 41 of the algorithm, and by the induction hypothesis,  $ss_k(r)[k] = T_k(ST'(k, r-1), \text{trueRcvd}_k(r), r)$ , and so  $ST'(k, r) = T_k(ST'(k, r-1), [MS'(1, k, r), \dots, MS'(n, k, r)], r)$ .  $\square$

By Lemmas 23 and 26, the function  $\text{Sim}$  is well defined, and consistent with Proposition 20.

**Lemma 27.**  $R'$  is a run of  $\Pi_s$ .

**Proof.** Lemmas 23 and 24 show that  $R'$  is in  $\text{PSR}$  if  $R$  is in  $M$ , where  $M \in \{\text{Crash}, \text{Omission}, \text{General}\}$ . Lemma 26 shows that the functions  $ST'$ ,  $MS'$  and  $MR'$  consistently define a run of  $\Pi_s$ , with input value  $I'$ .  $\square$

**Lemma 28.** Let  $x$  be any phase,  $p_i$  any process and  $r$  any round. If  $r \in ST(i, x).\text{states}$  then  $ST(i, x).ss[r] = ST'(i, r)$ .

**Proof.** Each time a round  $r$  is added to  $\text{states}_i$  (line 23),  $ss_i[r] = T_i(ST'(i, r-1), \text{trueRcvd}_i(r), r)$  (line 29).  $\square$

**Lemma 29.** For any round  $r$  and any process  $p_i$  in  $\text{correct}(R')$ , if it exists  $c$  such that  $r \in ST(i, c).\text{states}$  then  $ST(i, c).ss[r] = ST'(i, r)$ .

**Proof.** If such a  $c$  exists, then  $p_i$  has decided in the  $r$ -th instance of IC, and by the algorithm in Fig. 5, in each previous instance. By Lemma 25, we have  $ST(i, c).ss[r] = ST'(i, r)$ .  $\square$

**Lemma 30.** *Let  $p_i$  be any process in  $\text{correct}(R')$  and  $r$  be any round. There exists a unique  $c \leq r + S$  such that  $r \in ST(i, c).states$  and  $ST(i, c).ss[r] = ST'(i, r)$ .*

**Proof.** By Lemma 10,  $p_i$  decides in all instance of IC. In particular, for the  $r$ -th instance,  $p_i$  decides in phase  $c$ . We pose  $S$  as the number of phases for an instance of IC to decide, thus  $c \leq S + r$ . We have  $r \in ST(i, c).states$ , and by Lemma 18,  $ST(i, c).ss[r] = ST'(i, r)$ .  $\square$

**End of proof of Proposition 21.** We show that our function  $Sim$ , as defined by Proposition 20, satisfies the seven properties of Definition 3, with  $f(r) = r + S$ , where  $S$  is the number of phases to solve IC in  $M_w$ .

Lemma 27 shows that  $sim(R)$  is a run of  $\Pi_s$ , which implies Property (i). Lemma 25 shows that  $\text{correct}(R) \subseteq \text{correct}(R')$ , which implies Property (ii). By the definition of the simulation, we have  $I = I'$ , which implies Property (iii). Property (iv) follows from Lemma 28. By Lemma 22, any correct process  $p_i$  of  $F$  decides in any instance of IC. In particular, in the  $r$ -th instance of IC, it decides at most at phase  $r + S$ . Lemma 30 shows that for any round  $r$  and any process  $p_i$  in  $\text{correct}(R')$ , there exists  $c \leq r + S$  such that  $r \in ST(i, c).states$  and  $ST(i, c).ss[r] = ST'(i, r)$  (Property (v)) and that this  $c$  is unique (Property (vi)). If a process adds any round  $r$  in its set  $states$ , by the algorithm, it has decided the  $r$ -th instance of IC, and all previous instances of IC as well. This implies Property (vii).

## 7. Complexity

We analyze the performance of our transformation technique in terms of message and phase complexities. For the rest of this section, we need to make a distinction between the number of phases before any process decides in a non-uniform, respectively uniform, IC algorithm, and the number of phases before any process halts in the same algorithm.

From the IC algorithms presented in Section 4, we have the following results:

- For non-uniform IC, and for any of the models *Crash*, *Omission*, and *General*, the number of phases for all processes to decide is  $\delta_{\text{non-uniform}} = f + 1$  in any run  $R$  with at most  $f$  failures, whereas the number of phases for all processes to halt is  $\tau_{\text{non-uniform}} = \min(f + 2, t + 1)$  in any run  $R$  with at most  $f$  failures.
- For uniform IC, and for any of the models *Crash*, *Omission*, and *General-MAJ*, the number of phases for all processes to decide and halt is  $\delta_{\text{uniform}} = \tau_{\text{uniform}} = t + 1$ , in any run  $R$ .

### 7.1. Message complexity

In terms of messages, the first transformation generates at most a  $n \log_2 |I|$ -bit message per process, per phase, and per IC instance. As there are  $\tau_{\text{non-uniform}}$  non-uniform, respectively  $\tau_{\text{uniform}}$  uniform, IC instances running in parallel, any process sends a  $n \tau \log_2 |I|$ -bit message in any phase, in any run, where  $\tau$  is either  $\tau_{\text{non-uniform}}$  or  $\tau_{\text{uniform}}$ , whether we consider the non-uniform or the uniform transformation.

In the extended, uniform transformation, any process maintains at least  $2^{n\tau_{\text{uniform}}}$  states for a round simulation. A state (tuple) is coded using  $\sigma = 2\log_2|\mathcal{S}| + (n+1)\log_2|\mathcal{M}|$  bits. As there are  $\tau_{\text{uniform}}$  instances of the uniform IC algorithm running in parallel, any process sends a  $\tau_{\text{uniform}}n\sigma 2^{n\tau_{\text{uniform}}}$ -bit message in any phase, in any run.

Determining the exact overhead in terms of message size complexity is an open issue, as is the tight lower bound on the message overhead for a automatic shifting transformation technique.

## 7.2. Phase complexity

The *phase complexity overhead* is defined as the number of additional phases executed by the transformed protocol  $\Pi_w$  in  $M_w$ , compared with the original protocol  $\Pi_s$  in  $M_s$ .

In any of our shifting transformation algorithms, the simulation of consecutive rounds is overlapped, such that the simulation of two consecutive rounds start with an interval of a single phase. Thus, for any of our shifting transformation algorithm, the only phase complexity overhead is the number of phases for obtaining the outcome of the simulation for the first round, corresponding to  $\delta_{\text{non-uniform}} - 1$  or  $\delta_{\text{uniform}} - 1$ , depending on the transformation considered.

We observe that for the non-uniform transformation with  $\delta_{\text{non-uniform}} = f + 1$ , the phase complexity overhead is just  $\delta_{\text{non-uniform}} - 1 = f$ . In any failure-free ( $f = 0$ ) run  $R$ , there is thus no phase complexity overhead for the non-uniform transformation. This transformation provides the outcome of the simulation in real-time, as the run executes. Had we try to simulate a weaker model than *PSR* by using our non-uniform shifting transformation, we would not have gained any improvement in the phase complexity in failure-free runs.

## 8. Concluding remarks

In this paper, we have concentrated on models *Crash*, *Omission*, *General* and *General-MAJ*, and *Byzantine* and have presented different shifting transformation techniques to translate protocols from the perfectly synchronous model *PSR* into each of these weaker models. We first presented a simple shifting transformation algorithm (in which each process knows that state machine executed by any other process) that allows for both a uniform and a non-uniform transformation, and we have then extended the uniform transformation to a more sophisticated transformation algorithm (in which any process does not know the state machine executed by any other process).

We show in the paper that the complexity of the transformation in terms of rounds is optimal, in two precise senses. First the round overhead to simulate a single IC instance is optimal since we need just a single *PSR*. Second, the non-uniform transformation provides real-time outputs of the simulation, and thus had we try to simulate a weaker model than *PSR* by using our shifting transformation, we would not have gained any improvement in the round complexity.

We leave open the question of the message complexity. We have characterized the message complexity obtained in our solution, but the optimal message complexity for a shifting transformation is open, as is the question of finding the corresponding shifting transformation.

## A. Correctness of the interactive consistency algorithms

### A.1. Non-uniform interactive consistency in General

We give in Fig. 1 an algorithm solving non-uniform IC in *General*, *Omission*, and *Crash*, for  $t < n$ . We prove in this section that the algorithm in Figure 1 satisfies the specification of non-uniform IC, through a series of lemmas (Lemmas 31–33).

**Lemma 31.** *In the algorithm in Fig. 1, all the correct processes decide on a vector of values by round  $f + 1$ , and halt by round  $\min(f + 2, t + 1)$ .*

**Proof.** If  $f = 0$ , all the processes are correct. Any correct process  $p_i$  sends its vector  $V_i$  of values in round 1 with its value  $v_i$  (line 4), receives the vector from every other process at the end of round 1 (line 7), decides on its vector of values at the end of round 1 (line 12) and halts in round 2 (line 5).

Now assume  $1 \leq f \leq t$ . We distinguish two cases:

- (1) Some correct process  $p_i$  decides on its vector  $V_i$  in round  $f$ . In this case, process  $p_i$  sends its vector  $V_i$  to all the processes in round  $f + 1$  (line 4), thus every correct process decides at the end of round  $f + 1$  (line 12) and halts by round  $\min(f + 2, t + 1)$  (line 5).
- (2) No correct process decides on its vector in round  $f$ . Thus no correct process halts by round  $f + 1$ . In all rounds  $1 \leq k \leq f + 1$ , every correct process sends a message to all processes, in particular to all correct processes (line 4). Consider any correct process  $p_i$  at the end of round  $f + 1$ . As all the correct processes send a message to all the processes in round  $1, \dots, f + 1$ , the set  $quiet_i$  for  $p_i$  at the end of round  $f + 1$  contain faulty processes only. Thus,  $quiet_i \leq f < f + 1$  at the end of round  $f + 1$ . By the algorithm,  $p_i$  fills the missing values in  $V_i$  with  $\perp$  (line 11), then decides on its vector  $V_i$  at the end of round  $f + 1$  (line 12), and halts in round  $\min(f + 2, t + 1)$  (line 5).  $\square$

**Lemma 32.** *In the algorithm in Fig. 1, for any decision vector  $V$ , the  $j$ th component of  $V$  is either the value proposed by  $p_j$  or  $\perp$ . For any vector  $V$  decided by any correct process, the  $j$ th component of  $V$  is  $\perp$  only if  $p_j$  fails.*

**Proof.** If any process  $p_i$  is correct,  $p_i$  sends its vector with its value  $v_i$  to all the processes in round 1 (line 4). Every correct process receives  $p_i$ 's vector and copies the value  $v_i$  into its own vector (line 7). Thereafter, the value does not change anymore.  $\square$

**Lemma 33.** *In the algorithm in Fig. 1, if any correct process decides on a vector of values  $V$ , then every correct process eventually decides on the same vector of values  $V$ .*

**Proof.** We proceed through a series of claims.  $\square$

**Claim 1.** *Any process decides on at most one vector.*

**Proof.** From the algorithm in Fig. 1, it is clear that any process may decide at most on a single vector.  $\square$

**Claim 2.** *If some correct process decides on a vector  $V$  such that  $V[k] = v_k \neq \perp$  for some  $k$ , then some correct process copies  $v_k$  (line 7 in Fig. 1) in its own vector before round  $t + 1$ .*

**Proof.** Suppose a process decides on a vector  $V$  at the end of round  $t + 1$  such that  $V[k] = v_k \neq \perp$  for some  $k$ . This means that there exist  $t + 1$  processes  $q_1, \dots, q_{t+1}$ , such that  $q_i$  sends a vector  $Q_i$  in round  $i$  such that  $Q_i[k] = v_k$ . One of them, say  $q_k$ , is necessarily correct. If  $k = 1$ ,  $q_1$  has  $v_1 = v_k$  as its initial value, at the beginning of round 1. For  $2 \leq k \leq t + 1$ ,  $q_k$  copies  $v_k$  in round  $k - 1$ . In both cases,  $q_k$  copies  $v_k$  into its own vector before round  $t + 1$ .  $\square$

**Claim 3.** *If any correct process copies  $v_k$  into its vector  $V_i$  for some  $k$ , then no correct process ever sets  $V_j[k] := \perp$ .*

**Proof.** Let  $p_i$  be the first correct process that sets  $V_i[k] := v_k$  into its vector  $V_i$  of values, for some  $k$ . By Claim 2,  $p_i$  copies  $v_k$  in some round  $l$ ,  $l < t + 1$ . This implies the existence of  $l$  processes  $q_1, \dots, q_l$ , such that  $q_x$  sends a vector  $Q_x$  in round  $x$  such that  $Q_x[k] = v_k$ . We distinguish two subcases, namely rounds 1 to  $l$ , and rounds  $l + 1$  to  $t + 1$ :

- (1) Suppose now by contradiction that some processes set  $V[k] := \perp$  in their vector of values, and denote  $p_j$  the first process that sets  $V_j[k] := \perp$  in some round  $m \leq l$ . In any round  $i$ ,  $1 \leq i \leq m$ ,  $p_j$  does not receive the vector of values  $Q_i$  from  $q_i$ . Thus, at the end of round  $m$ , the set  $|quiet_j|$  includes processes  $q_1, \dots, q_m$ , and thus  $|quiet_j| \geq m$  at the end of round  $m$ . As  $p_j$  is the first process that sets  $V_j[k] := \perp$ ,  $p_j$  necessarily does it because  $|quiet_j| < m$  at the end of round  $m$  — a contradiction.
- (2) As  $p_i$  copies  $v_k$  in round  $l$ ,  $p_i$  sends its vector of values  $V_i$  such that  $V_i[k] = v_k$  in round  $l + 1$ . From subcase 1 above, no process sends a vector of values  $V$  such that  $V[k] = \perp$  in round  $l + 1$ . Thus, all the correct processes receive  $V_i$  in round  $l + 1$  and set  $V[k] = v_k$ . Thereafter, the value  $V[k]$  at any correct process may not change anymore. Thus no correct process  $p_j$  ever sets  $V_j[k] := \perp$ .  $\square$

By Lemma 31 and Claim 1, every correct process eventually decides on at most one vector of values. By Lemma 32, for any decided vector  $V$  and for any  $k$ ,  $V[k]$  is either  $\perp$  or  $p_k$ 's value. For any two vectors of values  $V_i$  and  $V_j$ , respectively, decided by two correct processes  $p_i$  and  $p_j$ , and for any  $k$ ,  $V_i[k] = V_j[k]$  then directly follows from Claims 2 and 3.

#### A.2. Uniform interactive consistency in Omission and Crash

The algorithm in Fig. 2 solves IC in *Omission* and *Crash*. In both models we assume  $t < n$ . In the algorithm, all the processes that decide, decide after  $t + 1$  rounds. We show in this section that the algorithm in Fig. 2 satisfies the specification of Uniform IC, through a series of lemmas (Lemma 34–36).

**Lemma 34.** *In the algorithm in Fig. 2, every correct process eventually decides.*

**Proof.** It is clear that the algorithm runs for exactly  $t + 1$  rounds at each process, and no process may ever block while executing the algorithm. Thus every process that does not crash, including any correct process, decides at the end of round  $t + 1$ .  $\square$

**Lemma 35.** *In the algorithm in Fig. 2, for any decision vector  $V$ , the  $j$ th component of  $V$  is either the value proposed by  $p_j$  or  $\perp$ , and is  $\perp$  only if  $p_j$  fails.*

**Proof.** We observe that the processes only decide at line 6, and each process  $p_i$  decides on its vector  $V_i$  of estimate values. Throughout the algorithm, the coordinate  $V_x[j]$  at any process  $p_x$ , corresponding to the estimate value proposed by any process  $p_j$ , is assigned with either  $\perp$  or  $p_j$ 's initial value (at line 1), or  $p_k$ 's component from another vector of estimate values (at line 5). Hence, any  $V_x[j]$  component at any process  $p_x$  may only contain  $p_j$ 's initial value or  $\perp$ .

If  $p_j$  is a correct process and thus never crashes,  $p_j$  sends its initial vector of estimate values (which contains its initial value  $v_j$ ) to all the other processes in round 1, and thereafter, every process  $p_i$  that reaches the end of round 1, receives  $p_j$ 's message at line 5 and assigns  $V_i[j]$  with  $v_j$ . (In the algorithm, the expression  $(j - r + 1)$  is evaluated to  $j$  at any process  $p_i$ .) In any subsequent round  $r$ , the  $j$ th component of any vector  $V_i$  for any process  $p_i$  does not change.  $\square$

**Lemma 36.** *In the algorithm in Fig. 2, no two processes decide on different vectors of estimate values.*

**Proof.** We show that all the processes that decide a vector of estimate values at line 6, decide on the same component for any process  $p_i$ .

Suppose by contradiction that there exist two processes  $p_a$  and  $p_b$  that, respectively, decide on the vectors of estimate values  $V_a$  and  $V_b$ , such that  $V_a[i] \neq V_b[i]$ . Furthermore, assume without loss of generality that  $V_a[i] = v_i$  and  $V_b[i] = \perp$ .

According to the algorithm in Fig. 2, and in particular to line 5, all processes copy  $p_i$ 's estimate value from  $p_i$ 's vector in round 1, from  $p_{i+1}$ 's vector in round 2, ..., from  $p_{(i+t) \bmod n}$ 's vector in round  $t + 1$ . Among processes  $p_i$  to  $p_{(i+t) \bmod n}$ , one is necessarily correct, say  $p_{(i+x) \bmod n}$  ( $0 \leq x \leq t$ ), as there are at most  $t$  processes that may ever commit a send omission in any execution.

In round  $x + 1$  ( $1 \leq x + 1 \leq t + 1$ ), every process  $p_j$ , and in particular  $p_a$  and  $p_b$ , receives  $p_{(i+x) \bmod n}$ 's vector of estimate values, and copies  $p_i$ 's estimate value from  $V_{(i+x) \bmod n}$  to  $V_j$ . Thereafter, all the processes have the same estimate value for  $p_i$ 's value. There are two cases, whether  $V_{(i+x) \bmod n}[i]$  is  $v_i$  or is  $\perp$ :

- $V_{(i+x) \bmod n}[i] = v_i$ . Every process, in particular  $p_b$ , receives  $p_{(i+x) \bmod n}$ 's message. Thus  $p_b$  assigns  $V_b[i]$  with  $v_i$ . Thereafter,  $p_b$  only receives  $v_i$  in the  $i$ th component of any vector of estimate values. A contradiction with the fact that  $p_b$  decides on  $V_b$  and  $V_b[i] = \perp$ .
- $V_{(i+x) \bmod n}[i] = \perp$ . Every process, in particular  $p_a$ , receives  $p_{(i+x) \bmod n}$ 's message. Thus  $p_a$  assigns  $V_a[i]$  with  $\perp$ . Thereafter,  $p_a$  only receives  $\perp$  in the  $i$ th component of any vector of estimate values. A contradiction with the fact that  $p_a$  decides on  $V_a$  and  $V_a[i] = v_i$ .  $\square$

### A.3. Uniform interactive consistency in General-MAJ

The algorithm in Fig. 3 solves IC in *General-MAJ*. We show that the algorithm satisfies termination, validity and uniform agreement.

**Lemma 37.** *In the algorithm in Figure 3, for any decision vector  $V$ , the  $j$ th component of  $V$  is either the value proposed by  $p_j$  or  $\perp$ , and is  $\perp$  only if  $p_j$  fails.*

**Proof.** We observe that the processes only decide at line 14, and each process  $p_i$  decides on its vector  $V_i$  of estimate values. Throughout the algorithm, the coordinate  $V_x[j]$  at any process  $p_x$ , corresponding to the estimate value of any process  $p_j$ , is assigned with  $p_j$ 's initial value (at line 3), or  $p_k$ 's component from another vector of estimate values (at line 12). Hence, any  $V_x[j]$  component may only contain  $p_j$ 's initial value, or  $\perp$ .

If  $p_j$  is a correct process and thus never crashes,  $p_j$  sends its initial vector of estimate values (which contains its initial value  $v_j$ ) to all the other processes in round 1, and thereafter, every process  $p_i$  that reaches the end of round 1, receives  $p_j$ 's message at line 6 and assigns  $V_i[j]$  with  $v_j$ . In any subsequent round  $r$ , the  $j$ th component of any vector  $V_i$  for any process  $p_i$  does not change.  $\square$

**Lemma 38.** *In the algorithm in Figure 3, every correct process eventually decides.*

**Proof.** Suppose by contradiction that there is a correct process  $p_i$  that does not decide. Since  $p_i$  is correct,  $p_i$  completes round  $t + 1$ . Furthermore, since  $p_i$  does not decide,  $|halt_i^{t+1} \cup suspect_i^{t+1}| \geq t + 1$ . Consider any process  $p_j$  in the set  $halt_i^{t+1} \cup suspect_i^{t+1}$ . There is at least one round  $k \leq t + 1$  in which either  $p_i$  does not receive a message from  $p_j$ , or  $p_j$  does not receive a message from  $p_i$ . Since  $p_i$  is correct, (1) if  $p_i$  does not receive a message from  $p_j$  in round  $k$ , then either  $p_j$  crashes in round  $r$ , or commits a send-omission in round  $k$ , or (2) if  $p_j$  does not receive a message from  $p_i$  in round  $k$ , then  $p_j$  commits a receive omission in round  $k$ . In both cases,  $p_j$  is a faulty process. Therefore,  $|halt_i^{t+1} \cup suspect_i^{t+1}| \geq t + 1$  implies there are more than  $t$  faulty processes. A contradiction.  $\square$

**Lemma 39.** *In the algorithm in Figure 3, no two processes decide on different vectors of estimate values.*

**Proof.** Suppose by contradiction that there are two distinct processes  $p_a$  and  $p_b$  which decide on distinct vectors, respectively,  $V_a$  and  $V_b$ . Without loss of generality, consider that there exists  $k$ , such that  $V_a[k] = c \neq V_b[k] = d$ . Hence, at the end of round  $t + 1$ , we have  $V_a[k] = c$  and  $V_b[k] = d$ , and  $|halt_a \cup suspect_a| \leq t$  and  $|halt_b \cup suspect_b| \leq t$ .

For any process  $p_i$ , the  $k$ th component of  $V_i$ ,  $V_i[k]$ , is only assigned with the value  $V_j[k]$  from another process  $p_j$ , and is assigned with  $p_k$ 's initial value by  $p_k$  at the initialization of the algorithm. Hence,  $V_i[k]$  may only contain  $v_k$  or  $\perp$ . Assume without loss of generality that  $V_a[k] = v_k$  and that  $V_b[k] = \perp$ . For every run of the algorithm in Figure 3, let  $C_0 = \{p_k\}$  and  $C_x$  ( $1 \leq x \leq t + 1$ ) be the set of every process  $p_l$  such that  $V_l[k] = v_k$  at the end of round  $x' \leq x$ . From the definition of  $C_x$ , we immediately observe that:

- (1) O1: For  $0 \leq x \leq t + 1$ ,  $C_x \subseteq C_{x+1}$ . This follows from the fact that as soon as any process  $p_l$  assigns  $v_k$  to  $V_l[k]$ ,  $p_l$  keeps  $V_l[k]$  unchanged.
- (2) O2: For  $0 \leq x \leq t + 1$ ,  $\forall p_l \in C_x$ , if  $p_l$  sends its vector of estimate values in round  $x' > x$ , then  $V_l[k] = v_k$ .

In the following, we prove five lemmas (Lemma 40 to Lemma 44) based on these assumptions. Lemma 44 contradicts Lemma 41.  $\square$

**Lemma 40.** *Consider any process  $p_l$  after completing round  $k$  ( $1 \leq k \leq t + 1$ ). Let  $senderMsg_l^k$  be the processes from which  $p_l$  receives a message in round  $k$  and that do not already belong to  $halt_l^{k-1}$ . Then  $senderMsg_l^k = \Omega \setminus halt_l^k$ .*

**Proof.** Consider any process  $p_l$  that reaches the end of round  $k$ , for  $1 \leq k \leq t + 1$ , and consider any other process  $p_m \in \Omega$  distinct of  $p_l$ . There are three exhaustive and mutually exclusive cases

regarding the message from  $p_m$  to  $p_l$  in round  $k$ : (1) if  $p_l$  does not receive any message from  $p_m$ , then  $p_l$  inserts  $p_m$  in  $halt_l$  (line 10); (2) if  $p_l$  receives a message from  $p_m$  in round  $k$  and  $p_m \notin halt_l^{k-1}$ , then  $p_m \in senderMsg_l^k$  and  $p_m \notin halt_l^k$ ; (3) if  $p_l$  receives a message from  $p_m$  in round  $k$  and  $p_m \in halt_l^{k-1}$ , then  $p_m \notin senderMsg_l^k$  and  $p_m \in halt_l^k$  (line 6). Thus any process  $p_m$  is either in  $senderMsg_l^k$  or in  $halt_l^k$ .  $\square$

**Lemma 41.**  $|C_t| \leq t$ .

**Proof.** Suppose by contradiction that  $|C_t| > t$ . Consider the message sent by any process  $p_m \in C_t$  to  $p_b$ . From the observation O2 here above, it follows that either  $p_b$  receives from  $p_m$  a message with  $V_m[k] = v_k$  in round  $t + 1$ , or  $p_b$  does not receive any message from  $p_m$  in round  $t + 1$  due to some failure.

Now consider the messages received by  $p_b$  in round  $t + 1$ . The set  $senderMsg_b^{t+1}$  does not contain the message from  $p_m$ , otherwise  $p_b$  sets  $V_b[k] = v_k$ . Therefore, from Lemma 40,  $p_m \in halt_b^{t+1}$ . As a result,  $C_t \subseteq halt_b^{t+1}$  and  $|halt_b \cup suspect_b| \geq |halt_b| \geq |C_t| > t$ : a contradiction.  $\square$

**Lemma 42.**  $p_a \in C_{t+1}$  and  $p_a \notin C_{t-1}$ .

**Proof.** Suppose by contradiction that  $p_a \in C_{t-1}$ . Consider any process  $p_m \in \Omega \setminus C_t$  which sends a message to  $p_a$  in round  $t + 1$ . From the definition of  $\Omega \setminus C_t$ ,  $V_m[k] = \perp$ . Therefore,  $p_m$  does not receive any message from  $p_a$  in round  $t$ . Hence, from Lemma 40,  $p_a \in halt_m^t$ . Therefore, from each process in  $\Omega \setminus C_t$ ,  $p_a$  either receives a message  $[V', Halt']$  in round  $t + 1$ , such that  $p_a \in Halt'$ , or  $p_a$  does not receive a message due to some failure. Thus, every process in  $\Omega \setminus C_t$  is either in  $suspect_a^{t+1}$  or is in  $halt_a^{t+1}$ . Consequently,  $\Omega \setminus C_t \subseteq halt_a^{t+1} \cup suspect_a^{t+1}$ . From Lemma 41, it follows that  $|\Omega \setminus C_t| \geq n - t > t$  (recall that  $t < \frac{n}{2}$  in *Omission*( $n, t$ )). So  $|halt_a^{t+1} \cup suspect_a^{t+1}| \geq |\Omega \setminus C_t| > t$ : a contradiction.  $\square$

**Lemma 43.** For all  $k$  such that  $0 \leq k \leq t - 1$ ,  $C_k \subset C_{k+1}$ .

**Proof.** Consider any  $0 \leq k \leq t - 1$ . From the observation O2 here above,  $C_k \subseteq C_{k+1}$ . So, either  $C_k \subset C_{k+1}$ , or  $C_k = C_{k+1}$ . Suppose by contradiction  $C_k = C_{k+1}$ .

For any process  $p_m \in \Omega \setminus C_{k+1}$ ,  $p_m$  does not receive any message from any process in  $C_k$ ; otherwise,  $p_m$  sets  $V_m[k] = v_k$  and  $p_m \in C_{k+1}$ . Therefore, from Lemma 40,  $C_k \subseteq halt_m^{k+1}$ . Since  $C_k = C_{k+1}$ , so for every process  $p_m \in \Omega \setminus C_{k+1}$ ,  $C_{k+1} \subseteq halt_m^{k+1}$ . Thus, in every round higher than  $k + 1$ , the processes in  $\Omega \setminus C_{k+1}$  ignore all the messages from any process in  $C_{k+1}$  while updating their vector of estimate values. For any process  $p_m \in \Omega \setminus C_{k+1}$ ,  $V_m[k] = \perp$ . Thus, after  $k + 1$  rounds, the set  $C$  never changes (no process in  $\Omega \setminus C$  ever assigns  $V[k]$  with  $v_k$ ), i.e.,  $C_{k+1} = C_{k+2} = \dots = C_{t+1}$ . A contradiction with Lemma 42.  $\square$

**Lemma 44.**  $|C_t| \geq t + 1$ .

**Proof.** Lemma 43 implies that for every  $0 \leq k \leq t - 1$ ,  $|C_{k+1}| - |C_k| \geq 1$ . We know that  $C_0 = \{p_k\}$  and thus  $|C_0| \geq 1$ . Therefore,  $|C_t| \geq t + 1$ . (A contradiction with Lemma 41.)  $\square$

## References

- [1] R.A. Bazzi, G. Neiger, Simplifying fault-tolerance: providing the abstraction of crash failures, *Journal of the ACM* 48 (3) (2001) 499–554.



- [2] T. Chandra, S. Toueg, Unreliable failure detectors for reliable distributed systems, *Journal of the ACM* 43 (2) (1996) 225–267.
- [3] B. Charron-Bost, R. Guerraoui, A. Schiper, Synchronous system and perfect failure detector: solvability and efficiency issues, in: *Proceedings of the IEEE International Conference on Dependable Systems and Networks*, 2000, pp. 523–532.
- [4] D. Dolev, R. Reischuk, H.R. Strong, ‘Eventual’ is earlier than ‘Immediate’, in: *Proceedings of the 23rd IEEE Symposium on Foundations of Computer Science (FOCS’82)*, 1982, pp. 196–203.
- [5] D. Dolev, R. Reischuk, H.R. Strong, Early stopping in Byzantine agreement, *Journal of the ACM* 37 (4) (1990) 720–741.
- [6] P. Dutta, A short note on uniform consensus in message omission model, Technical Report EPFL/IC, EPFL, 2002.
- [7] M.J. Fischer, N.A. Lynch, A lower bound for the time to assure interactive consistency, *Information Processing Letters* 14 (4) (1982) 183–186.
- [8] M.J. Fischer, N.A. Lynch, M.S. Paterson, Impossibility of distributed consensus with one faulty process, *Journal of the ACM* 32 (2) (1985) 374–382.
- [9] J. Garay, Y. Moses, Fully polynomial byzantine agreement for  $n > 3t$  processors in  $t + 1$  rounds, *SIAM Journal of Computing (SICOMP)* 27 (1) (1998).
- [10] R. Guerraoui, On the hardness of failure-sensitive agreement problems, *Information Processing Letters* 79 (2) (2001) 99–104.
- [11] R. Guerraoui, P. Kouznetsov, B. Pochon, A note on set agreement with omission failures, *Electronic Notes in Theoretical Computing Science* 81 (2003).
- [12] V. Hadzilacos, Byzantine agreement under restricted types of failures (not telling the truth is different from telling lies). Technical Report 18-83, Department of Computer Science, Harvard University, 1983.
- [13] L. Lamport, Time, clocks, and the ordering of events in a distributed system, *Communications of the ACM* 21 (7) (1978).
- [14] L. Lamport, M. Fischer, Byzantine generals and transaction commit protocols. Technical Report 62, SRI International, 1982.
- [15] L. Lamport, R. Shostak, L. Pease, The Byzantine Generals problem, *ACM Transactions on Programming Languages and Systems* 4 (3) (1982) 382–401.
- [16] Y. Moses, M.R. Tuttle, Programming simultaneous actions using common knowledge, *Algorithmica* 3 (1) (1988) 121–169.
- [17] G. Neiger, S. Toueg, Automatically increasing the fault-tolerance of distributed algorithms, *Journal of Algorithms* 11 (3) (1990) 374–419.
- [18] L. Pease, R. Shostak, L. Lamport, Reaching agreement in presence of faults, *Journal of the ACM* 27 (2) (1980) 228–234.
- [19] M. Raynal, Consensus in synchronous systems: a concise guided tour, in: *Proceedings of the 2002 Pacific Rim International Symposium on Dependable Computing (PRDC’02)*, 2002.
- [20] M.C. Roşu, Early-stopping terminating reliable broadcast protocol for general-omission failures, in: *Proceedings of the 15th ACM Symposium on Principles of Distributed Computing (PODC’96)*, ACM Press, New York, NY, USA, 1996, pp. 209.
- [21] F.B. Schneider, Replication management using the state machine approach, in: S. Mullender (Ed.), *Distributed Systems*, Addison-Wesley, Reading, MA, 1993.