# Formal Verification of Multi-Paxos for Distributed Consensus

Saksham Chand, Yanhong A. Liu, and Scott D. Stoller

Stony Brook University, Stony Brook, New York, 11794, USA {schand, liu, stoller}@cs.stonybrook.edu

January 4, 2019

#### Abstract

This paper describes formal specification and verification of Lamport's Multi-Paxos algorithm for distributed consensus. The specification is written in TLA<sup>+</sup>, Lamport's Temporal Logic of Actions. The proof is written and automatically checked using TLAPS, a proof system for TLA<sup>+</sup>. The proof checks safety property of the specification. Building on Lamport, Merz, and Doligez's specification and proof for Basic Paxos, we aim to facilitate the understanding of Multi-Paxos and its proof by minimizing the difference from those for Basic Paxos, and to demonstrate a general way of proving other variants of Paxos and other sophisticated distributed algorithms. We also discuss our general strategies for proving properties about sets and tuples that helped the proof check succeed in significantly reduced time.

**Keywords.** Distributed Algorithms, Formal Methods, Verification

## 1 Introduction

Distributed consensus is a fundamental problem in distributed computing. It requires that a set of processes agree on some value or values. Consensus is essential when distributed services are replicated for fault-tolerance, because non-faulty replicas must agree. Examples include leader election, atomic broadcast, and state machine replication in replicated data storage services like Google File System, Apache ZooKeeper, Amazon DynamoDB, etc. Unfortunately, consensus is difficult when processes or communication channels may fail.

Paxos [19] is an important algorithm, developed by Lamport, for solving distributed consensus. Basic Paxos is for agreeing on a single value, such as whether to commit a database transaction. Multi-Paxos is for agreeing on a continuing sequence of values, for example, a stream of commands to execute. Multi-Paxos has been used in many important distributed services, for example, Google's Chubby [1, 3] and Microsoft's Autopilot [13]. There are other Paxos variants, for example, variants that reduce a message delay [22] or add preemption [20], but Multi-Paxos is the most important in making Paxos practical for distributed services that must perform a continuing sequence of operations.

Paxos handles processes that run concurrently without shared memory, where processes may crash and may later recover, and messages may be lost or delayed indefinitely. In Basic Paxos, each process may repeatedly propose some value, and wait for appropriate replies from appropriate subsets of the processes while also replying appropriately to other processes; consensus is reached eventually if enough processes and channels are non-faulty to vote on some proposal thus agreeing on the proposed value. In Multi-Paxos, many more different attempts, proposals, and replies may happen in overlapping fashions to reach consensus on values in different slots in the continuing sequence.

Paxos has often been difficult to understand since it was created in the late 1980s [24]. Lamport later wrote a much simpler description of the phases of the algorithm but only for Basic Paxos [20]. Lamport et al. [25] wrote a formal specification and proof of Basic Paxos in TLA<sup>+</sup> [21] and TLAPS [34]. Many efforts, especially in recent years, have been spent on formal specification and verification of Multi-Paxos, but they use more restricted or less direct language models, some mixed in large systems with many unrelated functionalities, or handle other variants of Paxos than Multi-Paxos, as discussed in Section 7. What is lacking is formal specification and proof of the exact phases of Multi-Paxos, in a most direct and general language like TLA<sup>+</sup> [21], with a complete proof that is mechanically checked, and a general method for doing such specifications and proofs in a more feasible way.

This article addresses this challenge. We describe a formal specification of Multi-Paxos written in TLA<sup>+</sup>, and a complete proof written and automatically checked using TLAPS. Building on Lamport et al.'s specification and proof for Basic Paxos, we aim to facilitate the understanding of Multi-Paxos and its proof by minimizing the difference from those for Basic Paxos. The key change in the specification is to replace operations involving two numbers with those involving a set of 3-tuples, for each of a set of processes, exactly capturing the minimum conceptual difference between Basic Paxos and Multi-Paxos. However, the proof becomes significantly more difficult because of the handling of sets and tuples in place of two numbers.

This work also aims to show the minimum-change approach as a general way of specifying and verifying other variants of Paxos, and more generally of specifying and verifying other sophisticated algorithms by starting from the basics. We demonstrate this by further showing the extension of the specification and proof of Multi-Paxos to add preemption—letting processes abandon proposals that are already preempted by other proposals [20, 37]. We also extended the specification and proof of Basic Paxos with preemption, which is even easier.

Finally, we discuss a general method that we followed to tackle tedious and difficult proof obligations involving sets and tuples, a well-known significant complication in general. For difficult properties involving sets, we use induction and direct the prover to focus on the changes in the set values. For properties involving tuples, we change the ways of accessing and testing the elements to yield significantly reduced proof-checking time. Overall, we were able to keep the specification minimally changed, and keep the proof-checking time to about 3 minutes for both specifications while the prover checks the proofs for 779 obligations for Multi-Paxos and over 825 obligations for Multi-Paxos with Preemption.

This article is a corrected, improved, and extended version of [2]. The main changes are as follows.

- 1. The claim of a complete proof in [2] was incorrect, and the problem discovered is now fixed. The problem was due to an undocumented bug [29] in TLAPS that we discovered after the work in [2], which made us realize that the proof of Multi-Paxos with Preemption was incomplete. We fixed this by adding the missing proof. This is described in the new Section 5.1.
- 2. The proofs are simplified, shortened by 19% for Multi-Paxos and 17% for Multi-Paxos

with Preemption, even with the added proof to overcome the TLAPS bug discovered. In fact, it was during simplification of the proof that we discovered the bug. The simplifications are described in the new Section 5.2.

- 3. Sections 4.1, 4.3 and 4.4 are extended to define and explain all auxiliary predicates, acceptor invariants, and message invariants, respectively, that are used in the proof. Section 6 is extended with detailed results about the different versions of proofs. Section 7 is expanded with additional related works and more details about other proofs.
- 4. Section 5 is extended with a summary of the overall proof, in the new Section 5.3.
- 5. The complete, revised, and simplified TLA<sup>+</sup> specification and TLAPS-checked proof of Multi-Paxos with Preemption are added in the new Appendix C, including about 1.5 pages of specification, 1 page of invariants, and 9 pages of proof.

The rest of the paper is organized as follows. Section 2 covers preliminaries: distributed consensus (Section 2.1), Paxos (Section 2.2), TLA<sup>+</sup> (Section 2.3), and TLA Proof System, TLAPS (Section 2.4). Section 3 presents the TLA<sup>+</sup> specification of Multi-Paxos and compares it with Lamport et al.'s specification of Basic Paxos [25]. Section 4 presents the auxiliary predicates used throughout the proof (Section 4.1), the invariants proved (Sections 4.2, 4.3 and 4.4), and an overview of the main strategy used in the proof (Section 4.5). Section 5 describes the parts added to the specification of Multi-Paxos to support Preemption, and the changes enumerated above. Section 6 summarizes the results from our specification and proof. Section 7 discusses related work and concludes. The complete, cleaned up specification, invariants, and proof can be found in Appendices A, B, and C, respectively.

#### 2 Preliminaries

#### 2.1 Distributed consensus

A distributed system is a set of processes that process data locally and communicate with each other by sending and receiving messages. The processes may crash and may later recover, and the messages may be delayed indefinitely or lost. The basic consensus problem, called single-value consensus, is to ensure that at most a single value is chosen from among the values proposed by the processes. Formally, it is defined as

$$Safe_{basic} \triangleq \forall v_1, v_2 \in \mathcal{V} : Chosen(v_1) \wedge Chosen(v_2) \Rightarrow v_1 = v_2$$
 (1)

where  $\mathcal{V}$  is the set of possible proposed values, and *Chosen* is a predicate that given a value v evaluates to true iff v was chosen by the algorithm. The specification of *Chosen* is part of the algorithm.

The more general consensus problem, called multi-value consensus, is to choose a sequence of values, instead of a single value. Here we have

$$Safe_{multi} \triangleq \forall v_1, v_2 \in \mathcal{V}, s \in \mathcal{S} : Chosen(s, v_1) \land Chosen(s, v_2) \Rightarrow v_1 = v_2$$
 (2)

where V is as above, S is a set of *slots* used to index the sequence of chosen values, and *Chosen* is a predicate that given a slot s and a value v evaluates to true iff for slot s, value v was chosen by the algorithm.

#### 2.2 Basic Paxos and Multi-Paxos

Paxos solves the problem of consensus. Two main roles of the algorithm are performed by two kinds of processes:

- $\bullet$   $\mathcal{P}$ , the set of proposers that propose values that can be chosen.
- $\bullet$   $\mathcal{A}$ , the set of acceptors that vote for proposed values. A value is chosen when there are enough votes for it.

A set  $\mathcal{Q}$  of subsets of the acceptors, that is,  $\mathcal{Q} \subseteq 2^{\mathcal{A}}$ , is used as a quorum system. It must satisfy the following properties:

- Q is a set cover for A, that is,  $\bigcup_{Q \in Q} Q = A$ .
- Any two quorums overlap, that is,  $\forall Q_1, Q_2 \in \mathcal{Q} : Q_1 \cap Q_2 \neq \emptyset$ .

The most commonly used quorum system Q takes any majority of acceptors as an element in Q.

Basic Paxos solves the problem of single-value consensus. It defines predicate *Chosen* as

$$Chosen(v) \triangleq \exists Q \in Q : \forall a \in Q : \exists b \in \mathcal{B} : sent("2b", a, b, v)$$
(3)

where  $\mathcal{B}$  is the set of proposal numbers, also called ballot numbers, which is any set that can be totally ordered. sent("2b", a, b, v) means that a message of type 2b with ballot number b and value v was sent by acceptor a. An acceptor votes by sending such a message.

Multi-Paxos solves the problem of multi-value consensus. It extends predicate *Chosen* to decide a value for each slot s in S:

$$Chosen(s, v) \triangleq \exists Q \in \mathcal{Q} : \forall a \in Q : \exists b \in \mathcal{B} : sent("2b", a, b, s, v)$$

$$(4)$$

To satisfy the Safe property, S can be any set. In practice, S is usually the natural numbers. Figure 1 shows Lamport's description of Basic Paxos [20]. It uses any majority of acceptors as a quorum. Following Lamport et al. [25], in the specifications presented in this paper, the prepare requests and responses have been renamed to 1a and 1b messages, respectively, the accept requests and responses have been renamed to 2a and 2b messages, respectively, and the number n is renamed to b and bal.

Multi-Paxos can be built from Basic Paxos by carefully adding slots. In Basic Paxos, acceptors cache the value they have accepted with the highest ballot number. In Multi-Paxos, we have a sequence of these values indexed by slot.

- 1. Phase 1a is unchanged.
- 2. In Phase 1b, the acceptors now respond with a set of triples in  $\mathcal{B} \times \mathcal{S} \times \mathcal{V}$  as opposed to just one ballot in  $\mathcal{B}$  and one value in  $\mathcal{V}$ .
- 3. In Phase 2a, the proposers now propose a set of pairs in  $\mathcal{S} \times \mathcal{V}$  instead of just one value in  $\mathcal{V}$ . Similar to Basic Paxos, the proposer executes Phase 2a once it has received a set of responses for its 1a message from a quorum of acceptors and picks the value with highest ballot number. But this is now performed separately for each slot in the set of the triples received in the responses.
- 4. In Phase 2b, the acceptors now respond with a set of pairs in  $\mathcal{S} \times \mathcal{V}$  as opposed to just one value in  $\mathcal{V}$ .

Putting the actions of the proposer and acceptor together, we see that the algorithm operates in the following two phases.

- **Phase 1.** (a) A proposer selects a proposal number n and sends a *prepare* request with number n to a majority of acceptors.
- (b) If an acceptor receives a *prepare* request with number n greater than that of any prepare request to which it has already responded, then it responds to the request with a promise not to accept any more proposals numbered less than n and with the highest-numbered proposal (if any) that it has accepted.
- **Phase 2.** (a) If the proposer receives a response to its *prepare* requests (numbered n) from a majority of acceptors, then it sends an accept request to each of those acceptors for a proposal numbered n with a value v, where v is the value of the highest-numbered proposal among the responses, or is any value if the responses reported no proposals.
- (b) If an acceptor receives an accept request for a proposal numbered n, it accepts the proposal unless it has already responded to a prepare request having a number greater than n.

A proposer can make multiple proposals, so long as it follows the algorithm for each one. ... It is probably a good idea to abandon a proposal if some proposer has begun trying to issue a high-numbered one. Therefore, if an acceptor ignores a *prepare* or *accept* request because it has already received a *prepare* request with a higher number, then it should probably inform the propose, who should then abandon its proposal. This is a performance optimization that does not affect correctness.

To learn that a value has been chosen, a learner must find out that a proposal has been accepted by a majority of acceptors. The obvious algorithm is to have each acceptor, whenever it accepts a proposal, respond to all learners, sending them the proposal.

Figure 1: Lamport's description of Basic Paxos in English [20].

5. Learning, as described in the last part of Figure 1, is changed to consider different slots separately—a process learns that a value is chosen for a slot if a quorum of acceptors accepted it for that slot in 2b messages.

#### 2.3 TLA $^+$

The specifications presented in this paper are written in  $TLA^+$ , an extension of the Temporal Logic of Actions (TLA) [28], a logic for specifying concurrent and distributed programs and reasoning about their properties. In TLA, a *state* is an instantiation of the variables of the program to values. An *action* is a relation between a current state and a new state, specifying the effect of executing a sequence of instructions. For example, the instruction x := x + 1 is represented in TLA and TLA<sup>+</sup> by the action x' = x + 1. An action is represented by a formula over unprimed and primed variables where unprimed variables refer to the values of the variables in the current state and primed variables refer to the values of the variables in the new state.

A program is specified by its actions and initial states. Formally, a program is specified

as  $Spec \triangleq Init \wedge \Box [Next]_{vars}$  where Init is a predicate that holds for initial states of the program, Next is a disjunction of all the actions of the program, and vars is the tuple of all the variables. The expression  $[Next]_{vars}$  is true if either Next is true, implying some action is true and therefore executed, or vars stutters, that is, the values of the variables are same in the current and new states.  $\Box$  is the temporal operator always.

As a simple example, consider this specification of a clock based on Lamport's logical clock [27] but on a shared memory system:

VARIABLE 
$$c$$
 $Max(S) \triangleq CHOOSE \ e \in S : \forall f \in S : e \geq f$ 
 $Init \triangleq c = [p \in \{0,1\} \mapsto 0]$ 
 $LocalEvent(p) \triangleq c' = [c \ EXCEPT \ ![p] = c[p] + 1]$ 
 $ReceiveEvent(p) \triangleq c' = [c \ EXCEPT \ ![p] = Max(\{c[p], c[1-p]\}) + 1]$ 
 $Next \triangleq \exists \ p \in \{0,1\} : LocalEvent(p) \lor ReceiveEvent(p)$ 
 $Spec \triangleq Init \land \Box[Next]_{\langle c \rangle}$ 
 $(5)$ 

The system has two processes numbered 0 and 1. Variable c stores their current clock values as a function from process numbers to clock values. Both processes start with clock value 0, as specified in Init. LocalEvent(p) specifies that process p has executed some local action and therefore increments its clock value. The expression c' = [c EXCEPT ! [p] = c[p] + 1] means that function c' is the same as function c except that c'[p] is c[p] + 1. ReceiveEvent(p) specifies that process p updates its clock value to 1 greater than the higher of its and the other process' clock value. We define operator Max to obtain the highest of a set of values. Choose denotes Hilbert's  $\epsilon$  operator that returns some nondeterministically chosen term satisfying the body of the Choose expression if it exists, otherwise an error is raised.

#### 2.4 TLAPS

TLA<sup>+</sup> Proof System (TLAPS) is a tool that mechanically checks proofs of properties of systems specified in TLA<sup>+</sup>. Proofs are written in a hierarchical style [26], and are transformed to individual proof obligations that are sent to backend theorem provers. The primary backend provers are Isabelle and Zenon, with the SMT solvers CVC3, Z3, veriT, and Yices as backups. Temporal formulas are proved using LS4, a PTL (Propsitional Temporal Logic) prover. Users can specify which prover they want to use by using its name and can specify the timeout for each obligation separately.

As an example, we present the proof of a simple type invariant about the clock specification in (5) — It is always the case that  $c \in [\{0,1\} \to \mathbb{N}]$ :

$$TypeOK \triangleq c \in [\{0,1\} \to \mathbb{N}]$$

$$THEOREM \ Inv \triangleq Spec \Rightarrow \Box (TypeOK)$$

$$\langle 1 \rangle. \ USE \ DEF \ TypeOK$$

$$\langle 1 \rangle 1. \ Init \Rightarrow TypeOK \ BY \ DEF \ Init$$

$$\langle 1 \rangle 2. \ TypeOK \land [Next]_{\langle c \rangle} \Rightarrow TypeOK' \ BY \ DEF \ Next, LocalEvent, ReceiveEvent$$

$$\langle 1 \rangle. \ QED \ BY \ \langle 1 \rangle 1, \langle 1 \rangle 2, \ PTL \ DEF \ Spec$$

$$(6)$$

The proof of theorem Inv is written in a hierarchical fashion. It is proved by two steps, named  $\langle 1 \rangle 1$  and  $\langle 1 \rangle 2$ , and RuleINV1 by Lamport [28]. Proof steps in TLAPS are typically written as:

$$\langle x \rangle y$$
. Assertion BY  $e_1, \dots, e_m$  DEF  $d_1, \dots, d_n$  (7)

Basic Paxos	Multi-Paxos
$Phase1a(b \in \mathcal{B}) \triangleq$	$Phase1a(p \in \mathcal{P}) \triangleq$
$\wedge \nexists m \in msgs : \wedge m.type = "1a"$	$\wedge \nexists \ m \in msgs : \wedge m.type = \text{``1a''}$
$\land m.bal = b$	$\wedge m.bal = pBal[p]$
$\land Send([type \mapsto "1a",$	$\land Send([type \mapsto "1a", from \mapsto p,$
$bal \mapsto b$ )	$bal \mapsto pBal[p]])$
$\land$ UNCHANGED $\langle maxVBal, maxBal, maxVal \rangle$	$\land$ UNCHANGED $\langle pBal, aBal, aVoted \rangle$

Figure 2: Phase 1a of Basic Paxos and Multi-Paxos

which states that step number  $\langle x \rangle y$  proves Assertion by assuming  $e_1, \ldots, e_m$ , and expanding the definitions of  $d_1, \ldots, d_n$ . For example, step  $\langle 1 \rangle 1$  proves  $Init \Rightarrow TypeOK$  by expanding the definition of Init. The QED step for  $\langle 1 \rangle$  requires us to invoke a PTL prover because Inv is a temporal formula.

## 3 Specification of Multi-Paxos

We give a formal specification of Multi-Paxos by minimally extending that of Basic Paxos by Lamport et al [25].

Variables. The specification of Multi-Paxos has four global variables.

msgs: the set of messages that have been sent. Processes read from or add to this set. This is the same as in the specification of Basic Paxos except that the contents of messages are more complex.

pBal: per proposer, the current ballot number of the proposer. This is not in the specification of Basic Paxos; it is added to support preemption.

aBal: per acceptor, the highest ballot number seen by the acceptor. This is named maxBal in the specification of Basic Paxos.

a Voted: per acceptor, a set of triples in  $\mathcal{B} \times \mathcal{S} \times \mathcal{V}$  voted by the acceptor. For each slot only the triple with the highest ballot number is stored. This contrasts with two numbers per acceptor, in two variables, maxVBal and maxVal, in the specification of Basic Paxos.

**Algorithm steps.** The algorithm consists of repeatedly executing two phases.

**Phase 1a.** Figure 2 shows the specifications of Phase 1a for Basic Paxos and Multi-Paxos, which are in essence the same. Parameter ballot number b in Basic Paxos is replaced with proposer p executing this phase in Multi-Paxos, to allow extensions such as Preemption that need to know the proposer of a ballot number; uses of b are changed to pBal[p]; and  $from \mapsto p$  is added in Send. Send is a macro that adds its argument to msgs, i.e.,  $Send(m) \triangleq msgs' = msgs \cup \{m\}$ . In this specification, 1a messages do not have a receiver, making them accessible to all processes. However, this is not required. For safety, it is enough to send this message to any subset of A, even  $\emptyset$ . For liveness, the receiving set should contain at least one quorum.

```
Basic Paxos
                                                               Multi-Paxos
Phase1b(a \in \mathcal{A})
                                                               Phase1b(a \in \mathcal{A})
\exists m \in msqs:
                                                               \exists m \in msqs:
   \land m.type = "1a"
                                                                  \land m.type = "1a"
  \land m.bal > maxBal[a]
                                                                  \land m.bal > aBal[a]
  \land Send([type \mapsto "1b",
                                                                  \land Send([type \mapsto "1b",
                                                                     from \mapsto a,
      acc \mapsto a,
      bal \mapsto m.bal,
                                                                     bal \mapsto m.bal,
     maxVBal \mapsto maxVBal[a],
                                                                     voted \mapsto aVoted[a]
      maxVal \mapsto maxVal[a]
   \land maxBal' = [maxBal \ EXCEPT \ ![a] = m.bal]
                                                                  \wedge aBal' = [aBal \ EXCEPT \ ![a] = m.bal]
   \land UNCHANGED \langle maxVBal, maxVal \rangle
                                                                  \land UNCHANGED \langle pBal, aVoted \rangle
```

Figure 3: Phase 1b of Basic Paxos and Multi-Paxos

**Phase 1b.** Figure 3 shows the specifications of Phase 1b. Parameter acceptor a executes this phase. The only key difference between the specifications is the set aVoted[a] of triples in Send of Multi-Paxos vs. the two numbers maxVBal[a] and maxVal[a] in Basic Paxos.

**Phase 2a.** Figure 4 shows Phase 2a. The key difference is, in Send, the bloating of a single value v in  $\mathcal{V}$  in Basic Paxos to a set of pairs in  $\mathcal{S} \times \mathcal{V}$  given by PropSV in Multi-Paxos. A proposal is a  $\langle s, v \rangle$  pair. The operation of finding the value with the highest ballot in Basic Paxos is performed for each slot by MaxSV in Multi-Paxos; MaxSV takes a set T of triples in  $\mathcal{B} \times \mathcal{S} \times \mathcal{V}$  and returns a set of pairs in  $\mathcal{S} \times \mathcal{V}$ . NewSV generates a set of pairs in  $\mathcal{S} \times \mathcal{V}$  where values are proposed for slots not in MaxSV. This is significantly more sophisticated than running Basic Paxos for each slot, because the ballots are shared and changing for all slots, and slots are paired with values dynamically where slots that failed to reach consensus values earlier are also detected and reused.

**Phase 2b.** Figure 5 shows Phase 2b. In Basic Paxos, the acceptor replies with the value received in the 2a message whereas in Multi-Paxos, it replies with a set of pairs in  $\mathcal{S} \times \mathcal{V}$  received in the 2a message. Also, in Basic Paxos, the acceptor updates its voted pair maxVBal[a] and maxVal[a] upon receipt of a 2a message of the highest ballot; in Multi-Paxos, this is performed for each slot. The acceptor updates aVoted to have all proposals in the received 2a message and all previous values in aVoted for slots not mentioned in that message.

Complete algorithm specification. To complete the algorithm specification, we define vars, Init, Next, and Spec, typical TLA<sup>+</sup> macro names for the set of variables, the initial state, possible actions leading to the next state, and the system specification, respectively:

```
vars \triangleq \langle msgs, pBal, aBal, aVoted \rangle
Init \triangleq msgs = \emptyset \land pBal = [p \in \mathcal{P} \mapsto 0] \land aBal = [a \in \mathcal{A} \mapsto -1] \land aVoted = [a \in \mathcal{A} \mapsto \emptyset]
Next \triangleq (\exists p \in \mathcal{P} : Phase1a(p) \lor Phase2a(p)) \lor (\exists a \in \mathcal{A} : Phase1b(a) \lor Phase2b(a))
Spec \triangleq Init \land \Box [Next]_{vars}
(8)
```

```
Basic Paxos
                                                                       Multi-Paxos
Phase2a(b \in \mathcal{B}) \triangleq
                                                                       Phase2a(p \in \mathcal{P}) \triangleq
\wedge \nexists m \in msqs : \wedge m.type = "2a"
                                                                       \wedge \nexists m \in msqs : \wedge m.type = "2a"
                  \wedge m.bal = b
                                                                                         \wedge m.bal = pBal[p]
\land \exists v \in \mathcal{V}:
   \land \exists Q \in \mathcal{Q}, S \subseteq \{m \in msgs : m.type = \text{``1b''} \land
                                                                       \land \exists Q \in \mathcal{Q}, S \subseteq \{m \in msgs : m.type = \text{``1b''} \land
      m.bal = b:
                                                                          m.bal = pBal[p]:
      \land \forall \ a \in Q : \exists \ m \in S : m.acc = a
                                                                          \land \forall \ a \in Q : \exists \ m \in S : m.from = a
      \land \lor \forall \, m \in S \, : \, m.maxVBal = -1
                                                                          \land Send([type \mapsto "2a",
          \forall \exists c \in 0..(b-1):
                                                                             from \mapsto p,
             \land \forall \ m \in S : m.maxVBal \le c
                                                                              bal \mapsto pBal[p],
                                                                              propSV \mapsto PropSV(UNION)
             \land \exists m \in S : (m.maxVBal = c)
                                                                                 \{m.voted : m \in S\}\}
                    \wedge m.maxVal = v
   \land Send([type \mapsto "2a", bal \mapsto b, val \mapsto v])
\land UNCHANGED \langle maxBal, maxVBal, maxVal \rangle
                                                                       \land UNCHANGED \langle pBal, aBal, aVoted \rangle
                                                                       MaxBSV(T) \triangleq \{t \in T : \forall t2 \in T :
                                                                          t2.slot = t.slot \Rightarrow t2.bal < t.bal
                                                                       MaxSV(T) \triangleq \{[slot \mapsto t.slot,
                                                                          val \mapsto t.val: t \in MaxBSV(T)
                                                                       UnusedS(T) \triangleq \{s \in S : \nexists t \in T : t.slot = s\}
                                                                       NewSV(T) \triangleq CHOOSE D \subset [slot:]
                                                                          UnusedS(T), val : \mathcal{V}] : \forall d1, d2 \in D:
                                                                              d1.slot = d2.slot \Rightarrow d1 = d2 \land D \neq \emptyset
                                                                       PropSV(T) \triangleq MaxSV(T) \cup NewSV(T)
```

Figure 4: Phase 2a of Basic Paxos and Multi-Paxos

```
Basic Paxos
                                                                Multi-Paxos
Phase2b(a \in \mathcal{A}) \triangleq
                                                                Phase2b(a \in \mathcal{A}) \triangleq
\exists m \in msqs:
                                                                \exists m \in msqs:
   \wedge m.type = "2a"
                                                                   \land m.type = "2a"
   \land m.bal \ge maxBal[a]
                                                                   \land m.bal \geq aBal[a]
  \land Send([type \mapsto "2b",
                                                                   \land Send([type \mapsto "2b",
      acc \mapsto a,
                                                                      from \mapsto a,
                                                                      bal \mapsto m.bal,
      bal \mapsto m.bal,
                                                                      propSV \mapsto m.propSV
      val \mapsto m.val
  \land maxBal' = [maxBal \ EXCEPT \ ![a] = m.bal]
                                                                   \wedge aBal' = [aBal \ EXCEPT \ ![a] = m.bal]
   \wedge maxVBal' =
                                                                   \wedge a Voted' = [a Voted EXCEPT ! [a] =
      [maxVBal \ EXCEPT \ ![a] = m.bal]
                                                                      \cup \{[bal \mapsto m.bal, slot \mapsto d.slot,
                                                                         val \mapsto d.val]: d \in m.propSV}]
  \wedge maxVal' =
      [maxVal \text{ EXCEPT } ![a] = m.val]
                                                                      \cup \{e \in aVoted[a] :
                                                                         \nexists r \in m.propSV : e.slot = r.slot \}
                                                                   \landUNCHANGED \langle pBal \rangle
```

Figure 5: Phase 2b of Basic Paxos and Multi-Paxos

The complete specification of Multi-Paxos with Preemption is given in Appendix A. Preemp-

tion is discussed in Section 5. We only provide specification of Multi-Paxos with Preemption because it is an extension of Multi-Paxos, thus also giving specification of Multi-Paxos would be redundant.

# 4 Verification of Multi-Paxos and Proof Strategy

We first define the auxiliary predicates and invariants used, by extending those for the proof of Basic Paxos with slots, and then describe our proof strategy, which proves *Safe* of Multi-Paxos.

#### 4.1 Auxiliary predicates

These predicates are used throughout the proof. Chosen(s, v) is true if there exists some ballot b such that ChosenIn(b, s, v) holds.

$$Chosen(s \in \mathcal{S}, v \in \mathcal{V}) \triangleq \exists b \in \mathcal{B} : ChosenIn(b, s, v)$$
(9)

ChosenIn(b, s, v) is true if there exists some quorum of acceptors such that for each acceptor in the quorum, VotedForIn(a, b, s, v) holds.

$$ChosenIn(b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V}) \triangleq \exists Q \in \mathcal{Q} : \forall a \in Q : VotedForIn(a, b, s, v)$$
 (10)

VotedForIn(a,b,s,v) is true if acceptor a has voted value v for slot s in ballot b. This is realized in the algorithm by sending a 2b message with ballot b and pair  $\langle s,v \rangle$  in the message's proposals. Putting everything together, the algorithm chooses value v for slot s if there exist some quorum Q and ballot b such that every acceptor in Q has voted value v for slot s in ballot b:

$$VotedForIn(a \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V}) \triangleq \exists m \in msgs : \\ m.type = \text{``2b''} \land m.from = a \land m.bal = b \land \exists d \in m.propSV : d.slot = s \land d.val = v$$

$$(11)$$

Predicate  $SafeAt(b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V})$  means that no value except perhaps v has been or will be chosen in any ballot lower than b for slot s. This is realized by asserting that for each ballot b2 < b, there exists a quorum of acceptors such that for each acceptor in the quorum, either VotedForIn(a, b2, s, v) holds or WontVoteIn(a, b2, s) holds.

$$SafeAt(b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V}) \triangleq \forall b2 \in 0..(b-1) : \exists Q \in \mathcal{Q} : \\ \forall a \in Q : VotedForIn(a, v, b2, s) \lor WontVoteIn(a, b2, s)$$

$$(12)$$

WontVoteIn(a, b, s) holds if acceptor a has seen a higher ballot than b, and did not and will not vote any value in b for slot s

$$WontVoteIn(a \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}) \triangleq aBal[a] > b \land \forall v \in \mathcal{V} : \neg VotedForIn(a, b, s, v)$$
 (13)

Predicate  $MaxBalInSlot(S \subseteq [bal : \mathcal{B}, slot : \mathcal{S}], s \in \mathcal{S})$  selects among set of elements in S with slot s, the highest ballot, or -1 if no element has slot s.

$$Max(T) \triangleq \text{CHOOSE } e \in T : \forall f \in T : e \geq f$$
 $MaxBalInSlot(T \subseteq [bal : \mathcal{B}, slot : \mathcal{S}], s \in \mathcal{S}) \triangleq$ 

LET  $E \triangleq \{e \in T : e.slot = s\}$ 

IN IF  $E = \emptyset$  THEN  $-1$  ELSE  $Max(\{e.bal : e \in E\})$ 

To prove the Safe property for the algorithm, we prove two properties:

- 1. Lemma VotedInv. If any acceptor votes any triple  $\langle b, s, v \rangle$ , then the predicate SafeAt(b, s, v) holds. That is,  $\forall a \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : VotedForIn(a, b, s, v) \Rightarrow SafeAt(b, s, v)$ .
- 2. Lemma VotedOnce. If acceptor a1 votes triple  $\langle b, s, v1 \rangle$  and acceptor a2 votes triple  $\langle b, s, v2 \rangle$ , then v1 = v2. That is,  $VotedForIn(a1, b, s, v1) \wedge VotedForIn(a2, b, s, v2) \Rightarrow v1 = v2$ .

In fact, for other consensus algorithms, either Paxos extensions like Fast Paxos [22] and Byzantine Paxos [23], or Paxos alternatives like Viewstamped Replication [30] and Raft [35], safety can also be proven by asserting these two properties.

To prove these properties, we write invariants about the algorithm. We, following Lamport et al. [25], use three kinds of invariants: type invariants, process invariants, and message invariants. These are sufficient because a distributed algorithm handles two kinds of data: messages communicated and process local data. Correspondingly, we obtain message invariants for the messages passed between processes and process invariants over the local data that these processes maintain. Type invariants ensure that the specification always uses data with correct types.

#### 4.2 Type invariants

Type invariants are captured by TypeOK. They specify the sets of values that the variables of the system can hold. For example,  $pBal \in [\mathcal{P} \to \mathcal{B}]$  specifies the set of values pBal can hold: for any proposer p in  $\mathcal{P}$ , pBal[p] is in  $\mathcal{B}$ .

```
\begin{aligned} & Messages \; \triangleq \\ & \cup [type : \{\text{``1a''}\}, bal : \mathcal{B}, from : \mathcal{P}] \\ & \cup [type : \{\text{``1b''}\}, bal : \mathcal{B}, voted : \text{SUBSET} [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}], from : \mathcal{A}] \\ & \cup [type : \{\text{``2a''}\}, bal : \mathcal{B}, propSV : \text{SUBSET} [slot : \mathcal{S}, val : \mathcal{V}], from : \mathcal{P}] \\ & \cup [type : \{\text{``2b''}\}, bal : \mathcal{B}, from : \mathcal{A}, propSV : \text{SUBSET} [slot : \mathcal{S}, val : \mathcal{V}]] \\ & \cup [type : \{\text{``preempt''}\}, bal : \mathcal{B}, to : \mathcal{P}, maxBal : \mathcal{B}] \end{aligned} 
& TypeOK \; \triangleq \\ & \land msgs \subseteq Messages \land pBal \in [\mathcal{P} \rightarrow \mathcal{B}] \land aBal \in [\mathcal{A} \rightarrow \mathcal{B} \cup \{-1\}] \\ & \land aVoted \in [\mathcal{A} \rightarrow \text{SUBSET} [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}]] \end{aligned}
```

## 4.3 Invariants about acceptors

The following predicate specifies invariants about acceptor processes. For each acceptor a, the first conjunct establishes the initial condition. The second conjunct says that a has voted each triple in aVoted[a] and aBal[a] is higher than or equal to the ballot number of each triple in aVoted[a]. The third conjunct states that if acceptor a has voted any value v in ballot b for slot s, then there is some triple t in aVoted[a] such that  $t.bal \geq b$  and t.slot = s. The last conjunct says that acceptor a does not vote for a value in any ballot higher than the highest

it has seen per slot.

```
AccInv \triangleq \forall \ a \in \mathcal{A} :
\land (aBal[a] = -1) \Rightarrow (aVoted[a] = \emptyset)
\land \forall \ t \in aVoted[a] : aBal[a] \geq t.bal \land VotedForIn(a, t.bal, t.slot, t.val)
\land \forall \ b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : VotedForIn(a, b, s, v) \Rightarrow \exists \ t \in aVoted[a] : t.bal \geq b \land t.slot = s
\land \forall \ b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b > MaxBalInSlot(aVoted[a], s) \Rightarrow \neg VotedForIn(a, b, s, v)
(16)
```

#### 4.4 Invariants about messages

$$MsgInv1b(m) \triangleq \\ \land m.bal \leq aBal[m.from] \\ \land \forall t \in m.voted : VotedForIn(m.from, t.bal, t.slot, t.val) \\ \land \forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b \in MaxBalInSlot(m.voted, s) + 1..m.bal - 1 \\ \Rightarrow \neg VotedForIn(m.from, b, s, v)$$
 (17)

The following invariant is for a 2a message m. The first conjunct establishes safety for each proposal d in m. The second conjunct says that two proposals in m with the same slot must be the same proposal. That is, for each slot, there is at most one proposal in m. The third conjunct says that there is at most one 2a message for each ballot.

$$MsgInv2a(m) \triangleq \\ \wedge \forall d \in m.propSV : SafeAt(m.bal, d.slot, d.val) \\ \wedge \forall d1, d2 \in m.propSV : d1.slot = d2.slot \Rightarrow d1 = d2 \\ \wedge \forall m2 \in msqs : (m2.type = "2a") \wedge (m2.bal = m.bal) \Rightarrow m2 = m$$

$$(18)$$

The following invariant is for a 2b message m sent by acceptor m.from. The first conjunct states that there exists a 2a message with the same ballot and same set of proposals as m. The second conjunct asserts that the ballot of m is no higher than the highest ballot seen by the sender of m.

```
MsgInv2b(m) \triangleq 
 \land \exists \ m2 \in msgs : m2.type = "2a" \land m2.bal = m.bal \land m2.propSV = m.propSV  (19)
 \land m.bal \leq aBal[m.from]
```

The complete message invariant is the conjunction of MsgInv1b, MsgInv2a, and MsgInv2b:

$$MsgInv \triangleq \forall m \in msgs : \land (m.type = "1b") \Rightarrow MsgInv1b(m)$$
  
  $\land (m.type = "2a") \Rightarrow MsgInv2a(m)$   
  $\land (m.type = "2b") \Rightarrow MsgInv2b(m)$  (20)

The complete invariants, auxiliary operators, and the safety property to be proved can be found in Appendix B.

#### 4.5 Proof strategy

The main theorem to prove is Safety as defined in Equation (21). For this, we define Inv and first prove  $Inv \Rightarrow Safe$ . Then, we prove  $Spec \Rightarrow \Box Inv$  which by temporal logic, concludes  $Spec \Rightarrow \Box Safe$ . Note that property Safety is called Consistent, and invariant Safe is called Consistency by Lamport et al [25].

$$Inv \triangleq TypeOK \land AccInv \land MsgInv$$

$$Safe \triangleq \forall v1, v2 \in \mathcal{V}, s \in \mathcal{S} : Chosen(v1, s) \land Chosen(v2, s) \Rightarrow v1 = v2$$

$$\mathsf{THEOREM} \ Safety \triangleq Spec \Rightarrow \Box Safe$$

$$(21)$$

The proof is developed following a standard hierarchical structure and uses proof by induction and contradiction. To prove  $Spec \Rightarrow \Box Inv$ , we employ a systematic proof strategy that works well for algorithms described in the event driven paradigm. Distributed algorithms are prime examples of this paradigm as they can be specified in blocks of code that are triggered upon receiving a certain set of messages and finish by sending a set of messages. We demonstrate the strategy for some invariants in Inv.

First, consider invariant TypeOK. The goal to prove is  $Spec \Rightarrow \Box TypeOK$ . Recall  $Spec \triangleq Init \land \Box [Next]_{vars}$ :

- The induction basis,  $Init \Rightarrow TypeOK$ , is trivial in this case because the set of sent messages is empty. TLAPS handles it automatically. This should be the case for almost all distributed algorithms.
- The induction step is to prove  $TypeOK \wedge [Next]_{vars} \Rightarrow TypeOK'$ , where the left side is the induction hypothesis, and right side is the goal to be proved.  $[Next]_{vars}$  is a disjunction of phases, as for any distributed algorithm, and TypeOK' is a conjunction of smaller invariants, as for many invariants.

Now, the hypothesis of the induction step can be stripped down to each disjunct of Next separately, and each smaller goal needs to be proved from all smaller disjuncts. This process is mechanical, and TLAPS provides a feature for precisely this expansion into smaller proof obligations. This breakdown is the first step in our proof strategy. For TypeOK, this expands to 4 smaller assertions; with 4 phases in Next, we obtain 16 small proof obligations.

Proving complex invariants using invariance lemmas and increments. AccInv and MsgInv are more involved. We proceed as we did for TypeOK and create a proof tree, each branch of which aims to prove an invariant for a disjunct in Next. To explain the rest of our strategy, we show a complex combination: MsgInv and Phase 1b. Equation (22) shows the skeleton of the proof. The goal is to prove  $\langle 4 \rangle 2$ , which states that MsgInv' holds if acceptor a executes Phase 1b.  $\langle 6 \rangle 1$  proves MsgInv1b(m)' and  $\langle 6 \rangle 2$  proves MsgInv2a(m)' and MsgInv2b(m)'.

Phase 1b sends a 1b message. Intuitively,  $\langle 6 \rangle$ 2 should be easy because its goals are not invariants about 1b messages. However, this is not the case because of predicate SafeAt, which is used in MsgInv2a. At this point the prover needs an  $invariance\ lemma$ .

Invariance lemma. An invariance lemma is a lemma that asserts that a predicate continues to hold (or not hold) as the system goes from one state to the next in a single step. This happens when a step does not affect the part of system state asserted by the predicates, and requires more complex proofs otherwise. For example, the invariance lemma for SafeAt states that SafeAt continues to hold for any disjunct in Next, which includes Phase 1b. The

characteristic property of such lemmas is their reuse. In our proof of Multi-Paxos, we defined 5 invariance lemmas which are used in 27 places.

Lastly, we need to prove  $\langle 6 \rangle 1$ . Since  $\langle 6 \rangle 1$  is about 1b messages, and Phase 1b generates such a message, the proof is more complicated, and the prover needs more manual intervention. Here we split the set of messages in the new state into two sets:  $\langle 7 \rangle 1$  for the old messages, and  $\langle 7 \rangle 2$  for the *increment*, the new message sent in this step. For the old messages, we use invariance lemmas. The most challenging case is for the increment.

Increment. An increment is a new message sent in a phase or generally a new element in a set. In the example of Equation (22), the increment is m1 and we focus on the cause of the increment—available here in the definition of Phase 1b—and prove each conjunct of the goal MsgInv1b(m1)' separately in  $\langle 5 \rangle 2$ , 4, 12. The prover proves  $\langle 5 \rangle 2$  by just the definition of Phase 1b. For  $\langle 5 \rangle 4$ , along with the definition of Phase 1b, the prover also needs invariance lemma for VotedForIn for Phase1b.  $\langle 5 \rangle 12$  requires, along with the definition of Phase 1b and invariance lemmas, case-specific manual intervention. Specifically, we helped the prover understand the change in limits of the set MaxBalInSlot(m1.voted, s) + 1..m1.bal - 1. The proof extended one more level for case  $\langle 7 \rangle 2$ .

```
 \langle 4 \rangle \text{2.ASSUME NEW } a \in \mathcal{A}, Phase1b(a) \text{ PROVE } MsgInv' \\ \langle 5 \rangle. \text{DEFINE } m1 \triangleq [type \mapsto \text{``1b''}, from \mapsto a, bal \mapsto m.bal, voted \mapsto aVoted[a]] \\ \langle 5 \rangle 2.(m1.bal \leq aBal[m1.from])' \dots \\ \langle 5 \rangle 4.(\forall r \in m1.voted : VotedForIn(m1.from, r.bal, r.slot, r.val))' \dots \\ \langle 5 \rangle 12.(\forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b \in MaxBalInSlot(m1.voted, s) + 1..m1.bal - 1 \Rightarrow \\ \neg VotedForIn(m1.from, b, s, v))' \dots \\ \vdots \\ \langle 6 \rangle \text{ASSUME NEW } m2 \in msgs' \text{ PROVE } MsgInv' \\ \langle 6 \rangle 1.(m2.type = \text{``1b''} \Rightarrow MsgInv1b(m2))' \\ \langle 7 \rangle 1.\text{CASE } m2 \in msgs \dots \\ \langle 7 \rangle 2.\text{CASE } m2 \in msgs' \setminus msgs \dots \\ \langle 6 \rangle 2.((m2.type = \text{``2a''} \Rightarrow MsgInv2a(m2)) \wedge (m2.type = \text{``2b''} \Rightarrow MsgInv2b(m2)))' \dots \end{aligned}
```

Induction for properties over sets, and ways of accessing elements of tuples. With the above strategy, we were still faced with certain assertions that were difficult to prove. One of the main difficulties lay in proving properties about tuples and sets of tuples for each of a set of processes in Multi-Paxos, as opposed to one or two values for each of a set of processes in Basic Paxos. It may appear that, in many places, this requires simply adding an extra parameter for the slot, but the proof became significantly more difficult: even in places where an explicit inductive proof is not needed, auxiliary facts had to be added to help TLAPS succeed or proceed sufficiently fast.

For example, adding slots to the proof of theorem Safety for Basic Paxos caused the prover to take about 90 seconds to check it. To aid the proof, we added  $\exists a \in \mathcal{A}$ :  $VotedForIn(a, b1, s, v1) \wedge VotedForIn(a, b1, s, v2)$  as an intermediate fact derivable from  $b1 = b2 \wedge ChosenIn(b1, s, v1) \wedge ChosenIn(b2, s, v2)$  (step  $\langle 3 \rangle 1$  of step  $\langle 2 \rangle 1$  in the proof of theorem Safety). Following this, the prover asserted the conclusion  $v_1 = v_2$  in a few milliseconds, a

10,000 time speedup.

Tuples have only a fixed number of components and therefore do not require separate inductive proofs, but they often turn out to be tricky and require special care in choosing the ways to access and test their elements, to sufficiently reduce TLAPS's proof-checking time and observe progress. For example, consider the definition of VotedForIn in Equation (11). Originally  $[slot \mapsto s, val \mapsto v] \in m.propSV$  was written for the last conjunct with existential quantification, because it was natural, but it had to be changed to  $\exists d \in m.propSV : d.slot = s \land d.val = v$ , because the prover made more observable progress. With the original version, the proof did not carry through after 1 or 2 minutes. After the change, the proof proceeded quickly. One minute of waiting for such simple, small tests felt very long, making it uncertain whether the proof would carry through.

## 5 Multi-Paxos with Preemption, and overall proof

Preemption is described informally in Lamport's description of Basic Paxos in Figure 1, in the paragraph about abandoning a proposal number. Preemption has an acceptor reply to a proposer, in both Phases 1b and 2b, if the proposer's ballot is preempted i.e., the acceptor has seen a higher ballot than the one just received from the proposer. This reply informs the proposer of the highest ballot the acceptor has seen, and the proposer can abandon its lower ballot number.

```
NewBal(b2 \in \mathcal{B}) \triangleq \text{CHOOSE } b \in \mathcal{B}: b > b2 \land \nexists m \in msgs: m.type = "1a" \land m.bal = b
Preempt(p \in \mathcal{P}) \triangleq \exists m \in msgs :
                             \land m.type = "preempt"
                             \wedge m.to = p
                             \land m.bal > pBal[p]
                             \wedge pBal' = [pBal \ EXCEPT \ ![p] = NewBal(m.bal)]
                             \land UNCHANGED \langle msgs, aBal, aVoted \rangle
                                                              Phase 1b with Preemption
Phase 1b without Preemption
Phase1b(a \in \mathcal{A})
                                                              Phase1b(a \in \mathcal{A})
\exists m \in msgs:
                                                              \exists m \in msgs :
   \land m.type = "1a"
                                                                 \land m.type = "1a"
   \land m.bal > aBal[a]
                                                                 \wedgeIF m.bal > aBal[a] THEN
   \land Send([type \mapsto "1b",
                                                                    \land Send([type \mapsto "1b",
      from \mapsto a,
                                                                       from \mapsto a,
      bal \mapsto m.bal,
                                                                       bal \mapsto m.bal,
                                                                        voted \mapsto aVoted[a])
      voted \mapsto aVoted[a])
   \wedge aBal' = [aBal \text{ EXCEPT } ! [a] = m.bal]
                                                                    \wedge aBal' = [aBal \ EXCEPT \ ![a] = m.bal]
   \land UNCHANGED \langle pBal, aVoted \rangle
                                                                    \land UNCHANGED \langle pBal, aVoted \rangle
                                                                    ELSE
                                                                    \land Send([type \mapsto "preempt"],
                                                                        to \mapsto m.from, bal \mapsto aBal[a])
                                                                    \land UNCHANGED \langle pBal, aBal, aVoted \rangle
```

Figure 6: Extension of Multi-Paxos to Multi-Paxos with Preemption

To specify Preemption, each of Phases 1b and 2b adds a new case for when the acceptor receives a lower ballot than the highest ballot it has seen. We also define predicate *Preempt* 

that specifies how proposers update pBal upon receiving a **preempt** message. Figure 6 shows Phase 1b with and without the modifications to add Preemption. Modifications to Phase 2b are similar.

Preemption adds a new phase, Preempt, as a passive action in Next, modifies definitions of existing phases, and adds a new type of message. This new phase increases the width of the proof tree. Except for TypeOK, this new branch of the proof was proven by asserting invariance lemmas established earlier. The entire task of adding the new parts of specification and proof, except for the proof of TypeOK, took less than an hour.

#### 5.1 Remedial proof for TypeOK

The invariance of TypeOK, i.e.,  $\Box TypeOK$ , was automatically checked previously [2] by instructing TLAPS to use its PTL (Propositional Temporal Logic) backend prover. However, we later discovered that due to an undocumented bug in TLAPS [29], using PTL in this way causes incorrect obligations to be marked correct. This bug lets one prove that a primed formula is true if its unprimed form is an assumption. A primed formula is one with only primed variables and constants. For example, x' = 0 is primed but x' = x + 1 and x = 0 are not. The unprimed form of x' = 0 is x = 0. Thus, TLAPS incorrectly verifies TypeOK' under the assumption  $TypeOK \land [Next]_{vars}$ .

As a remedy, we here provide the missing proof of invariance of TypeOK. Using our proof strategy, writing this proof of invariance took a matter of minutes for Multi-Paxos. The remedial proof of TypeOK is longer by 39 lines (72%), increasing from 54 for the previous incomplete proof [2] to 93, and took 3 seconds (30%) more for TLAPS to check, from 10 to 13 seconds here.

For Multi-Paxos with Preemption, our strategy failed at first because the invariants were not strong enough. Preemption adds the conjunct  $pBal \in [\mathcal{P} \to \mathcal{B}]$  to TypeOK. Upon preemption, proposer p changes pBal[p] to ballot b such that (i) no 1a message has been sent yet with b as ballot, and (ii) b is higher than the ballot of the preempting message. To prove that  $\Box pBal \in [\mathcal{P} \to \mathcal{B}]$ , we need to establish that such a b indeed exists.

To this end, we strengthen our invariants by adding the fact that msgs is always a finite set, i.e., IsFiniteSet(msgs). We add this as a conjunct in TypeOK. Now, it can be proven that only a finite number of 12 messages exist. Thus, only a finite number of ballots can ever be used. Because  $\mathcal{B}$  is infinite, there will always be some ballot available to be chosen. We prove this constructively by providing the prover with a witness: a ballot that is 1 greater than the highest ballot in 12 messages. This remedial proof of TypeOK is 63 lines (84%) longer, increasing from 75 for the previous incomplete proof [2] to 138, and took 29 seconds (242%) more for TLAPS to check, from 12 to 41 seconds here.

Table 1 summarizes the results.

## 5.2 Overall proof improvement

To make the proof easier to read and understand, we made three main kinds of improvements. These improvements apply to the proofs for both Multi-Paxos and Multi-Paxos with Preemption, but the numbers presented are for the proof of Multi-Paxos with Preemption.

1. **Refined invariants.** MsgInv is now defined as a conjunction of three new predicates MsgInv1b, MsgInv2a, and MsgInv2b, in (17)-(19), which were not separate predicates in the proof for [2], even though the names were used in the text for ease of presentation. This introduction of intermediate predicates reduced the proof size by 75 lines, because,

- at 12 places, the new names are used in place of their expanded definitions. This also reduced the proof-checking time by about 2 seconds, because now only the portions of MsqInv that one needed for examination in the proof are instructed to be expanded.
- 2. Merged proof cases. Cases with similar proofs are merged into a single proof. For example, to prove that AccInv continues to hold when acceptor a executes Phase 1b with Preemption, there are three cases to consider for each acceptor a2 in A: (1) if a2 is not a, (2) if a2 is a, and it sends a preempt message in this action, and (3) if a2 is a, and it sends a 1b message in this action. The proofs for cases (1) and (2) are similar because in both cases, a2's state does not change. We found 5 instances in the proof described in [2] where such cases were handled separately, creating unnecessarily longer proof. Merging them resulted in an overall reduction of 27 lines of proof and slightly reduced (by less than 2 seconds) proof-checking times.
- 3. Removed unnecessary obligations. When TLAPS fails to prove an assertion, a proof must be manually written to be checked automatically by TLAPS. We discovered that parts of these manually written proofs described in [2] are unnecessary because they can be found automatically by TLAPS. About 25 instances of these were removed causing 110 lines of proof to be removed and the proof-checking times to reduce by around 20 seconds. An additional 22 instances were removed from the remedial proof for *TypeOK*, totalling to 61 lines of proof and reducing its checking time by 5 seconds.

		Multi-Paxo	OS	Multi- w/ Preemption			
Metric	Old [2]	Remedial	Improved	Old [2]	Remedial	Improved	
		Proof	Remedial		Proof	Remedial	
			Proof			Proof	
Proof size	54^^	93	54	75^^	138	77	
CPU check time (seconds)	10^^	13	4	12^^	41	12	
Elapsed check time (seconds)	20^^	22	8	21^^	35	25	

Table 1: Proof statistics for remedial proof of TypeOK and improvements on it. Proof size is measured as non-empty lines excluding comments.

## 5.3 Overall proof summary

The proof for Multi-Paxos with Preemption spans about 9 pages, and is summarized below:

- 1. Auxiliary predicates and invariants span about 1 page.
- 2. Helper lemmas and their proofs are about 1.5 pages. They are used to prove helpful properties of operators MaxBalInSlot, NewSV, and MaxSV. This also includes proofs for lemmas VotedInv and VotedOnce.
- 3. The invariance lemmas and their proofs are less than 1 page.
- 4. The proof of type invariant TypeOK is 1 page, using only a 1-level proof for each action.

<sup>^^</sup> indicates an incorrect number from the incomplete proof due to the TLAPS bug [29].

- 5. The proof of acceptor invariant AccInv is 1.5 pages, using a 5-level proof for action Phase2b because only Phase2b changes variable aVoted which affects AccInv. The proofs for other actions are only 1-level because of their independence from AccInv.
- 6. The proof of message invariant MsgInv is a little over 4 pages, using 4-level proofs for actions Phase1b and Phase2b taking about 1 page each, and a 6-level proof for Phase2a taking 2 pages, because MsgInv is over messages sent in these actions. The proof is long and complex in particular because each message is a record and its contents may contain sets of records.
- 7. The proof of theorem Safety using  $Spec \Rightarrow \Box Inv$  is less than half a page with a straightforward argument of  $Inv \Rightarrow Safe$ .

The complete TLAPS-checked proof is given in Appendix C.

# 6 Results of TLAPS-checked proofs

Table 2 summarizes the results from our specification and proof. Some mistakes in the counting logic of the result generating script caused incorrect numbers to be reported in [2] for specification and proof sizes and number of proofs by contradiction. These errors have been corrected here.

- The specification size grew by only 3 lines (6%), from 52 lines for Basic Paxos to 55 lines for Multi-Paxos; another 19 lines (35%) are added for Preemption.
- Overall, the proof size increased significantly by 477 lines (154%), from 310 for Basic Paxos to 787 for Multi-Paxos, due to the complex interaction between slots and ballots; only 44 more lines (6%) were added for Preemption, thanks to the reuse of all lemmas, especially invariance lemmas. As mentioned in Section 5.1, adding the remedial proof to the incomplete proof reported in [2] initially increased the proof size by 39 lines (4%) from 1010 to 1049 for Multi-Paxos and by 63 lines (6%) from 1054 to 1117 for Multi-Paxos with Preemption. However, the proof improvements decreased these numbers by 262 lines (25%) from 1049 to 787 for Multi-Paxos and by 286 lines (26%) from 1117 to 831 for Multi-Paxos with Preemption.
- The maximum level of proof tree nodes increased from 7 to 10 going from Basic Paxos to Multi-Paxos but remained 10 after adding Preemption; this contrast is even stronger for the maximum degree of proof tree nodes, consistent with the challenge of going to Multi-Paxos. The proof improvements reduced the maximum level of proof tree nodes by 1 for both Multi-Paxos and Multi-Paxos with Preemption compared with [2].
- The increase in number of lemmas is due to the change from Max in Basic Paxos to MaxBalInSlot in Multi-Paxos, defined in Equation (14). Five lemmas were needed for this predicate alone to aid the prover, as we moved from scalar values to a set of tuples for each acceptor.
- No proof by induction on set increment is used for Basic Paxos. Four such proofs are used for Multi-Paxos and for Multi-Paxos with Preemption.
- Proof by contradiction is used twice in the proof of Basic Paxos, and we extended them with slots in the proof of Multi-Paxos and Multi-Paxos with Preemption. We use contradiction proofs one more time in our proofs in lieu of longer constructive proofs.

- The number of proof obligations for the prover increased most significantly, by 540 (226%), from 239 for Basic Paxos to 779 for Multi-Paxos. Only another 46 (6%) proof obligations were generated for Multi-Paxos with Preemption. However, the number of obligations decreased by 139 (15%) from 918 in [2] to 779 for Multi-Paxos and by 134 (14%), from 959 in [2] to 825 for Multi-Paxos with Preemption.
- The checking time of invariance proof of TypeOK increased by 3 seconds (300%) from 1 for Basic Paxos to 4 for Multi-Paxos due to the more complicated structures involved in Multi-Paxos. It further increased by 8 seconds (200%) from 4 to 12 when Preemption was added. This is due to new aspects in the remedial proof of TypeOK as explained in Section 5.1.
- The checking time of the total proof increased by 135 seconds (338%), from 40 for Basic Paxos to 175 for Multi-Paxos, despite our continuous efforts to help the prover reduce it. This is because of the greatly increased size and complexity of inductions, leading to significantly more obligations for the prover. A small increase of 5 seconds (3%) is observed when Preemption is added. As a result of the proof improvements, we notice a 53 second decrease (23%, from 228 to 175) and a 42 second decrease (19%, from 222 to 180) in the proof checking time for Multi-Paxos and Multi-Paxos with Preemption.

#### 7 Related work and conclusion

We discuss closest related results on verification of Paxos, categorized by the verification techniques.

Model checking. Lamport has written TLA<sup>+</sup> specifications for Basic Paxos and its variants, e.g., Fast Paxos [22], and checked them using the TLA<sup>+</sup> model checker TLC [33], but not for Multi-Paxos or its variants; a number of MS students at our university have also done this in course projects, including for Multi-Paxos. Delano et al. [8] modeled Basic Paxos in Promela and checked it using the Spin model checker [40]. To reduce the state space, they use counting guards to track majority, reset local variables after state operations, and use sorted send instead of FIFO send (with random receive, to model non-FIFO channels). They checked Basic Paxos for pairs of numbers of proposers and acceptors up to (2,8), (3,5), (4,4), (5,3), and (8,2). Yabandeh et al. [44] checked a C++ implementation of Basic Paxos using CrystalBall, a tool built on Mace [16], which includes a model checker. Yang et al. [45] used their model checker MoDist to check a Multi-Paxos-based service system developed by a Microsoft product team [31]. With dynamic partial-order reduction [10], they found 13 bugs including 2 bugs in the Paxos implementation, with as few as 3 replicas and a few slots. In all cases, existing work in model checking either does not check Multi-Paxos or can check it for only a very small number of processes and slots.

**Deductive verification.** Kellomäki [15] formally specified and verified Basic Paxos using PVS [41]. Charron-Bost and Schiper [5] expressed Basic Paxos in the Heard-Of model, and Charron-Bost and Merz [4] verified it formally using Isabelle/HOL [42]. Drăgoi et al. [9] specified and verified a version of Basic Paxos in PSync, which is based on the Heard-Of model, so the specification and proof are similar to [5, 4]. Lamport et al. [25] give a formal specification of Basic Paxos in TLA<sup>+</sup> and a TLAPS-checked proof of its correctness. Lamport [23] wrote a TLA<sup>+</sup> specification of Byzantine Paxos, a variant of Basic Paxos that tolerates arbitrary

Metric	Basic		Multi-Paxos		Multi- w/ Preemption			
Metric	Paxos		Old [2]		New	Old [2]		New
Spec size (lines, excl. comments)	-,	52	-,	56	55	-,	75	74, 52*
Proof size (lines, excl. comments)	-,	310	-,	1010^^	787	-,	1054^^	831,528*
Spec size incl. comments (lines)	115^,	106	133^,	123	115	158^,	151	144, 76*
Proof size incl. comments (lines)	423^,	432	1106^	,1096^^	868	1136^	, 1143^^	915,613*
Max level of proof tree nodes	7		11		10	11		10
Max degree of proof tree nodes	3		17		17	17		17
# lemmas	4		11		11	12		12
# invariance lemmas	1		5		5	6		6
# uses of invariance lemmas	8		27		27	29		29
# proofs by induction on set increment	0		4		4	4		4
# proofs by contradiction	1^,	2	1^,	3	3	1^,	3	3
# obligations in TLAPS	239		918^^		779	959^^		825
TypeOK CPU check time (seconds)	-,	1	-,	10^^	4	-,	12	12
Total CPU check time (seconds)	-,	40	-,	228^^	175	-,	222^^	180
TypeOK elapsed check time (seconds)	-,	1	-,	20^^	7	-,	22^^	24
Total elapsed check time (seconds)	24**,	14	128**,	63^^	51	94**,	69^^	90

Table 2: Summary of results. An obligation is a condition that TLAPS checks. The check time is on a 4-core Intel i7-4720HQ 2.6 GHz CPU with 16 GB of memory, running 64-bit Ubuntu 17.10 and TLAPS 1.5.6.

<sup>-</sup> denotes an entry not in [2], followed by the now added measurement.

<sup>^</sup> indicates a number with a count oversight in [2], followed by the now correct count.

<sup>^^</sup> indicates an incorrect number from the incomplete proof due to the TLAPS bug [29].

<sup>\*\*</sup> indicates a number that used TLAPS version 1.5.3 in [2], followed by the number using the new version 1.5.6.

<sup>\*</sup> indicates a number for the specification or proof in Appendix C, after removing unnecessary line breaks from default latex generated by TLA<sup>+</sup> Tools.

failures, and a TLAPS-checked proof that it refines Basic Paxos. Küfner et al. [18] exhibit a methodology to develop machine-checkable parameterized proofs of correctness of fault-tolerant round-based distributed algorithms with Basic Paxos as a case study. Their proof is approximately 10,000 lines in Isabelle/HOL.

With IronFleet, Hawblitzel et al. [11] verified a state machine replication system that uses Multi-Paxos at its core. Their specification mimics TLA<sup>+</sup> models but is written in Dafny [38], which has no direct concurrency support but has more automated proof support than TLAPS. This work is superior to its peers by proving not only safety but also liveness properties. However, it is a complex system, with 3 levels and many components of specifications, over 1000 lines, and proofs, over 30,000 lines. Schiper et al. [39] used EventML [36] to specify Multi-Paxos and used NuPRL [7] to verify safety. Using the Verdi framework, Wilcox et al. [43] expressed Raft [35], an algorithm similar to Multi-Paxos, in OCAML and verified it using Coq [12]. The proof is over 50,000 lines and takes almost 30 minutes to verify. Padon et al. [32] specify many variants of Paxos including Basic and Multi-Paxos in first-order logic. They present a methodology aiming at automatic verification based on effectively propositional logic (EPR).

All these works either do not handle Multi-Paxos or handle it using more restricted or less direct language models than TLA<sup>+</sup>, some mixed in large systems, making the essence of the algorithm's proof harder to find and understand.

In contrast, our work is the first to specify the exact phases of Multi-Paxos in a most direct and general language model, TLA<sup>+</sup>, with a complete correctness proof automatically checked using TLAPS. Building on Lamport et al.'s specification and proof for Basic Paxos [25], we aim to facilitate the understanding of Multi-Paxos and its proof by minimizing the difference from those for basic Paxos. We further show this as a general way for specifying and proving variants of Multi-Paxos, by doing so for Multi-Paxos extended with preemption. We also discuss the significantly more complex but necessary subproofs by induction. Future work may automate inductive proofs and support the verification of variants that improve and extend Multi-Paxos, by extending specifications of variants of Paxos, e.g., Fast Paxos [22] and Byzantine Paxos [23], to Multi-Paxos and verifying these variants of Multi-Paxos as well as Raft [35].

Acknowledgements. We thank Stephan Merz for his helpful comments on the proofs and explanations of TLAPS. We thank anonymous reviewers for their helpful comments on the paper. We thank Leslie Lamport for his encouragement for this work. This work was supported in part by National Science Foundation grants CCF-1248184, CCF-1414078, CNS-1421893, and IIS-1447549, and Office of Naval Research grant N000141512208. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of these agencies.

#### References

- [1] Burrows, M.: The Chubby lock service for loosely-coupled distributed systems. In: Proceedings of the 7th USENIX Symposium on Operating Systems Design and Implementation. pp. 335–350. USENIX Association (2006)
- [2] Chand, S., Liu Y.A., Stoller, S.D.: Formal verification of multi-Paxos for distributed consensus. In: Proceedings of the 21st International Symposium on FormalMethods, pp. 119–136. Springer (2016)

- [3] Chandra, T.D., Griesemer, R., Redstone, J.: Paxos made live—An engineering perspective. In: Proceedings of the 26th Annual ACM Symposium on Principles of Distributed Computing. pp. 398–407 (2007)
- [4] Charron-Bost, B., Merz, S.: Formal verification of a consensus algorithm in the Heard-Of model. International Journal of Software and Informatics 3(2-3), 273–303 (2009)
- [5] Charron-Bost, B., Schiper, A.: The Heard-Of model: Computing in distributed systems with benign faults. Distributed Computing 22(1), 49–71 (2009)
- [6] Clarke, Jr., E.M., Grumberg, O., Peled, D.A.: Model Checking. MIT Press (1999)
- [7] Constable, R.L., Allen, S.F., Bromley, H.M., Cleaveland, W.R., Cremer, J.F., Harper, R.W., Howe, D.J., Knoblock, T.B., Mendler, N.P., Panangaden, P., Sasaki, J.T., Smith, S.F.: Implementing Mathematics with the Nuprl Proof Development System. Prentice-Hall (1986)
- [8] Delzanno, G., Tatarek, M., Traverso, R.: Model checking Paxos in Spin. In: Proceedings of the 5th International Symposium on Games, Automata, Logics and Formal Verification. pp. 131–146 (2014)
- [9] Drăgoi, C., Henzinger, T.A., Zufferey, D.: Psync: A partially synchronous language for fault-tolerant distributed algorithms. In: Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages. pp. 400–415 (2016)
- [10] Flanagan, C., Godefroid, P.: Dynamic partial-order reduction for model checking software. In: Proceedings of the 32nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages. pp. 110–121 (2005)
- [11] Hawblitzel, C., Howell, J., Kapritsos, M., Lorch, J.R., Parno, B., Roberts, M.L., Setty, S., Zill, B.: IronFleet: Proving practical distributed systems correct. In: Proceedings of the 25th Symposium on Operating Systems Principles. pp. 1–17 (2015)
- [12] INRIA: The Coq Proof Assistant. http://coq.inria.fr/ (Last released January 2016)
- [13] Isard, M.: Autopilot: Automatic data center management. ACM SIGOPS Operating Systems Review 41(2), 60–67 (2007)
- [14] Jamshidi, M., Betancourt, A. J., Gomez, J.: Cyber-physical control of unmanned aerial vehicles. Scientia Iranica 18(3). pp. 663–668. (2011)
- [15] Kellomäki, P.: An annotated specification of the consensus protocol of Paxos using superposition in PVS. Report 36, Institute of Software Systems, Tampere University of Technology (2004)
- [16] Killian, C.E., Anderson, J.W., Braud, R., Jhala, R., Vahdat, A.M.: Mace: language support for building distributed systems. In: Proceedings of the 28th ACM SIGPLAN Conference on Programming Language Design and Implementation. pp. 179–188 (2007)
- [17] Kirsch, J., Goose, S., Amir, Y., Wei, D., Skare, P.: Survivable SCADA via intrusion-tolerant replication. In: IEEE Transactions on Smart Grid 5(1). pp. 60–70. (2014)

- [18] Küfner P., Nestmann U., Rickmann C.: Formal verification of distributed algorithms. In: Theoretical Computer Science. pp. 209-224 (2012)
- [19] Lamport, L.: The part-time parliament. ACM Transactions on Computer Systems 16(2), 133–169 (1998)
- [20] Lamport, L.: Paxos made simple. SIGACT News (Distributed Computing Column) 32(4), 51–58 (2001)
- [21] Lamport, L.: Specifying Systems: The TLA<sup>+</sup> Language and Tools for Hardware and Software Engineers. Addison-Wesley (2002)
- [22] Lamport, L.: Fast Paxos. Distributed Computing 19(2), 79–103 (2006)
- [23] Lamport, L.: Byzantizing Paxos by refinement. In: Proceedings of the 25th International Symposium on Distributed Computing. pp. 211–224. Springer (2011)
- [24] Lamport, L.: My writings. http://research.microsoft.com/en-us/um/people/lamport/pubs/p (Accessed January 24, 2016), Lamport's history of paper [19]
- [25] Lamport, L., Merz, S., Doligez, D.: A TLA<sup>+</sup> specification of the Paxos Consensus algorithm described in Paxos Made Simple and a TLAPS-checked proof of its correctness. https://github.com/tlaplus/v1-tlapm/blob/master/examples/paxos/Paxos.tla (Last modified Fri Nov 28 10:39:17 PST 2014 by Lamport. Accessed Feb 6, 2018)
- [26] Lamport, L.: How to write a 21<sup>st</sup> century proof. In: Journal of Fixed Point Theory and Applications (JFPTA) 11(1). pp. 43–63. Springer (2012)
- [27] Lamport, L.: Time, clocks, and the ordering of events in a distributed system. In: Communications of the ACM 21(7). pp. 558–565. (1978)
- [28] Lamport, L.: The temporal logic of actions. In: ACM Transactions on Programming Languages and Systems (TOPLAS) 16(3). pp. 872–923. (1994)
- [29] Merz, S.: Coalescing bug in TLAPS. https://github.com/tlaplus/tlaplus/issues/40. Accessed March 16, 2018.
- [30] Liskov B.: From viewstamped replication to byzantine fault tolerance. In: Replication. pp. 121–149. Springer Berlin Heidelberg (2010)
- [31] Liu, X., Guo, Z., Wang, X., Chen, F., Lian, X., Tang, J., Wu, M., Kaashoek, M.F., Zhang, Z.: D3S: Debugging deployed distributed systems. In: Proceedings of the 5th USENIX Symposium on Networked Systems Design and Implementation. pp. 423–437. USENIX Association (2008)
- [32] Padon, O., Losa, G., Sagiv, M., Shoham, S.: Paxos made EPR: Decidable Reasoning about Distributed Protocols. In: Proceedings of the ACM on Programming Languages, 1(OOPSLA). pp. 108. (2017)
- [33] Microsoft Research: The TLA Toolbox. http://research.microsoft.com/en-us/um/people/lam (Last modified January 4, 2016)
- [34] Microsoft Research-Inria Joint Center: TLA<sup>+</sup> Proof System (TLAPS). http://tla.msr-inria.inria.fr/tlaps/ (Last released June 2015)

- [35] Ongaro, D., Ousterhout, J.: In search of an understandable consensus algorithm. In: 2014 USENIX Annual Technical Conference (USENIX ATC 14). pp. 305–319. USENIX Association (2014)
- [36] PRL Project: EventML. http://www.nuprl.org/software/#WhatisEventML (Last released September 21, 2012)
- [37] van Renesse, R., Altinbuken, D.: Paxos made moderately complex. ACM Computing Surveys 47(3), 42:1–42:36 (Feb 2015)
- [38] Microsoft Research: Dafny: A language and program verifier for functional correctness. http://research.microsoft.com/en-us/projects/dafny/ (Last released October 12, 2015)
- [39] Schiper, N., Rahli, V., van Renesse, R., Bickford, M., Constable, R.L.: Developing correctly replicated databases using formal tools. In: Proceedings of the 44th Annual IEEE/IFIP International Conference on Dependable Systems and Networks. pp. 395–406. IEEE CS Press (2014)
- [40] Spin Community: Verifying Multi-threaded Software with Spin. http://spinroot.com/spin/whatispin.html (Last released January 1, 2016)
- [41] SRI: PVS Specification and Verification System. http://pvs.csl.sri.com/ (Last released February 11, 2013)
- [42] University of Cambridge: Isabelle (a generic proof assistant). http://isabelle.in.tum.de/ (Last released May 25, 2015)
- [43] Wilcox, J.R., Woos, D., Panchekha, P., Tatlock, Z., Wang, X., Ernst, M.D., Anderson, T.: Verdi: A framework for implementing and formally verifying distributed systems. In: Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation. pp. 357–368 (2015)
- [44] Yabandeh, M., Knezevic, N., Kostic, D., Kuncak, V.: CrystalBall: Predicting and preventing inconsistencies in deployed distributed systems. In: Proceedings of the 6th USENIX Symposium on Networked Systems Design and Implementation. pp. 229–244. USENIX Association (2009)
- [45] Yang, J., Chen, T., Wu, M., Xu, Z., Liu, X., Lin, H., Yang, M., Long, F., Zhang, L., Zhou, L.: MoDist: Transparent model checking of unmodified distributed systems. In: Proceedings of the 6th USENIX Symposium on Networked Systems Design and Implementation. pp. 213–228. USENIX Association (2009)

# A TLA<sup>+</sup> specification of Multi-Paxos with Preemption

```
- MODULE MultiPaxosSpec -
This is a specification in TLA+ and machine checked proof in TLAPS of Multi-Paxos with Preemption.
EXTENDS Integers, TLAPS, FiniteSets, FiniteSetTheorems
CONSTANTS \mathcal{P}, \mathcal{A}, \mathcal{Q}, \mathcal{V}
                                      Sets of proposers, acceptors, quorums of acceptors, and values to propose
                             Set of sent messages
VARIABLES msqs,
                       For each proposer, the current ballot of the proposer
           pBal,
           aBal,
                        For each acceptor, the highest ballot seen by the acceptor
           a Voted
                        For each acceptor, a subset of (ballot, slot, value) triples that the acceptor has voted
ASSUME QuorumAssumption \triangleq \mathcal{Q} \subseteq \text{SUBSET } \mathcal{A} \land \forall Q1, Q2 \in \mathcal{Q} : Q1 \cap Q2 \neq \emptyset
\mathcal{B} \triangleq \mathbb{N}
             Set of ballots
S \triangleq \mathbb{N}
             Set of slots
vars \triangleq \langle msgs, pBal, aBal, aVoted \rangle
Send(m) \triangleq msgs' = msgs \cup \{m\}
```

Phase 1a: For a proposer p, this phase selects some ballot number pBal[p] with which a 1a message has not been sent, and sends it (to all processes).

```
\begin{array}{ll} Phase1a(p) & \triangleq \land \nexists \ m \in msgs: m.type = \text{``la''} \land m.bal = pBal[p] \\ & \land Send([type \mapsto \text{``la''}, from \mapsto p, \ bal \mapsto pBal[p]]) \\ & \land \text{UNCHANGED} \ \langle pBal, \ aBal, \ aVoted \rangle \end{array}
```

Phase 1b: For an acceptor a, if there is a 1a message m with ballot m.bal that is higher than the highest it has seen, a sends a 1b message with m.bal and with the set of highest-numbered triples it has voted for each slot, and it updates the highest ballot it has seen to be m.bal; otherwise it sends a preempt message back with the highest ballot it has seen.

```
\begin{array}{l} Phase1b(a) \; \triangleq \; \exists \; m \in msgs : m.type = \text{``1a''} \land \\ \text{IF } m.bal > aBal[a] \; \text{THEN} \\ \land Send([type \mapsto \text{``1b''}, from \mapsto a, \ bal \mapsto m.bal, \ voted \mapsto aVoted[a]]) \\ \land aBal' = [aBal \; \text{EXCEPT} \; ![a] = m.bal] \\ \land \; \text{UNCHANGED} \; \langle pBal, \ aVoted \rangle \\ \text{ELSE} \\ \land \; Send([type \mapsto \text{``preempt''}, \ to \mapsto m.from, \ bal \mapsto aBal[a]]) \\ \land \; \text{UNCHANGED} \; \langle pBal, \ aBal, \ aVoted \rangle \end{array}
```

Phase 2a: For a proposer p, if there is no 2a message with current ballot pBal[p], and a quorum of acceptors has sent a set S of 1b messages with pBal[p], p sends a 2a message with pBal[p] and a set of proposals PropSV(T), where T is the union of all voted triples in messages in S. PropSV(T) includes MaxSV(T), the set of slot-value pairs with the highest ballot for each slot in T, and NewSV(T), a set of new slot-value pairs for slots not in T.

```
\begin{array}{ll} \mathit{MaxBSV}(T) & \triangleq \{t \in T : \forall \, t2 \in T : t2.\mathit{slot} = t.\mathit{slot} \Rightarrow t2.\mathit{bal} \leq t.\mathit{bal} \} \\ \mathit{MaxSV}(T) & \triangleq \{[\mathit{slot} \mapsto t.\mathit{slot}, \, \mathit{val} \mapsto t.\mathit{val}] : t \in \mathit{MaxBSV}(T) \} \\ \mathit{UnusedS}(T) & \triangleq \{s \in \mathcal{S} : \nexists \, t \in T : t.\mathit{slot} = s \} \\ \mathit{NewSV}(T) & \triangleq \mathit{CHOOSE} \ D \subseteq [\mathit{slot} : \mathit{UnusedS}(T), \, \mathit{val} : \mathcal{V}] : \forall \, d1, \, d2 \in D : d1.\mathit{slot} = d2.\mathit{slot} \Rightarrow d1 = d2 \land D \neq \emptyset \\ \mathit{PropSV}(T) & \triangleq \mathit{MaxSV}(T) \cup \mathit{NewSV}(T) \\ \\ \mathit{Phase2a}(p) & \triangleq \\ \land \nexists \, m \in \mathit{msgs} : (\mathit{m.type} = \text{``2a''}) \land (\mathit{m.bal} = \mathit{pBal}[\mathit{p}]) \\ \land \exists \, \mathit{Q} \in \mathcal{Q}, \, \mathit{S} \subseteq \{\mathit{m} \in \mathit{msgs} : \mathit{m.type} = \text{``1b''} \land \mathit{m.bal} = \mathit{pBal}[\mathit{p}] \} : \\ \land \forall \, \mathit{a} \in \mathit{Q} : \exists \, \mathit{m} \in \mathit{S} : \mathit{m.from} = \mathit{a} \\ \land \, \mathit{Send}([\mathit{type} \mapsto \text{``2a''}, \, \mathit{from} \mapsto \mathit{p}, \, \mathit{bal} \mapsto \mathit{pBal}[\mathit{p}], \, \mathit{propSV} \mapsto \mathit{PropSV}(\mathit{UNION} \ \{\mathit{m.voted} : \mathit{m} \in \mathit{S} \})]) \\ \land \, \mathit{UNCHANGED} \ \langle \mathit{pBal}, \, \mathit{aBal}, \, \mathit{aVoted} \rangle \\ \end{array}
```

Phase 2b: For an acceptor a, if there is a 2a message m with ballot m.bal that is higher than or equal to the highest it has seen, a sends a 2b message with m.bal and m.propSV, updates the highest ballot it has seen to m.bal, and updates set of voted triples using m.propSV; otherwise it sends a preempt message back with the highest ballot it has seen.

```
Phase2b(a) \triangleq \exists m \in msgs : m.type = "2a" \land IF m.bal \geq aBal[a] THEN
```

Preempt: For a proposer p, if there is a preempt message m with ballot m.bal that is higher than p's current ballot, p updates its current ballot to a new ballot that is higher than m.bal and with which no 1a message has been sent.

```
\begin{aligned} & NewBal(b2) \ \triangleq \ \mathsf{CHOOSE} \ b \in \mathcal{B} : b > b2 \land \nexists \ m \in msgs : m.type = \text{``1a''} \land m.bal = b \\ & Preempt(p) \ \triangleq \ \exists \ m \in msgs : \\ & \land m.type = \text{``preempt''} \land m.to = p \land m.bal > pBal[p] \\ & \land pBal' = [pBal \ \mathsf{EXCEPT} \ ![p] = NewBal(m.bal)] \\ & \land \mathsf{UNCHANGED} \ \langle msgs, \ aBal, \ aVoted \rangle \end{aligned}
& Init \ \triangleq msgs = \emptyset \land pBal = [p \in \mathcal{P} \mapsto 0] \land aBal = [a \in \mathcal{A} \mapsto -1] \land aVoted = [a \in \mathcal{A} \mapsto \emptyset]
& Next \ \triangleq \ \lor \exists \ p \in \mathcal{P} : Phase1a(p) \lor Phase2a(p) \lor Preempt(p)
& \lor \exists \ a \in \mathcal{A} : Phase1b(a) \lor Phase2b(a)
& Spec \ \triangleq \ Init \land \Box[Next]_{vars} \end{aligned}
```

# B Safety property to prove for Multi-Paxos with Preemption and invariants used in proof

```
- MODULE MultiPaxosProp
VotedForIn(a, b, s, v) means that acceptor a has sent some 2b message m with m.bal equal to b and some
proposal in m.propSV with slot equal to s and val equal to v. This specifies that acceptor a has voted the
triple \langle b, s, v \rangle.
VotedForIn(a, b, s, v) \triangleq \exists m \in msgs:
 m.type = \text{``2b''} \land m.from = a \land m.bal = b \land \exists d \in m.propSV : d.slot = s \land d.val = v
ChosenIn(b, s, v) means that every acceptor in some quorum Q has voted the triple \langle b, s, v \rangle.
ChosenIn(b, s, v) \triangleq \exists Q \in Q : \forall a \in Q : VotedForIn(a, b, s, v)
Chosen(s, v) means that for some ballot b, ChosenIn(b, s, v) holds.
Chosen(s, v) \triangleq \exists b \in \mathcal{B} : ChosenIn(b, s, v)
Wont Vote In(a, b, s) means that acceptor a has seen a higher ballot than b, and did not and will not vote
any value with b for slot s.
WontVoteIn(a, b, s) \triangleq aBal[a] > b \land \forall v \in \mathcal{V} : \neg VotedForIn(a, b, s, v)
SafeAt(b, s, v) means that no value except perhaps v has been or will be chosen in any ballot lower than b
for slot s.
SafeAt(b, s, v) \triangleq \forall b2 \in 0 ... (b-1) : \exists Q \in Q : \forall a \in Q : VotedForIn(a, b2, s, v) \lor WontVoteIn(a, b2, s)
Safe states that at most one value can be chosen for each slot.
Safe \triangleq \forall v1, v2 \in \mathcal{V}, s \in \mathcal{S} : Chosen(s, v1) \land Chosen(s, v2) \Rightarrow v1 = v2
Messages defines the set of valid messages. TypeOK defines invariants for the types of the variables.
Messages \triangleq [type : \{"1a"\}, from : \mathcal{P}, bal : \mathcal{B}] \cup
             TypeOK \triangleq ("preempt"), to : \mathcal{P}, bal : \mathcal{B}]
TypeOK \triangleq (msgs \subseteq Messages \land IsFiniteSet(msgs) \land pBal \in [\mathcal{P} \rightarrow \mathcal{B}]
             \land aBal \in [A \rightarrow B \cup \{-1\}] \land aVoted \in [A \rightarrow SUBSET \ [bal : B, slot : S, val : V]]
Max(T) selects the largest element in nonempty set T.
Max(T) \triangleq CHOOSE \ e \in T : \forall f \in T : e > f
MaxBalInSlot(T, s) selects, among set of elements in T with slot s, the highest ballot, or -1 if no element
has slot s.
MaxBalInSlot(T, s) \triangleq \text{LET } E \triangleq \{e \in T : e.slot = s\} \text{ IN } \text{IF } E = \emptyset \text{ THEN } -1 \text{ ELSE } Max(\{e.bal : e \in E\})
MsgInv defines properties satisfied by the contents of messages, for 1b, 2a, and 2b messages.
MsgInv1b(m) \triangleq \land m.bal \leq aBal[m.from]
                  \land \forall r \in m.voted : VotedForIn(m.from, r.bal, r.slot, r.val)
                  \land \forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b \in MaxBalInSlot(m.voted, s) + 1 \dots m.bal - 1 \Rightarrow
                      \neg VotedForIn(m.from, b, s, v)
MsqInv2a(m) \triangleq \land \forall d \in m.propSV : SafeAt(m.bal, d.slot, d.val)
                  \land \forall d1, d2 \in m.propSV : d1.slot = d2.slot \Rightarrow d1 = d2
                  \land \forall m2 \in msgs: (m2.type = "2a" \land m2.bal = m.bal) \Rightarrow m2 = m
MsgInv2b(m) \triangleq \land \exists \ m2 \in msgs: m2.type = "2a" \land m2.bal = m.bal \land m2.propSV = m.propSV
                  \land m.bal \leq aBal[m.from]
MsgInv \triangleq \forall m \in msgs : \land (m.type = "1b") \Rightarrow MsgInv1b(m)
```

AccInv defines properties satisfied by the data maintained by the acceptors.

```
\begin{array}{l} AccInv \triangleq \forall \ a \in \mathcal{A}: \\ \land \ aBal[a] = -1 \Rightarrow \ aVoted[a] = \emptyset \\ \land \ \forall \ r \in \ aVoted[a]: \ aBal[a] \geq r.bal \land \ VotedForIn(a,\ r.bal,\ r.slot,\ r.val) \\ \land \ \forall \ b \in \mathcal{B}, \ s \in \mathcal{S}, \ v \in \mathcal{V}: \ VotedForIn(a,\ b,\ s,\ v) \Rightarrow \exists \ r \in \ aVoted[a]: r.bal \geq b \land r.slot = s \\ \land \ \forall \ b \in \mathcal{B}, \ s \in \mathcal{S}, \ v \in \mathcal{V}: \ b > MaxBalInSlot(aVoted[a],\ s) \Rightarrow \neg VotedForIn(a,\ b,\ s,\ v) \end{array}
```

#### Inv is the complete inductive invariant.

 $Inv \triangleq TypeOK \land AccInv \land MsgInv$ 

## C TLAPS checked proof of Multi-Paxos with Preemption

```
- MODULE MultiPaxosProof
The following 2 axioms and 10 lemmas are straightforward consequences of the predicates defined above.
AXIOM MaxInSet \triangleq \forall S \in (SUBSET \mathbb{N}) \setminus \emptyset : Max(S) \in S
AXIOM MaxOnNat \triangleq \forall S \in \text{SUBSET } \mathbb{N} : \nexists s \in S : Max(S) < s
LEMMA MaxOnNatS \triangleq \forall S1, S2 \in (SUBSET \mathbb{N}) \setminus \emptyset : S1 \subseteq S2 \Rightarrow Max(S1) \leq Max(S2) BY MaxInSet
LEMMA MaxBinSType \triangleq \forall S \in SUBSET [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}], s \in \mathcal{S} : MaxBalInSlot(\mathcal{S}, s) \in \mathcal{B} \cup \{-1\}
BY MaxInSet DEF MaxBalInSlot
LEMMA MaxBinSSubsets \triangleq \forall S1, S2 \in SUBSET [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}], s \in \mathcal{S} : S1 \subseteq S2 \Rightarrow
 \begin{array}{l} MaxBalInSlot(S1,\,s) \leq MaxBalInSlot(S2,\,s) \\ \langle 1 \rangle \text{ SUFFICES ASSUME NEW } S1 \in \text{SUBSET } [bal:\mathcal{B},\,slot:\mathcal{S},\,val:\mathcal{V}], \text{ NEW } s \in \mathcal{S}, \\ \text{NEW } S2 \in \text{SUBSET } [bal:\mathcal{B},\,slot:\mathcal{S},\,val:\mathcal{V}], S1 \subseteq S2 \\ \end{array}
                          PROVE MaxBalInSlot(S1, s) \leq MaxBalInSlot(S2, s) OBVIOUS
  \langle 1 \rangle 1. CASE \nexists d \in S1 : d.slot = s
     \langle 2 \rangle 1. MaxBalInSlot(S1, s) = -1 BY \langle 1 \rangle 1 DEF MaxBalInSlot
  \langle 2 \rangle QED BY \langle 2 \rangle1, MaxBinSType DEF \mathcal{B} \langle 1 \rangle2. CASE \exists d \in S1: d.slot = s
    \langle 2 \rangle1. CASE \nexists d \in S2 \setminus S1 : d.slot = s
\langle 3 \rangle1. MaxBalInSlot(S1, s) = MaxBalInSlot(S2, s) BY \langle 2 \rangle1, \langle 1 \rangle2 DEF MaxBalInSlot
    \langle 3 \rangle QED BY \langle 3 \rangle1, MaxBinSType DEF \mathcal{B} \langle 2 \rangle2. CASE \exists d \in S2 \setminus S1: d.slot = s BY \langle 2 \rangle2, \langle 1 \rangle2, MaxBinSType, MaxOnNatS DEF \mathcal{B}, MaxBalInSlot
     \langle 2 \rangle QED BY \langle 2 \rangle 1, \langle 2 \rangle \dot{2}
  \langle \dot{1} \rangle QED BY \langle \dot{1} \rangle \dot{1}, \langle \dot{1} \rangle \dot{2}
LEMMA MaxBinSNoSlot \triangleq \forall S \in SUBSET [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}], s \in \mathcal{S} :
                                              (\nexists d \in S : d.slot = s) \equiv MaxBalInSlot(S, s) = -1
BY MaxInSet DEF MaxBalInSlot, B
LEMMA MaxBinSExists \triangleq \forall S \in SUBSET [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}], s \in \mathcal{S} : MaxBalInSlot(S, s) \in \mathcal{B} \Rightarrow
 \exists d \in S : d.slot = s \land d.bal = MaxBalInSlot(S, s) \\ \langle 1 \rangle \text{ SUFFICES ASSUME NEW } S \in \text{SUBSET } [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}], \\ \text{NEW } s \in \mathcal{S}, MaxBalInSlot(S, s) \in \mathcal{B}
    PROVE \exists d \in S : d.bal = MaxBalInSlot(S, s) \land d.slot = s OBVIOUS \langle 1 \rangle 1. \exists d \in S : d.slot = s BY DEF MaxBalInSlot, \mathcal{B}
     \langle 1 \rangle 2. MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\}) BY \langle 1 \rangle 1 DEF MaxBalInSlot(S, s) = Max(\{d.bal : d \in \{d \in S : d.slot = s\}\})
  \langle 1 \rangle QED BY \langle 1 \rangle 1, \langle 1 \rangle 2, MaxInSet
\nexists \ d \in S: d.bal > \mathit{MaxBalInSlot}(S,\ s) \land d.slot = s \\ \langle 1 \rangle \ \mathsf{SUFFICES} \ \mathsf{ASSUME} \ \mathsf{NEW} \ S \in \mathsf{SUBSET} \ [\mathit{bal}: \mathcal{B},\ \mathit{slot}: \mathcal{S},\ \mathit{val}: \mathcal{V}], \ \mathsf{NEW} \ s \in \mathcal{S}
    PROVE \nexists d \in S : d.bal > MaxBalInSlot(S, s) \land d.slot = s OBVIOUS \langle 1 \rangle 1. CASE \nexists d \in S : d.slot = s BY \langle 1 \rangle 1 \langle 1 \rangle 2. CASE \exists d \in S : d.slot = s
       \begin{array}{c} (2)1. \ \nexists \ b \in \{d.bal: d \in \{d \in S: d.slot = s\}\}: b > MaxBalInSlot(S, s) \\ \text{BY } \langle 1 \rangle 2, \ MaxOnNat \ \ \text{DEF} \ \ MaxBalInSlot, \ \mathcal{B}, \ \mathcal{S} \end{array} 
      \langle 2 \rangle 2. \not\equiv d \in S: (d.slot = s \land \neg (d.bal \leq MaxBalInSlot(S, s))) BY \langle 2 \rangle 1
      \langle 2 \rangle QED BY \langle 2 \rangle 2
  \langle 1 \rangle QED BY \langle 1 \rangle 1, \langle 1 \rangle 2
LEMMA Misc \triangleq \forall S: \land NewSV(S) \in (SUBSET [slot: UnusedS(S), val: V]) \setminus \emptyset
                                        \land \forall t1, t2 \in NewSV(S) : t1.slot = t2.slot \Rightarrow t1 = t2
                                        \land \nexists t1 \in MaxSV(S), t2 \in NewSV(S) : t1.slot = t2.slot
 ⟨1⟩ SUFFICES ASSUME NEW S
PROVE \land NewSV(S) \in (SUBSET [slot : UnusedS(S), val : V]) \setminus \emptyset
                                         \land \forall t1, t2 \in NewSV(S) : t1.slot = t2.slot \Rightarrow t1 = t2
                                         \land \not\equiv t1 \in MaxSV(S), t2 \in NewSV(S) : t1.slot = t2.slot \ OBVIOUS
  \langle 1 \rangle 1. \exists T \in (SUBSET [slot : UnusedS(S), val : V]) \setminus \emptyset : \forall t1, t2 \in T : t1.slot = t2.slot <math>\Rightarrow t1 = t2
           BY DEF UnusedS
  \langle 1 \rangle 2. \ NewSV(S) \in ({\color{blue} SUBSET} \ [slot: UnusedS(S), \ val: \mathcal{V}]) \setminus \emptyset \ {\color{blue} BY} \ \langle 1 \rangle 1 \ \ {\color{blue} DEF} \ \ NewSV(1) 3. \ \forall \ t1, \ t2 \in NewSV(S): t1.slot = t2.slot \Rightarrow t1 = t2 \ {\color{blue} BY} \ \langle 1 \rangle 1 \ \ {\color{blue} DEF} \ \ \ NewSV(S)
  \langle 1 \rangle 4. \not\exists t1 \in MaxSV(S), t2 \in ([slot: UnusedS(S), val: V] \setminus \emptyset) : t1.slot = t2.slot
           BY DEF Max\dot{S}\dot{V}, MaxB\dot{S}V, UnusedS
  \langle 1 \rangle QED BY \langle 1 \rangle 2, \langle 1 \rangle 3, \langle 1 \rangle 4
```

VotedInv asserts that if any acceptor a voted any triple  $\langle b, s, v \rangle$ , then that triple is safe.

```
VotedForIn(a, b, s, v) \Rightarrow SafeAt(b, s, v) \land b \leq aBal[a]
BY DEF VotedForIn, MsgInv, Messages, TypeOK, MsgInv2a, MsgInv1b, MsgInv2b
VotedOnce asserts that if any acceptor a voted triple \langle b, s, v1 \rangle and acceptor a voted triple \langle b, s, v2 \rangle, then
v1 = v2.
LEMMA VotedOnce \triangleq MsgInv \Rightarrow \forall a1, a2 \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v1, v2 \in \mathcal{V}:
                                               VotedForIn(a1, b, s, v1) \land VotedForIn(a2, b, s, v2) \Rightarrow (v1 = v2)
BY DEF MsgInv, VotedForIn, MsgInv2a, MsgInv1b, MsgInv2b
VotedUnion asserts that, in any two 1b messages' voted field, triples with the same ballot and slot have the
same value.
LEMMA VotedUnion \triangleq MsgInv \land TypeOK \Rightarrow \forall m1, m2 \in msgs: m1.type = "1b" \land m2.type = "1b" \Rightarrow
                                                \forall d1 \in m1.voted, d2 \in m2.voted : (d1.bal = d2.bal \land d1.slot = d2.slot) \Rightarrow
                                                  d1.val = d2.val
  \langle 1 \rangle SUFFICES ASSUME MsgInv, TypeOK, NEW m1 \in msgs, NEW m2 \in msgs, m1.type = "1b", m2.type = "1b",
                                             NEW d1 \in m1.voted, NEW d2 \in m2.voted, d1.bal = d2.bal, d1.slot = d2.slot
                            PROVE d1.val = d2.val OBVIOUS
    \begin{array}{lll} \langle 1 \rangle 1. \ \textit{VotedForIn}(m1.\textit{from},\ d1.\textit{bal},\ d1.\textit{slot},\ d1.\textit{val}) \ \text{BY} & \ \text{DEF} \ \textit{MsgInv},\ \textit{MsgInv} 1b \\ \langle 1 \rangle 2. \ \textit{VotedForIn}(m2.\textit{from},\ d2.\textit{bal},\ d2.\textit{slot},\ d2.\textit{val}) \ \text{BY} & \ \text{DEF} \ \textit{MsgInv},\ \textit{MsgInv} 1b \\ \end{array}
  \langle 1 \rangle QED BY \langle 1 \rangle 1, \langle 1 \rangle 2, VotedOnce DEF TypeOK, Messages
The following 5 invariance lemmas assert that, for acceptor a, ballot b, slot s, and value v, and for all phases
and preempt, if VotedForIn(a, b, s, v) holds then VotedForIn(a, b, s, v)' holds; for all except Phase2b, the
inverse also holds.
LEMMA Phase1aVotedForInv \triangleq TypeOK \Rightarrow \forall p \in \mathcal{P} : Phase1a(p) \Rightarrow \forall a \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} :
                                                                  VotedForIn(a, b, s, v) \equiv VotedForIn(a, b, s, v)'
BY DEF VotedForIn, Send, TypeOK, Messages, Phase1a
LEMMA Phase1bVotedForInv \triangleq TypeOK \Rightarrow \forall a \in \mathcal{A} : Phase1b(a) \Rightarrow \forall a2 \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} :
                                                                 VotedForIn(a2, b, s, v) \equiv VotedForIn(a2, b, s, v)
BY DEF VotedForIn, Send, TypeOK, Messages, Phase1b
LEMMA Phase2aVotedForInv \triangleq TypeOK \Rightarrow \forall p \in \mathcal{P} : Phase2a(p) \Rightarrow \forall a \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} :
 \begin{array}{c} VotedForIn(a,\ b,\ s,\ v) \equiv VotedForIn(a,\ b,\ s,\ v)' \\ \text{BY } \forall\ p \in \mathcal{P}: Phase2a(p) \Rightarrow \forall\ m \in msgs' \setminus msgs: m.type = \text{``2a''} \ \text{DEF} \ VotedForIn, \end{array} 
Send, TypeOK, Messages, Phase2a
VotedForIn(a2, b, s, v) \Rightarrow VotedForIn(a2, b, s, v)'
BY DEF VotedForIn, Send, TypeOK, Messages, Phase2b
LEMMA Preempt VotedForInv \triangleq TypeOK \Rightarrow \forall p \in \mathcal{P} : Preempt(p) \Rightarrow \forall a \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} :
                                                                 VotedForIn(a, b, s, v) \equiv VotedForIn(a, b, s, v)'
BY DEF VotedForIn, Send, TypeOK, Messages, Preempt
Invariance lemma SafeAtStable asserts that if SafeAt(b, s, v) holds then SafeAt(b, s, v)' holds in the next
LEMMA SafeAtStable \triangleq Inv \land Next \land TypeOK' \Rightarrow \forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : SafeAt(b, s, v) \Rightarrow SafeAt(b, s, v)' \langle 1 \rangle SUFFICES ASSUME Inv, Next, TypeOK', NEW b \in \mathcal{B}, NEW s \in \mathcal{S}, NEW v \in \mathcal{V}, SafeAt(b, s, v) PROVE SafeAt(b, s, v)' OBVIOUS
(1) USE DEF Send, Inv, B
\langle 1 \rangle 1. CASE \exists p \in \mathcal{P} : Phase1a(p) BY \langle 1 \rangle 1 DEF SafeAt, Phase1a, VotedForIn, WontVoteIn
\langle 1 \rangle 2. CASE \exists a \in \mathcal{A} : Phase1b(a)
           \textbf{BY} \hspace{0.1cm} \langle 1 \rangle 2, \hspace{0.1cm} Quorum Assumption \hspace{0.1cm} \textbf{DEF} \hspace{0.1cm} TypeOK, \hspace{0.1cm} SafeAt, \hspace{0.1cm} WontVoteIn, \hspace{0.1cm} VotedForIn, \hspace{0.1cm} Phase1b 
\langle 1 \rangle 3. ASSUME NEW p \in \mathcal{P}, Phase 2a(p) PROVE Safe At(b, s, v)' \langle 2 \rangle 1. \forall a \in \mathcal{A}, b2 \in \mathcal{B}, s2 \in \mathcal{S} : Wont Vote In(a, b2, s2) \equiv Wont Vote In(a, b2, s2)'
            BY \langle 1 \rangle 3, Phase 2a Voted For Inv DEF Wont Vote In, Send, Phase 2a
   \langle 2 \rangle QED BY \langle 2 \rangle1, QuorumAssumption, Phase2aVotedForInv, \langle 1 \rangle3 DEF SafeAt
 \begin{array}{l} \langle 1 \rangle 4. \text{ ASSUME NEW } a \in \mathcal{A}, \ Phase2b(a) \ PROVE \\ \langle 2 \rangle 1. \ PICK \ m \in msgs: Phase2b(a)!(m) \ BY \ \langle 1 \rangle 4. \ DEF \ Phase2b \\ \langle 2 \rangle 2. \ \forall \ a2 \in \mathcal{A}, \ b2 \in \mathcal{B}: aBal[a2] > b2 \Rightarrow aBal'[a2] > b2 \ BY \ \langle 2 \rangle 1. \ DEF \ TypeOK \\ \langle 2 \rangle 3. \ ASSUME \ NEW \ a2 \in \mathcal{A}, \ NEW \ b2 \in \mathcal{B}, \ NEW \ c2 \in \mathcal{S}, \ neW \ c2 \in \mathcal{
                              VotedForIn(a2, b2, s2, v2)', NEW S \in SUBSET [slot : S \setminus \{s2\}, val : V]
             PROVE FALSE
     \begin{array}{c} \langle 3 \rangle 1. \ \exists \ m1 \in \mathit{msgs'} \setminus \mathit{msgs} : \land m1.\mathit{type} = \text{``2b''} \land m1.\mathit{bal} = \mathit{b2} \land \mathit{m1.from} = \mathit{a2} \\ \land \exists \ \mathit{d} \in \mathit{m1.propSV} : \mathit{d.slot} = \mathit{s2} \land \mathit{d.val} = \mathit{v2} \end{array} 
               BY \langle 2 \rangle 3 DEF VotedForIn, \hat{WontVoteIn}
```

**LEMMA** VotedInv  $\triangleq$  MsgInv  $\land$  TypeOK  $\Rightarrow \forall a \in \mathcal{A}, b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V}$ :

```
\langle 3 \rangle 2. a2 = a \wedge m.bal = b2 BY \langle 2 \rangle 1, \langle 3 \rangle 1 DEF TypeOK
   \begin{array}{c} \langle 3 \rangle \text{ QED BY } \langle 2 \rangle 1, \ \langle 2 \rangle 3, \ \langle 3 \rangle 2, \ \langle 3 \rangle 1 \quad \text{DEF } Phase2b, \ WontVoteIn, \ TypeOK \\ \langle 2 \rangle 4, \ \forall \ a2 \in \mathcal{A}, \ b2 \in \mathcal{B}, \ s2 \in \mathcal{S}: \ WontVoteIn(a2, b2, s2) \Rightarrow \ WontVoteIn(a2, b2, s2)' \end{array} 
    BY \langle 2 \rangle 2, \langle 2 \rangle 3 DEF Wont VoteIn
\langle 2 \rangle QED BY Phase 2bVoted For Inv, \langle 2 \rangle 4, Quorum Assumption, \langle 1 \rangle 4 DEF Safe At \langle 1 \rangle 5. CASE \exists p \in \mathcal{P}: Preempt(p) BY \langle 1 \rangle 5 DEF Safe At, Preempt, Voted For In, Wont Vote In
\langle 1 \rangle QED BY \langle 1 \rangle 1, \langle 1 \rangle 2, \langle 1 \rangle 3, \langle 1 \rangle 4, \langle 1 \rangle 5 DEF Next
Invariant asserts the temporal formula that if Spec holds then Inv always holds.
THEOREM Invariant \triangleq Spec \Rightarrow \Box Inv
\langle 1 \rangle USE DEF \mathcal{B}, \mathcal{S}
 \langle 1 \rangle 2. Inv \wedge [Next]_{vars} \Rightarrow Inv
   \langle 2 \rangle SUFFICES ASSUME Inv, [Next]_{vars}PROVE Inv' OBVIOUS \langle 2 \rangle USE DEF Inv
   \langle 2 \rangle 1. CASE Next
    \langle 3 \rangle1 proves TypeOK' for Next. Each of \langle 4 \rangle1-4 assumes the action of a phase and proves TypeOK' for that
   case.
   \langle 3 \rangle 1. TypeOK'
     \langle 4 \rangle 1. ASSUME NEW p \in \mathcal{P}, Phase 1a(p) PROVE
                                                                                                                                   TypeOK'
    BY \langle 4 \rangle 1, msgs' \setminus msgs \subseteq Messages, FS\_AddElement\ DEF\ Phase1a, TypeOK, Send, Messages\ \langle 4 \rangle 2. ASSUME NEW p \in \mathcal{P}, Phase2a(p)\ PROVE\ TypeOK'
\langle 5 \rangle 1. PICK Q \in \mathcal{Q}, S \in SUBSET\ \{m \in msgs : (m.type = "1b") \land (m.bal = pBal[p])\} : \land \forall a \in Q : \exists m \in S : m.from = a \land Send([type \mapsto "2a", from \mapsto p, bal \mapsto pBal[p], propSV \mapsto PropSV(UNION\ \{m.voted : m \in S\})])\ BY\ \langle 4 \rangle 2 DEF Phase2a
      propSV \mapsto PropSV \text{ (UNION } \{m.voted : m \in S\})] \text{ BY } \langle 4 \rangle 2 \text{ DEF } Phase 2a
\langle 5 \rangle 2. \ UnusedS \text{ (UNION } \{m.voted : m \in S\}) \subseteq S \text{ BY } \langle 5 \rangle 1 \text{ DEF } UnusedS, \ TypeOK, \ Messages
\langle 5 \rangle 3. \ MaxBSV \text{ (UNION } \{m.voted : m \in S\}) \subseteq [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}]
\text{BY DEF } MaxBSV, \ TypeOK, \ Messages
\langle 5 \rangle 4. \land NewSV \text{ (UNION } \{m.voted : m \in S\}) \subseteq [slot : \mathcal{S}, val : \mathcal{V}]
\land MaxSV \text{ (UNION } \{m.voted : m \in S\}) \subseteq [slot : \mathcal{S}, val : \mathcal{V}] \text{ BY } \langle 5 \rangle 3, \ \langle 5 \rangle 2, \ Misc \ DEF \ MaxSV
\langle 5 \rangle 5. \ PropSV \text{ (UNION } \{m.voted : m \in S\}) \subseteq [slot : \mathcal{S}, val : \mathcal{V}] \text{ BY } \langle 5 \rangle 4 \text{ DEF } PropSV
\langle 5 \rangle 6. \ \forall \ m2 \in msgs' \setminus msgs : \land m2.type = \text{"2a"} \land m2.from = p \land m2.bal = pBal[p]
\land m2 \text{ propSV} = PropSV \text{ (UNION } \{m.voted : m \in S\} \rangle
                                                              \wedge m2.propSV = PropSV(UNION \{m.voted : m \in S\})
       BY \langle 5 \rangle 1, \langle 5 \rangle 5 DEF Send, TypeOK, Messages \langle 5 \rangle 7. (msgs \subseteq Messages)' BY \langle 4 \rangle 2, \langle 5 \rangle 6, \langle 5 \rangle 1, \langle 5 \rangle 5, msgs' \setminus msgs \subseteq Messages DEF Phase2a,
                 TypeOK, Send, Messages
     \langle 5 \rangle QED BY \langle 5 \rangle 7, \langle 4 \rangle 2, FS_AddElement DEF Phase2a, TypeOK, Send \langle 4 \rangle 3. ASSUME NEW a \in \mathcal{A}, Phase1b(a) PROVE TypeOK'
      BY \langle 5 \rangle 1 DEF Send
       \langle 5 \rangle 4. (msgs' \subseteq Messages)' BY \langle 5 \rangle 1, \langle 5 \rangle 3 DEF TypeOK, Messages, Send
       \langle 5 \rangle 5. IsFiniteSet(msgs)' BY \langle 5 \rangle 2, FS\_AddElement DEF TypeOK \langle 5 \rangle QED BY \langle 5 \rangle 4, \langle 5 \rangle 5, \langle 4 \rangle 3 DEF Phase1b, TypeOK
     \langle 4 \rangle 4. ASSUME NEW a \in \mathcal{A}, Phase 2b(a) PROVE Type OK'
        \begin{array}{l} \langle 5 \rangle 1. \ \ \text{PICK} \ m \in msgs : Phase2b(a)! (m) \ \text{BY} \ \langle 4 \rangle 4 \ \ \ \text{DEF} \ Phase2b \\ \langle 5 \rangle 2. \lor msgs' = msgs \cup \{[type \mapsto \text{"2b"}, \ bal \mapsto m.bal, \ from \mapsto a, \ propSV \mapsto m.propSV]\} \\ \lor msgs' = msgs \cup \{[type \mapsto \text{"preempt"}, \ to \mapsto m.from, \ bal \mapsto aBal[a]]\} \ \text{BY} \ \langle 5 \rangle 1 \ \ \text{DEF} \ Send \\ \langle 5 \rangle 3. \ \ IsFiniteSet(msgs)' \ \text{BY} \ \langle 5 \rangle 2, \ FS\_AddElement \ \ \text{DEF} \ \ TypeOK \\ \end{array} 
       \langle 5 \rangle 4. \ \forall \ m2 \in \textit{msgs'} \setminus \textit{msgs} : \ \forall \ m2. \textit{type} = \text{``2b''} \land m2. \textit{from} = a \land m2. \textit{bal} = m. \textit{bal} \land m2. \textit{propSV} = m. \textit{propSV}
                                                              \lor m2.type = "preempt" \land m2.to = m.from \land m2.bal = aBal[a]
                 BY \langle 5 \rangle 1 DEF Send
       \langle 5 \rangle 5. (msgs \subseteq Messages)' BY \langle 5 \rangle 1, \langle 5 \rangle 4 DEF TypeOK, Send, Messages
       \langle 5 \rangle 6. \ \forall \ a Voted = a Voted'
                 \vee \wedge DOMAIN \ aVoted = DOMAIN \ aVoted'
                      \land aVoted'[a] = \{d \in aVoted[a] : \nexists d2 \in m.propSV : d.slot = d2.slot\} \cup d
                      \langle 5 \rangle 7. \ (aVoted \in [A \rightarrow SUBSET [bal : B, slot : S, val : V]])' BY \langle 5 \rangle 1, \langle 5 \rangle 6,
                 \{[bal \mapsto m.bal, slot \mapsto d.slot, val \mapsto d.val] : d \in m.propSV\} \subseteq [bal': \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}],
                   \{d \in \mathit{aVoted}[\mathit{a}] : \nexists \, d2 \in \mathit{m.propSV} : \mathit{d.slot} = \mathit{d2.slot}\} \subseteq [\mathit{bal} : \mathcal{B}, \, \mathit{slot} : \mathcal{S}, \, \mathit{val} : \mathcal{V}],
```

 $aVoted'[a] \subseteq [bal : \mathcal{B}, slot : \mathcal{S}, val : \mathcal{V}]$  DEF TypeOK, Messages

```
\langle 5 \rangle QED BY \langle 5 \rangle 5, \langle 5 \rangle 7, \langle 4 \rangle 4, \langle 5 \rangle 3 DEF Phase 2b, Type OK
 \langle 4 \rangle5 proves TupeOK' for Preempt. This is the new part of the remedial proof described in Section 5.1.
 ⟨4⟩5. ASSUME NEW p \in \mathcal{P}, Preempt(p) PROVE TypeOK' ⟨5⟩ DEFINE S \triangleq \{m1 \in msgs : m1.type = "1a"\} T \triangleq \{s.bal : s \in S\}f \triangleq [s \in S \mapsto s.bal] ⟨5⟩ HIDE DEF S, T, f
   \langle 5 \rangle1. PICK m \in msgs: Preempt(p)!(m) BY \langle 4 \rangle5 DEF Preempt \langle 5 \rangle2. T \subseteq \mathcal{B} BY DEF T, S, TypeOK, Messages
   \langle 5 \rangle 3. \; \exists \; b \in \mathcal{B} : b > m.bal \land b \notin T
  BY \langle 5 \rangle 2, MaxInSet, Max(T \cup \{m.bal\}) + 1 > m.bal, \langle 5 \rangle 1 DEF Max, TypeOK, Messages \langle 5 \rangle 4. NewBal(m.bal) \in \mathcal{B} BY \langle 5 \rangle 3 DEF NewBal, TypeOK, Messages, T, S \langle 5 \rangle 5. (pBal \in [\mathcal{P} \to \mathcal{B}])' BY \langle 5 \rangle 1, \langle 5 \rangle 4 DEF TypeOK, Messages \langle 5 \rangle QED BY \langle 4 \rangle 5, \langle 5 \rangle 5 DEF Preempt, TypeOK
 \langle \dot{4} \rangle QED BY \langle \dot{2} \rangle \dot{1}, \langle \dot{4} \rangle \dot{1}, \langle 4 \rangle \dot{2}, \langle 4 \rangle \dot{3}, \langle 4 \rangle \dot{4}, \langle 4 \rangle \dot{5} DEF Next
\langle 3 \rangle2 proves AccInv' for Next. Each of \langle 4 \rangle1-4 assumes the action of a phase and proves AccInv' for that case.
Only Phase2b is challenging because only it updates acceptor variables.
\langle 3 \rangle 2. AccInv'
 \langle 4 \rangle 1. CASE \exists p \in \mathcal{P} : Phase1a(p) BY \langle 4 \rangle 1, \langle 3 \rangle 1, Phase1aVotedForInv DEF AccInv, TypeOK, Phase1a, Send
 \langle 4 \rangle 2. ASSUME NEW p \in \mathcal{P}, Phase 2a(p) PROVE AccInv'
  \langle 5 \rangle 1. \ \forall \ a \in \mathcal{A}, \ b \in \mathcal{B}, \ s \in \mathcal{S}, \ v \in \mathcal{V} : VotedForIn(a, b, s, v) \equiv VotedForIn(a, b, s, v)'
          \textbf{BY} \ \langle 4 \rangle 2, \ Phase 2 \, a \, Voted For Inv
 (5) QED BY (3)1, (4)2, (5)1 DEF AccInv, TypeOK, Phase2a, Send, Messages (4)3. CASE \exists a \in \mathcal{A}: Phase1b(a) BY (4)3, (3)1, Phase1bVotedForInv DEF AccInv, TypeOK, Phase1b, Send (4)4. ASSUME NEW a \in \mathcal{A}, Phase2b(a) PROVE AccInv'
  \langle 5 \rangle SUFFICES ASSUME NEW a2 \in \mathcal{A}'
                      PROVE (\land aBal[a2] = -1 \Rightarrow aVoted[a2] = \emptyset
                                    \land \forall r \in aVoted[a2] : VotedForIn(a2, r.bal, r.slot, r.val) \land r.bal \leq aBal[a2]
                                    \land \forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V}:
                                          \land VotedForIn(a2, b, s, v) \Rightarrow \exists r \in aVoted[a2] : r.bal \geq b \land r.slot = s
                                          \land b > MaxBalInSlot(aVoted[a2], s) \Rightarrow \neg VotedForIn(a2, b, s, v))'
                      BY DEF AccInv
   \langle 5 \rangle 1. PICK m \in msqs : Phase2b(a)!(m) BY \langle 4 \rangle 4 DEF Phase2b
   \langle 5 \rangle 2 assumes that a received a 2a message with a ballot lower than the highest it has seen, thus triggering
   preemption. This case is simple as acceptor variables are unchanged.
   \langle 5 \rangle 2. CASE (a2 = a \land \neg (m.bal \geq aBal[a])) \lor a2 \neq a
     \begin{array}{l} \langle 6 \rangle 1. \ \forall \ b \in \mathcal{B}, \ s \in \mathcal{S}, \ v \in \mathcal{V}: \ VotedForIn(a2, b, s, v) \equiv VotedForIn(a2, b, s, v)' \\ \text{BY } \langle 3 \rangle 1, \ \langle 5 \rangle 1, \ \langle 5 \rangle 2 \ \ \text{DEF} \ Phase2b, \ TypeOK, \ \mathcal{B}, \ Messages, \ VotedForIn, \ Send \\ \langle 6 \rangle \ \text{QED BY } \langle 6 \rangle 1, \ \langle 5 \rangle 1, \ \langle 5 \rangle 2, \ \langle 4 \rangle 4, \ \langle 3 \rangle 1 \ \ \text{DEF} \ Phase2b, \ Send, \ AccInv, \ TypeOK, \ Messages \\ \end{array} 
   \langle 5 \rangle3 assumes that a received a 2a message with a ballot higher than or equal to the highest it has seen.
   Thus, it responds with a 2b message and updates its variables as specified. Each of \langle 6 \rangle 1-4 proves a conjunct
   of AccInv.
   \langle 5 \rangle 3. CASE a2 = a \wedge (m.bal \geq aBal[a])
    \langle 6 \rangle 1. (aBal[a2] = -1 \Rightarrow aVoted[a2] = \emptyset)' BY \langle 5 \rangle 3, \langle 4 \rangle 4, \langle 3 \rangle 1 DEF AccInv, Phase 2b, Send,
            TypeOK, Messages
    \langle 6 \rangle 2. (\forall r \in aVoted[a2] : VotedForIn(a2, r.bal, r.slot, r.val) \land r.bal \leq aBal[a2])'
      \langle 7 \rangle SUFFICES ASSUME NEW r \in (aVoted[a2])'
                         PROVE (VotedForIn(a2, r.bal, r.slot, r.val) \land r.bal \leq aBal[a2])' OBVIOUS
      \langle 7 \rangle 1 uses two cases. \langle 8 \rangle 1 is for r \in aVoted[a2] and uses invariance lemma for Phase 2b. \langle 8 \rangle 2 is for the
      increment r \in aVoted'[a2] \setminus aVoted[a2] and uses definition of Phase2b.
      \langle 7 \rangle 1. VotedForIn(a2, r.bal, r.slot, r.val)'
       \langle 8 \rangle 1. CASE r \in aVoted[a2]
         \langle 9 \rangle1. VotedForIn(a2, r.bal, r.slot, r.val) BY \langle 5 \rangle3, \langle 4 \rangle4, \langle 8 \rangle1 DEF AccInv \langle 9 \rangle QED BY \langle 9 \rangle1, Phase2bVotedForInv, \langle 3 \rangle1, \langle 4 \rangle4 DEF TypeOK, Messages
       \langle \hat{8} \rangle 2. CASE r \in a Voted'[a2] \setminus a Voted[a2]
         \langle 9 \rangle 1. \exists m2 \in msqs' : m2.type = "2b" \land m2.from = a2 \land m2.bal = m.bal \land m2.propSV = m.propSV
                BY \langle 3 \rangle 1, \langle 8 \rangle 2, \langle 5 \rangle 1, \langle 5 \rangle 3 DEF Send
         \langle 8 \rangle QED BY \langle 8 \rangle 1, \langle 8 \rangle 2
      \langle \dot{7} \rangle 2. (r.bal \leq aBal[a2])' BY \langle 5 \rangle 1, \langle 5 \rangle 3, \langle 3 \rangle 1, aBal[a] \leq aBal'[a], r \in aVoted[a] \Rightarrow r.bal \leq aBal'[a],
             aVoted'[a] = \{d \in aVoted[a] : \nexists d2 \in m.propSV : d.slot = d2.slot\} \cup d
             \{[bal \mapsto m.bal, slot \mapsto d.slot, val \mapsto d.val] : d \in m.propSV\}, r \in aVoted'[a2] \setminus aVoted[a2] \Rightarrow r.bal = m.bal
```

```
DEF AccInv, Send, TypeOK, Messages
            \langle 7 \rangle QED BY \langle 7 \rangle 1, \langle 7 \rangle 2
        \langle 6 \rangle 3. \ (\forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : VotedForIn(a2, b, s, v) \Rightarrow \exists \ r \in aVoted[a2] : r.bal \geq b \land r.slot = s)'   \langle 7 \rangle \text{ SUFFICES ASSUME NEW } b \in \mathcal{B}', \text{ NEW } s \in \mathcal{S}', \text{ NEW } v \in \mathcal{V}', \text{ VotedForIn}(a2, b, s, v)'   PROVE \quad (\exists \ r \in aVoted[a2] : r.bal \geq b \land r.slot = s)' \text{ OBVIOUS} 
           \langle 7 \rangle 1. CASE VotedForIn(a2, b, s, v)
              \begin{array}{l} \langle 8 \rangle 1. \ \text{PICK} \ r \in aVoted[a2]: r.bal \geq b \land r.slot = s \ \text{BY} \ \langle 7 \rangle 1 \ \ \text{DEF} \ AccInv, \ TypeOK, \ Messages \\ \langle 8 \rangle 2. \ m.bal \geq b \ \text{BY} \ \langle 8 \rangle 1, \ \langle 5 \rangle 3 \ \ \text{DEF} \ \ TypeOK, \ Messages, \ AccInv \end{array}
             a Voted'[a] = \{d \in a Voted[a] : \nexists d2 \in m.propSV : d.slot = d2.slot\} \cup \{d \in a Voted[a] : \nexists d2 \in m.propSV : d.slot = d2.slot\} \cup \{d \in a Voted[a] : \exists d Voted[a] : \exists d \in a Voted[a] : \exists d Voted
                              \{[bal \mapsto m.bal, slot \mapsto d.slot, val \mapsto d.val] : d \in m.propSV\}, \langle 8 \rangle 3 DEF TypeOK
              \langle 9 \rangle QED BY \langle 9 \rangle 1, \, \langle 5 \rangle 3, \, \langle 8 \rangle 2 \langle 8 \rangle 4. CASE \nexists \, d \in m.propSV: d.slot = s BY \langle 8 \rangle 4, \, \langle 8 \rangle 1, \, \langle 5 \rangle 1, \, \langle 5 \rangle 3, \, r \in aVoted'[a2]
              \langle 8 \rangle QED BY \langle 8 \rangle 3, \langle 8 \rangle 4
           \langle \dot{7} \rangle \dot{2}. CASE \neg \dot{VotedFor}In(a2, b, s, v)
              \langle 8 \rangle 2. \exists r \in aVoted'[a2] : r.bal = m.bal \land r.slot' = s BY <math>\langle 5 \rangle 1, \langle 8 \rangle 1, \langle 2 \rangle 1, \langle 5 \rangle 3 DEF Send, TypeOK,
           \langle 8 \rangle 3. \ b = m.bal \, \mathrm{BY} \ \langle 5 \rangle 1, \ \langle 7 \rangle 2, \ \langle 5 \rangle 3 \ \mathrm{DEF} \ VotedForIn, \ Send \ \langle 8 \rangle \, \mathrm{QED} \, \, \mathrm{BY} \ \langle 8 \rangle 2, \ \langle 8 \rangle 3 \ \langle 7 \rangle \, \, \mathrm{QED} \, \, \mathrm{BY} \ \langle 7 \rangle 1, \ \langle 7 \rangle 2
        \langle 6 \rangle 4. \ (\forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b > MaxBalInSlot(aVoted[a2], s) \Rightarrow \neg VotedForIn(a2, b, s, v))'
           \langle 7 \rangle SUFFICES ASSUME NEW b \in \mathcal{B}', NEW s \in \mathcal{S}', NEW v \in \mathcal{V}', (VotedForIn(a2, b, s, v))',
                                                                   (b > MaxBalInSlot(aVoted[a2], s))'
                                             PROVE FALSE OBVIOUS
          BY MaxBinSNoMore, \langle 3 \rangle 1 DEF TypeOK
           \langle 7 \rangle QED BY \langle 7 \rangle 1, \langle 6 \rangle 3, \exists r \in aVoted'[a2]: r.bal \geq b \wedge r.slot = s, MaxBinSType, \langle 3 \rangle 1 DEF Send,
 TypeOK, Messages

(6) QED BY (6)1, (6)2, (6)3, (6)4

(5) QED BY (5)2, (5)3

(4)5. CASE \exists p \in \mathcal{P} : Preempt(p)

(5)1. \forall a \in \mathcal{A}, s \in \mathcal{S} : MaxBalInSlot(aVoted[a], s) = MaxBalInSlot(aVoted[a], s)'
 BY \langle 3 \rangle 1, \langle 4 \rangle 5 DEF Preempt, MaxBalInSlot \langle 5 \rangle QED BY \langle 4 \rangle 5, \langle 3 \rangle 1, PreemptVotedForInv, \langle 5 \rangle 1 DEF AccInv, TypeOK, Preempt, Send \langle 4 \rangle.QED BY \langle 4 \rangle 1, \langle 4 \rangle 2, \langle 4 \rangle 3, \langle 4 \rangle 4, \langle 4 \rangle 5, \langle 2 \rangle 1 DEF Next
 \langle 3 \rangle 3 proves MsgInv' for Next. Each of \langle 4 \rangle 1-4 assumes the action of a phase and proves MsgInv' for that
 case.
\langle 3 \rangle 3. MsgInv'
  \langle 4 \rangle 1. CASE \exists p \in \mathcal{P} : Phase1a(p) \text{ BY } \langle 4 \rangle 1, Phase1aVotedForInv, <math>\langle 3 \rangle 1, SafeAtStable, \langle 2 \rangle 1
            DEF Phase1a, MsqInv, Send, TypeOK, Messages, MsqInv1b, MsqInv2a, MsqInv2b
   \langle 4 \rangle2 proves MsgInv' for Phase1b. \langle 5 \rangle 13,14 conclude the proof while \langle 5 \rangle 1-12 prove intermediate facts. Each
   of \langle 5 \rangle 2,4,12 proves a conjunct of MsgInv1b for the increment m1—the new 1b message sent in Phase1b(a).
  \langle 4 \rangle 2. ASSUME NEW a \in \mathcal{A}, Phase1b(a) PROVE MsgInv' \langle 5 \rangle DEFINE m1 \triangleq [type \mapsto "1b", from \mapsto a, bal \mapsto m.bal, voted \mapsto aVoted[a]] \langle 5 \rangle 1. PICK m \in msgs: Phase1b(a)!(m) BY \langle 4 \rangle 2 DEF Phase1b
     \langle 5 \rangle 2. \ (m1.bal \leq aBal[m1.from])'  BY \langle 5 \rangle 1, \langle 3 \rangle 1 DEF Phase1b, Send, TypeOK, MsgInv, Messages
     \langle 5 \rangle 3. m1.voted = aVoted[m1.from] BY <math>\langle 5 \rangle 1, \langle 3 \rangle 1 DEF Phase1b, Send, TypeOK, Messages
     \begin{array}{l} \langle 5 \rangle 4. \ (\forall \ r \in m1.voted: VotedForIn(m1.from, \ r.bal, \ r.slot, \ r.val))' \\ \text{BY} \ \langle 5 \rangle 1, \ \langle 4 \rangle 2, \ Phase1b \ VotedForInv, \ \langle 3 \rangle 1, \ \langle 5 \rangle 3 \quad \text{DEF} \quad TypeOK, \ Messages, \ AccInv \ AccIn
     \langle 5 \rangle 5. \ (\forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b > MaxBalInSlot(m1.voted, s) \Rightarrow \neg VotedForIn(m1.from, b, s, v) \rangle
                   \textbf{BY} \ \textit{Phase1bVotedForInv}, \ \langle 4 \rangle 2, \ \langle 5 \rangle 3, \ \langle 5 \rangle 1, \ \langle 3 \rangle 2, \ \langle 3 \rangle 1 \ \ \textbf{DEF} \ \textit{AccInv}, \ \textit{MsgInv}, \ \textit{Send}, \ \textit{TypeOK}, \ \textit{Messages} 
      \langle 5 \rangle 6. \ \forall \ s \in \mathcal{S}: MaxBalInSlot(m1.voted, \ s) \in \mathcal{B} \cup \{-1\} BY \langle 3 \rangle 1, MaxBinSType DEF TypeOK, Messages
      \langle 5 \rangle 7. \ \forall s \in \mathcal{S} : \land MaxBalInSlot(m1.voted, s) + 1 \in \mathcal{B}
                                           \land MaxBalInSlot(m1.voted, s) + 1 > MaxBalInSlot(m1.voted, s)
        \langle 6 \rangle SUFFICES ASSUME NEW s \in \mathcal{S}
                                                                              \land MaxBalInSlot(m1.voted, s) + 1 > MaxBalInSlot(m1.voted, s)
                                           PROVE
                                                                   \land MaxBalInSlot(m1.voted, s) + 1 \in \mathcal{B} \text{ OBVIOUS}
        \langle 6 \rangle 1. CASE MaxBalInSlot(m1.voted, s) = -1 BY \langle 6 \rangle 1
        \langle 6 \rangle 2. CASE MaxBalInSlot (m1.voted, s) \in \mathcal{B} BY \forall x \in \mathcal{B} : x+1 > x, \langle 6 \rangle 2
    \langle 6 \rangle SUFFICES ASSUME NEW b \in \mathcal{B}, NEW s \in \mathcal{S}, b \in MaxBalInSlot(m1.voted, s) + 1 ... <math>m1.bal - 1
```

```
PROVE b > MaxBalInSlot(m1.voted, s) OBVIOUS
   \langle 6 \rangle HIDE DEF m1
   \langle 6 \rangle DEFINE x \triangleq MaxBalInSlot(m1.voted, s) <math>y \triangleq m1.bal - 1
   \begin{array}{l} \langle 6 \rangle 1. \ x \in \mathcal{B} \cup \{-1\} \underset{\mathsf{BY}}{\mathsf{BY}} \ \langle 5 \rangle 6 \\ \langle 6 \rangle 2. \ y \in \mathcal{B} \cup \{-1\} \underset{\mathsf{BY}}{\mathsf{BY}} \ \langle 5 \rangle 9 \end{array} 
   \langle 6 \rangle HIDE DEF x, y
   \langle 6 \rangle 3. CASE x+1>y
    \langle 7 \rangle 1. \ \forall \ e \in \mathcal{B} : e \notin x + 1 \dots y \text{ BY } \langle 6 \rangle 3, \langle 6 \rangle 1, \langle 6 \rangle 2
    \langle 7 \rangle QED BY \langle 6 \rangle 3, \langle 7 \rangle 1 DEF x, y
   \langle 6 \rangle 4. CASE x + 1 = y BY \langle 6 \rangle 4, \langle 5 \rangle 7, \langle 3 \rangle 1, \langle 5 \rangle 9 DEF x, y
   \langle 6 \rangle 5. CASE x + 1 < y
    \langle 7 \rangle 1. \ \forall \ e \in \mathcal{B} : e \in x+1 \dots y \Rightarrow e > x \text{ BY } \langle 6 \rangle 5, \langle 6 \rangle 1, \langle 6 \rangle 2
    \begin{array}{l} \langle 7 \rangle 2. \ b \in x + 1 \dots y \text{ BY DEF } x, y \\ \langle 7 \rangle 3. \ b > x \text{ BY } \langle 7 \rangle 2, \langle 7 \rangle 1 \\ \langle 7 \rangle \text{ QED BY } \langle 7 \rangle 3 \text{ DEF } x \end{array}
   \langle \dot{6} \rangle QED BY \langle \dot{6} \rangle \dot{3}, \langle 6 \rangle 4, \langle 6 \rangle 5, \langle 6 \rangle 1, \langle 6 \rangle 2
 (5)11. \ (\forall b \in \mathcal{B}, s \in \mathcal{S}: b \in MaxBalInSlot(m1.voted, s) + 1...m1.bal - 1 \Rightarrow
            b > MaxBalInSlot(m1.voted, s))' BY \langle 5 \rangle 10, \langle 5 \rangle 1
 \langle 5 \rangle 12. \ (\forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b \in MaxBalInSlot(m1.voted, s) + 1 \dots m1.bal - 1 \Rightarrow
             \neg VotedForIn(m1.from, b, s, v))'
  (6) SUFFICES ASSUME NEW b \in \mathcal{B}', NEW s \in \mathcal{S}', NEW v \in \mathcal{V}'
                     PROVE (b \in MaxBalInSlot(m1.voted, s) + 1 \dots m1.bal - 1 \Rightarrow
                                     \neg VotedForIn(m1.from, b, s, v))' OBVIOUS
  \langle 6 \rangle 1. CASE \nexists x \in \mathcal{B} : x \in (MaxBalInSlot(m1.voted, s) + 1 \dots m1.bal - 1)' BY \langle 6 \rangle 1 \langle 6 \rangle 2. CASE \exists x \in \mathcal{B} : x \in (MaxBalInSlot(m1.voted, s) + 1 \dots m1.bal - 1)' BY \langle 6 \rangle 2, \langle 5 \rangle 5, \langle 5 \rangle 11
   \langle 6 \rangle QED BY \langle 6 \rangle 1, \langle 6 \rangle 2
  \langle 5 \rangle 13 is for 1a message m having a higher ballot than the highest seen, thus generating a 1b message.
 \langle 5 \rangle 13. CASE m.bal > aBal[a]
  \langle 6 \rangle SUFFICES ASSUME NEW m2 \in msqs'
                      PROVE \quad \big( \land (m2.type = "1b") \Rightarrow MsgInv1b(m2) \land (m2.type = "2a") \Rightarrow MsgInv2a(m2) 
                                   \land (m2.type = "2b") \Rightarrow MsgInv2b(m2))' BY DEF MsgInv
   Proves MsgInv1b using two cases. \langle 7 \rangle 1 is for m2 \in msgs. \langle 7 \rangle 2 is for the increment m2 \in msgs' \setminus msgs.
   \langle 6 \rangle 1. \ (m2.type = "1b" \Rightarrow MsgInv1b(m2))'
    \langle 7 \rangle 1. CASE m2 \in msgs BY \langle 7 \rangle 1, \langle 5 \rangle 13, \langle 5 \rangle 1, Phase 1 b Voted For Inv, \langle 4 \rangle 2 DEF MsgInv, MsgInv 1 b,
           TypeOK, Messages
    \langle 7 \rangle 2. CASE m2 \in msgs' \setminus msgs
   Proves MsgInv2a and MsgInv2b using invariance lemma for Phase1b because it does not send 2a or 2b
   \langle 6 \rangle 2. ((m2.type = "2a" \Rightarrow MsgInv2a(m2)) \land (m2.type = "2b" \Rightarrow MsgInv2b(m2)))'
         BY \langle 5 \rangle 13, Phase 1 b Voted For Inv, \langle 5 \rangle 1, \langle 4 \rangle 2, \langle 3 \rangle 1, Safe At Stable, \langle 2 \rangle 1 DEF Send, Type OK,
         MsqInv, Messages, MsqInv2a, MsqInv2b
   \langle 6 \rangle QED BY \langle 6 \rangle 1, \langle 6 \rangle 2
  \langle 5 \rangle 14 is the simple case of preemption and uses the invariance lemmas for Phase 1b and SafeAt.
 \langle 5 \rangle 14. CASE \neg (m.bal > aBal[a]) BY \langle 5 \rangle 14, Phase1bVotedForInv, \langle 4 \rangle 2, \langle 5 \rangle 2, \langle 5 \rangle 4, \langle 5 \rangle 12, \langle 5 \rangle 1,
         SafeAtStable, (3)1, (2)1 DEF Send, TypeOK, MsgInv, Messages, MsgInv1b, MsgInv2a, MsqInv2b
 \langle 5 \rangle QED BY \langle 5 \rangle 13, \langle 5 \rangle 14
\langle 4 \rangle3 proves MsgInv' for Phase2a. Each of \langle 5 \rangle4-6 proves a conjunct of MsgInv.
\langle 4 \rangle 3. ASSUME NEW p \in \mathcal{P}, Phase2a(p) PROVE MsgInv'
 \langle 5 \rangle SUFFICES ASSUME NEW m \in msgs'
                   PROVE (\land (m.type = "1b" \stackrel{\cdot}{\Rightarrow} MsgInv1b(m)) \land (m.type = "2a" \Rightarrow MsgInv2a(m))
                                 \land (m.type = "2b" \Rightarrow MsgInv2b(m)))' BY DEF MsgInv
 \langle 5 \rangle DEFINE b \triangleq pBal[p]
 \langle 5 \rangle1. PICK Q \in \mathcal{Q}, S \in \text{SUBSET} \{ m2 \in msgs : (m2.type = "1b") \land (m2.bal = b) \} : \land \forall a \in Q : \exists m2 \in S : m2.from = a
          \land Send([type \mapsto "2a", bal \mapsto b, from \mapsto p, propSV \mapsto PropSV(UNION \{m2.voted : m2 \in S\})])
        BY \langle 4 \rangle 3 DEF Phase 2a
 \langle 5 \rangle 2.\ b = pBal'[p] \land b \in \mathcal{B}\ \text{BY}\ \langle 4 \rangle 3\ \ \text{DEF}\ Phase 2a,\ Type OK \ \langle 5 \rangle 3.\ \forall\ m2 \in msgs' \setminus msgs: m2.type = "2a" \land m2.bal = b\ \text{BY}\ \langle 5 \rangle 1\ \ \text{DEF}\ Send
 \langle 5 \rangle4 proves MsqInv1b'. It uses invariance lemma for Phase2a, because Phase2a does not send 1b messages.
```

```
\langle 5 \rangle 4. \ (m.type = "1b" \Rightarrow MsgInv1b(m))'
  (6) SUFFICES ASSUME (m.type = "1b")' PROVE MsgInv1b(m)' OBVIOUS
  \langle 6 \rangle 1. \ (m.bal \leq aBal[m.from])' \stackrel{\text{BY}}{\text{BY}} \langle 4 \rangle 3, \langle 5 \rangle 3, \langle 3 \rangle 1, Phase 2 a Voted For Inv DEF Type OK, Messages,
           MsgInv, Phase2a, MsgInv1b
  \langle 6 \rangle 2. \ (\forall r \in m.voted : VotedForIn(m.from, r.bal, r.slot, r.val))' BY \langle 4 \rangle 3, \langle 5 \rangle 3, \langle 3 \rangle 1,
           Phase 2 a Voted For Inv DEF Type OK, Messages, MsgInv, MsgInv1b
  \langle 6 \rangle 4. (\forall b2 \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b2 \in MaxBalInSlot(m.voted, s) + 1 \dots m.bal - 1 \Rightarrow
                \neg VotedForIn(m.from, b2, s, v))'
  BY \langle 4 \rangle 3, \langle 5 \rangle 3, \langle 3 \rangle 1, Phase2aVotedForInv, \langle 6 \rangle 3 DEF TypeOK, Messages, MsgInv, MsgInv1b \langle 6 \rangle QED BY \langle 6 \rangle 1, \langle 6 \rangle 2, \langle 6 \rangle 4 DEF MsgInv1b
 \langle 5 \rangle5 proves MsgInv2a'. Each of \langle 6 \rangle2-4 proves a conjunct of MsgInv2a using the increment approach. The
increment is m2 in \langle 8 \rangle 2 of \langle 7 \rangle 9.
\langle 5 \rangle 5. (m.type = "2a" \Rightarrow MsgInv2a(m))'
  \langle 6 \rangle SUFFICES ASSUME (m.type = \text{``2a''})' PROVE MsgInv2a(m)' OBVIOUS
  \langle 6 \rangle DEFINE VS \triangleq \text{UNION } \{m2.voted : m2 \in S\}
  \begin{array}{l} \langle 6 \rangle 1. \ \forall \ a \in Q: \ aBal[a] \geq b \ \text{BY} \ \langle 5 \rangle 1, \ \langle 3 \rangle 2, \ \langle 3 \rangle 1 \ \ \text{DEF} \ \ MsgInv, \ TypeOK, \ Messages, \ MsgInv1b \\ \langle 6 \rangle 2. \ (\forall \ d \in m.propSV: \ SafeAt(m.bal, \ d.slot, \ d.val))' \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 6 \rangle 2. \ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.val) \\ \langle 7 \rangle 1. \ \forall \ d \in [slot: \ UnusedS(VS), \ val: \ \mathcal{V}] \setminus \emptyset: \ SafeAt(b, \ d.slot, \ d.slot,
      \langle 8 \rangle SUFFICES ASSUME NEW d \in [slot: Unused S(VS), val: V] \setminus \emptyset PROVE SafeAt(b, d.slot, d.val)
      \langle 8 \rangle 1. \ \forall \ m2 \in S: \nexists \ d2 \in m2.voted: d.slot = d2.slot \ \texttt{BY} \ \ \ \texttt{DEF} \ \ UnusedS
      \langle 8 \rangle 2. \ \forall \ m2 \in S: MaxBalInSlot(m2.voted, \ d.slot) + 1 = 0 \ \text{BY} \ \langle 8 \rangle 1 \ \ \text{DEF} \ MaxBalInSlot
      \langle 8 \rangle 4. \ \forall \ v \in \mathcal{V}, \ b2 \in \mathcal{B}, \ a \in Q: b2 \in 0 ... \ b-1 \Rightarrow \neg VotedForIn(a, b2, d.slot, v)
      BY \langle 5 \rangle 1, \langle 8 \rangle 2, \langle 8 \rangle 3 DEF UnusedS, TypeOK, Messages \langle 8 \rangle QED BY \langle 8 \rangle 4, \langle 3 \rangle 1, \langle 6 \rangle 1 DEF SafeAt, NewSV, UnusedS, WontVoteIn, TypeOK, Messages
     \begin{array}{l} \langle 7 \rangle 2. \ \forall \ d \in [slot: UnusedS(VS), \ val: \mathcal{V}] \setminus \emptyset: SafeAt(b, \ d.slot, \ d.val)' \ \text{BY} \ \langle 7 \rangle 1. \ SafeAtStable, \\ \langle 3 \rangle 1, \ \langle 2 \rangle 1, \ \langle 5 \rangle 2. \ \text{DEF} \ NewSV, \ UnusedS, \ TypeOK, \ Messages \\ \langle 7 \rangle 3. \ \forall \ d \in NewSV(VS): SafeAt(b, \ d.slot, \ d.val)' \ \text{BY} \ \langle 7 \rangle 2. \ Misc \ \text{DEF} \ NewSV \\ \end{array} 
    \langle 7 \rangle 4. \ \forall \ d \in MaxBSV(VS) : SafeAt(b, \ d.slot, \ d.val)
      \langle 8 \rangle SUFFICES ASSUME NEW d \in MaxBSV(VS), NEW b2 \in \mathcal{B}, b2 \in 0...(b-1)
                               PROVE \exists Q2 \in \mathcal{Q} : \forall a \in Q2 : \lor VotedForIn(a, b2, d.slot, d.val)
                                                                                          \vee WontVoteIn(a, b2, d.slot) BY DEF SafeAt
      \langle 8 \rangle DEFINE max \triangleq MaxBalInSlot(VS, d.slot)
      (8) USE DEF MaxBSV
      \langle 8 \rangle 1. \ max \in \mathcal{B} \text{ BY } MaxBinSType, MaxBinSNoSlot DEF TypeOK, Messages
      \langle 8 \rangle 2. \ \forall \ m2 \in S: MaxBalInSlot(m2.voted, \ d.slot) \leq max
                BY \forall m2 \in S : m2.voted \subseteq VS, MaxBinSSubsets DEF MaxBSV, TypeOK, Messages
      \langle 8 \rangle 3. \not\exists d2 \in VS : (d2.bal > d.bal \land d2.slot = d.slot)
               BY \forall d2 \in V\dot{S}: \neg(\neg(d2.bal \leq d.bal) \land d2.slot = d.slot) DEF MaxBSV, TypeOK, Messages
      \langle 8 \rangle 4. \ VS \subseteq [bal: \mathcal{B}, slot: \mathcal{S}, val: \mathcal{V}] \ \mathbf{BY} \ \langle 3 \rangle 1 \ \mathbf{DEF} \ TypeOK, Messages
      \langle 8 \rangle 5. \ max = d.bal
        \langle 9 \rangle SUFFICES ASSUME max \neq d.bal PROVE FALSE OBVIOUS
        \langle 9 \rangle 1. CASE max > d.bal
          \langle 10 \rangle HIDE DEF VS
          \langle 10 \rangle 1. \ \exists \ d2 \in VS: d2.bal = max \land d2.slot = d.slot \ BY \ \langle 8 \rangle 4, \langle 8 \rangle 1, MaxBinSExists
          \langle 10 \rangle QED BY \langle 10 \rangle 1, \langle 8 \rangle 3, \langle 9 \rangle 1
        \langle \dot{9} \rangle2. CASE max < d.\dot{bal} BY MaxBinSNoMore, \langle 9 \rangle2 DEF MaxBSV, TypeOK, Messages \langle 9 \rangle QED BY \langle 9 \rangle1, \langle 9 \rangle2, \langle 8 \rangle1 DEF \mathcal{B}, TypeOK, Messages
      \langle 8 \rangle 6. CASE b2 \in max + 1 \cdot b - 1
        \langle 9 \rangle HIDE DEF max
        \langle 9 \rangle 1. \ \forall \ m2 \in S, \ b3 \in \mathcal{B}, \ v \in \mathcal{V}: b3 \in MaxBalInSlot(m2.voted, \ d.slot) + 1 \ldots b - 1 \Rightarrow
                    \neg VotedForIn(m2.from, b3, d.slot, v)
                  BY DEF MsgInv, TypeOK, Messages, MsgInv1b
        \langle 9 \rangle 2. \ \forall \ m2 \in S, \ v \in \mathcal{V} : \neg VotedForIn(m2.from, \ b2, \ d.slot, \ v)
        BY \langle 8 \rangle 6, \langle 9 \rangle 1, \langle 8 \rangle 2, \langle 8 \rangle 1, MaxBinSType DEF TypeOK, Messages, Send \langle 9 \rangle QED BY \langle 5 \rangle 1, \langle 8 \rangle 6, \langle 6 \rangle 1, \langle 9 \rangle 2 DEF MsgInv, MsgInv1b, TypeOK, Messages, WontVoteIn,
              MaxBSV
      \langle 8 \rangle 7. CASE b2 = max
          \langle 9 \rangle 1. \exists a \in \mathcal{A}, m2 \in S : m2.from = a \land \exists d2 \in m2.voted :
                      d2.bal = d.bal \land d2.slot = d.slot \land d2.val = d.val
                    BY DEF MaxBalInSlot, TypeOK, Messages
          \langle 9 \rangle 2. \exists a \in \mathcal{A} : VotedForIn(a, b2, d.slot, d.val)
                    BY \langle 8 \rangle 7, \langle 9 \rangle 1, \langle 8 \rangle 5 DEF MsgInv, TypeOK, Messages, MsgInv1b
          \langle 9 \rangle 3. \ \forall \ q \in Q, \ v2 \in \mathcal{V}: VotedForIn(q, b2, d.slot, v2) \Rightarrow v2 = d.val
```

```
BY (9)2, VotedOnce, QuorumAssumption DEF TypeOK, Messages
           \begin{array}{l} \langle 9 \rangle 4. \ \forall \ q \in Q: aBal[q] > b2 \ \text{BY} \ \langle 5 \rangle 1 \ \ \text{DEF} \ MsgInv, \ TypeOK, \ Messages, \ MsgInv1b \\ \langle 9 \rangle \ \ \text{QED} \ \ \text{BY} \ \langle 8 \rangle 7, \ \langle 9 \rangle 3, \ \langle 9 \rangle 4 \ \ \text{DEF} \ \ WontVoteIn \\ \end{array} 
       (8)8. \text{ CASE } b2 \in 0... max - 1
          \langle 9 \rangle 1. \exists a \in \mathcal{A} : VotedForIn(a, d.bal, d.slot, d.val)
                  BY \langle 8 \rangle 8, \langle 8 \rangle 2 DEF MsgInv, TypeOK, Messages, MsgInv1b
          \langle 9 \rangle2. SafeAt(d.bal, d.slot, d.val) BY \langle 9 \rangle1, VotedInv DEF TypeOK, Messages \langle 9 \rangle QED BY \langle 8 \rangle8, \langle 9 \rangle2, \langle 8 \rangle5 DEF SafeAt, MsgInv, TypeOK, Messages, MaxBalInSlot
      \begin{array}{l} \langle 7 \rangle \text{8.} \ \forall \ d \in \textit{MaxSV}(\textit{VS}) : \textit{SafeAt}(\textit{k}, \ d.\textit{salot}, \ d.\textit{val})' \ \text{BY} \ \langle 7 \rangle 3 \\ \langle 7 \rangle \text{9.} \ (\forall \ m2 \in \textit{msgs} : m2.\textit{type} = "2a") \Rightarrow \forall \ d \in \textit{m2.propSV} : \textit{SafeAt}(\textit{m2.bal}, \ d.\textit{slot}, \ d.\textit{val}))' \\ \langle 8 \rangle \ \text{SUFFICES ASSUME NEW} \ \ m2 \in \textit{msgs'}, \ (m2.\textit{type} = "2a")', \ \text{NEW} \ \ d \in \textit{m2.propSV} \\ \text{PROVE} \ \ \ (\textit{SafeAt}(\textit{m2.bal}, \ d.\textit{slot}, \ d.\textit{val}))' \ \ \text{OBVIOUS} \\ (\text{SafeAt}(\textit{m2.bal}, \ d.\textit{slot}, \ d.\textit{val}))' \ \ \text{OBVIOUS} \\ \end{array} 
       \langle 8 \rangle 1. CASE m2 \in msgs BY \langle 3 \rangle 1, SafeAtStable, \langle 8 \rangle 1, \langle 2 \rangle 1 DEF MsgInv, MsgInv2a, Messages, TypeOK
       \langle 8 \rangle 2. CASE m2 \in msgs' \setminus msgs
         \langle 9 \rangle1. SafeAt(m2.bal, d.slot, d.val)' BY \langle 7 \rangle8, \langle 8 \rangle2, \langle 5 \rangle1, \langle 5 \rangle2 DEF Send, PropSV
      \langle 9 \rangle QED BY \langle 3 \rangle 1, \langle 9 \rangle 1 DEF Send, TypeOK, Messages \langle 8 \rangle QED BY \langle 8 \rangle 1, \langle 8 \rangle 2
     \langle \dot{7} \rangle QED BY \langle \dot{7} \rangle \dot{9}
   The increment is m2 in \langle 7 \rangle 1.
   \langle 6 \rangle 3. \ (\forall d1, d2 \in m.propSV : d1.slot = d2.slot \Rightarrow d1 = d2)'
     \langle 7 \rangle 1. \ \forall \ m2 \in msgs' \setminus msgs : \forall \ d1, \ d2 \in m2.propSV : d1.slot = d2.slot \Rightarrow d1 = d2
       \begin{array}{l} \langle 8 \rangle 1. \ VS \in \textcolor{red}{SUBSET} \ [bal: \mathcal{B}, \ slot: \mathcal{S}, \ val: \mathcal{V}] \ \textcolor{red}{BY} \ \ \textcolor{red}{DEF} \ \ \textit{Messages}, \ \textit{TypeOK} \\ \langle 8 \rangle 2. \ \forall \ r1, \ r2 \in \textit{MaxBSV}(\textit{VS}) : r1.slot = r2.slot \Rightarrow r1.bal = r2.bal \ \textcolor{red}{BY} \ \langle 8 \rangle 1 \ \ \textcolor{red}{DEF} \ \ \textit{MaxBSV} \\ \langle 8 \rangle 3. \ \ \textit{MaxBSV}(\textit{VS}) \subseteq \textcolor{red}{VS} \ \ \textcolor{red}{BY} \ \langle 8 \rangle 1 \ \ \textcolor{red}{DEF} \ \ \textit{MaxBSV} \\ \end{array} 
       \langle 8 \rangle 4. \ \forall \ r1, \ r2 \in Max \overline{B}SV(VS) \ \dot{:} \ \dot{r}1.bal = r2.bal \land r1.slot = r2.slot \Rightarrow r1.val = r2.val
              BY \langle 8 \rangle 3, VotedUnion
       \langle 8 \rangle 5. \ \forall \ r1, \ r2 \in MaxBSV(VS): r1.slot = r2.slot \Rightarrow r1.bal = r2.bal \land r1.val = r2.val
       \begin{array}{c} \mathbf{BY} \ \langle 8 \rangle 4, \ \langle 8 \rangle 2, \ \langle 8 \rangle 3, \ \langle 8 \rangle 1 \\ \langle 8 \rangle 6. \ \forall \ r1, \ r2 \in \mathit{MaxSV}(\mathit{VS}) : r1.\mathit{slot} = r2.\mathit{slot} \Rightarrow r1 = r2 \ \mathbf{BY} \ \langle 8 \rangle 5 \ \ \mathbf{DEF} \ \mathit{MaxSV} \end{array}
     \langle 8 \rangle QED BY \langle 8 \rangle 6, Misc, \langle 5 \rangle 1 DEF PropSV, Send \langle 7 \rangle QED BY \langle 7 \rangle 1 DEF MsgInv, MsgInv2a
   The increment is m1, m2 in \langle 7 \rangle 2.
 (3)1, (4)3 DEF TypeOK, Messages, MsgInv, Phase2a, Send, MsgInv2b
 \langle 5 \rangle QED BY \langle 5 \rangle 4, \langle 5 \rangle 5, \langle 5 \rangle 6
\langle 4 \rangle 4 proves MsgInv' for Phase2b. Each of \langle 5 \rangle 2-4 proves a conjunct of MsgInv.
\langle 4 \rangle 4. ASSUME NEW a \in \mathcal{A}, Phase 2b(a) PROVE MsgInv'
 \langle 5 \rangle SUFFICES ASSUME NEW m \in msgs'
                      PROVE (\land (m.type = "1b") \Rightarrow MsgInv1b(m) \land (m.type = "2a") \Rightarrow MsgInv2a(m)
                                      \land (m.type = "2b") \Rightarrow MsgInv2b(m))' BY DEF MsgInv
 \langle 5 \rangle 1. PICK m1 \in msgs : Phase2b(a)!(m1) BY <math>\langle 4 \rangle 4 DEF Phase2b
  \langle 5 \rangle2 proves MsgInv1b' for Phase2b. Invariance lemmas do not apply because the 3rd conjunct in MsgInv1b
  quantifies over 2b messages negatively -VotedForIn(a, b, s, v) means acceptor a has sent a 2b message
  voting \langle b, s, v \rangle.
 \langle 5 \rangle 2. ((m.type = "1b") \Rightarrow MsqInv1b(m))'
   (\acute{b}) SUFFICES ASSUME (m.type = "\acute{1}b")' PROVE MsgInv1b(m)' OBVIOUS
   \langle 6 \rangle 1. \ (m.bal \leq aBal[m.from] \land \forall r \in m.voted : VotedForIn(m.from, r.bal, r.slot, r.val))'
           BY \langle 5 \rangle 1, \langle 3 \rangle 1, \langle 4 \rangle 4, Phase 2 b Voted For Inv DEF MsgInv, MsgInv1b, Type OK,
           Messages, Send
   \langle 6 \rangle 2. \ (\forall b \in \mathcal{B}, s \in \mathcal{S}, v \in \mathcal{V} : b \in MaxBalInSlot(m.voted, s) + 1 \dots m.bal - 1 \Rightarrow
               \neg VotedForIn(m.from, b, s, v))'
     \langle 7 \rangle SUFFICES ASSUME NEW b \in \mathcal{B}', NEW s \in \mathcal{S}', NEW v \in \mathcal{V}',
                                        (b \in MaxBalInSlot(m.voted, s) + 1 \dots m.bal - 1)'
     PROVE (\neg VotedForIn(m.from, b, s, v))' OBVIOUS (7)1. \neg VotedForIn(m.from, b, s, v) BY (5)1 DEF Send, MsgInv, TypeOK, Messages, MsgInv1b
     \langle 7 \rangle 2. CASE m.from \neq a \lor \neg (m1.bal \ge aBal[a]) BY \langle 5 \rangle 1, \langle 3 \rangle 1, \langle 7 \rangle 2, \langle 7 \rangle 1 DEF VotedForIn,
```

```
TypeOK, Messages, Send
          \langle 7 \rangle 3. CASE m.from = a \land (m1.bal \ge aBal[a])
          \begin{array}{l} \langle 8 \rangle 1. \ \forall \ m2 \in msgs' \setminus msgs : m2.bal = m1.bal \ \text{BY} \ \langle 5 \rangle 1, \ \langle 7 \rangle 3, \ \langle 6 \rangle 1, \ \langle 7 \rangle 3 \ \ \text{DEF} \ \ Send, \ TypeOK \\ \langle 8 \rangle 2. \ \forall \ m2 \in msgs' \setminus msgs : m2.bal \neq b \ \text{BY} \ \langle 5 \rangle 1, \ \langle 7 \rangle 3, \ \langle 6 \rangle 1, \ \langle 7 \rangle 3, \ \langle 8 \rangle 1 \ \ \text{DEF} \ \ TypeOK, \ Messages \\ \langle 8 \rangle \ \ \text{QED BY} \ \ \langle 7 \rangle 3, \ \langle 7 \rangle 1, \ \langle 8 \rangle 2 \ \ \ \text{DEF} \ \ \ VotedForIn, \ TypeOK, \ Messages \\ \langle 7 \rangle \ \ \ \text{QED BY} \ \ \langle 7 \rangle 2, \ \langle 7 \rangle 3 \ \ \ \text{DEF} \ \ \ TypeOK, \ Messages \\ \end{array} 
      \begin{array}{l} \langle 6 \rangle \ _{\rm QED\ BY} \ \langle 6 \rangle 1, \ \langle 6 \rangle 2 \ \ _{\rm DEF} \ MsgInv1b \\ \langle 5 \rangle 3. \ ((m.type="2a") \Rightarrow MsgInv2a(m))' \ {\rm BY} \ SafeAtStable, \ \langle 3 \rangle 1, \ \langle 4 \rangle 4, \ \langle 2 \rangle 1 \ \ _{\rm DEF} \ MsgInv2a, \end{array}
                TypeOK, Messages, Phase2b, Send
       \langle 5 \rangle4 proves MsqInv2b'. It uses two cases: the second case, \langle 6 \rangle2, is for the increment m.
      \langle 5 \rangle 4. ((m.type = "2b") \Rightarrow MsgInv2b(m))'
        \langle 6 \rangle 1. CASE \neg (m1.bal \geq aBal[a]) \lor m \in msgs BY <math>\langle 5 \rangle 1, \langle 3 \rangle 1, \langle 6 \rangle 1 DEF TypeOK, Messages,
                Send, MsgInv, MsgInv2b
        \langle 6 \rangle 2. CASE m1.bal \geq aBal[a] \land m \in msgs' \setminus msgs BY \langle 5 \rangle 1, \langle 3 \rangle 1, \langle 6 \rangle 2 DEF TypeOK, Send, MsgInv2b
        \langle 6 \rangle QED BY \langle 6 \rangle 1, \langle \overline{6} \rangle 2
      \langle 5 \rangle QED BY \langle 5 \rangle 2, \langle 5 \rangle 3, \langle 5 \rangle 4 DEF MsgInv2b
    \langle 4 \rangle5 proves MsgInv' for Preempt. It uses the invariance lemma for Preempt, since Preempt does not send
    \begin{array}{l} \langle 4 \rangle 5. \ \ \textbf{CASE} \ \exists \ p \in \mathcal{P}: Preempt(p) \ \textbf{BY} \ \langle 4 \rangle 5, \ PreemptVotedForInv, \ \langle 3 \rangle 1, \ SafeAtStable, \\ \langle 2 \rangle 1 \ \ \textbf{DEF} \ Preempt, \ MsgInv, \ TypeOK, \ Messages, \ MsgInv1b, \ MsgInv2a, \ MsgInv2b \\ \langle 4 \rangle \ \ \textbf{QED} \ \ \textbf{BY} \ \langle 4 \rangle 1, \ \langle 4 \rangle 2, \ \langle 4 \rangle 3, \ \langle 4 \rangle 4, \ \langle 4 \rangle 5, \ \langle 2 \rangle 1 \ \ \textbf{DEF} \ \ Next \\ \end{array} 
  \langle 3 \rangle QED BY \langle 3 \rangle 1, \langle 3 \rangle 2, \langle 3 \rangle 3 DEF Inv, vars, Next
  \langle 2 \rangle2. CASE UNCHANGED vars BY \langle 2 \rangle2 DEF vars, Inv, TypeOK, AccInv, MsgInv, VotedForIn,
           SafeAt,\ WontVoteIn,\ MaxBalInSlot,\ MsgInv1b,\ MsgInv2a,\ MsgInv2b
\langle 2 \rangle QED BY \langle 2 \rangle 1, \langle 2 \rangle 2
\langle 1 \rangle QED BY \langle 1 \rangle 1, \langle 1 \rangle 2, PTL DEF Spec
Safety asserts that Spec implies that Safe always holds.
THEOREM Safety \triangleq Spec \Rightarrow \Box Safe
\langle 1 \rangle USE DEF \mathcal{B}
\langle 1 \rangle 1. Inv \Rightarrow Safe
  \langle 2 \rangle SUFFICES ASSUME Inv, NEW v1 \in \mathcal{V}, NEW v2 \in \mathcal{V}, NEW s \in \mathcal{S}, NEW b1 \in \mathcal{B}, NEW b2 \in \mathcal{B},
                                       ChosenIn(b1, s, v1), ChosenIn(b2, s, v2), b1 \leq b2
                        PROVE
                                            v1 = v2 BY DEF Safe, Chosen
  \langle 2 \rangle 1. CASE b1 = b2
    \langle 3 \rangle 1. \exists a \in \mathcal{A} : VotedForIn(a, b1, s, v1) \land VotedForIn(a, b1, s, v2)
            BY \langle 2 \rangle 1, QuorumAssumption DEF ChosenIn
    \langle 3 \rangle QED BY \langle 3 \rangle 1, VotedOnce DEF Inv
  \langle 2 \rangle 2. CASE b1 < b2
    \langle \dot{3} \rangle 1. SafeAt(b2, s, v2) BY VotedInv, QuorumAssumption DEF ChosenIn, Inv
    \langle 3 \rangle 2. PICK Q1 \in \mathcal{Q}: \forall a \in Q1: VotedForIn(a, b1, s, v1) BY DEF ChosenIn
    \langle 3 \rangle 3. PICK Q2 \in Q: \forall a \in Q2: VotedForIn(a, b1, s, v2) \lor WontVoteIn(a, b1, s) BY <math>\langle 3 \rangle 1, \langle 2 \rangle 2 DEF SafeAt
    \langle 3 \rangle QED BY \langle 3 \rangle 2, \langle 3 \rangle 3, Quorum Assumption, Voted Once DEF Wont Vote In, Inv
  \langle 2 \rangle QED BY \langle 2 \rangle 1, \langle 2 \rangle 2
\langle 1 \rangle QED BY Invariant, \langle 1 \rangle 1, PTL
```